

All-to-All Communication: Each worker maintains a copy of the model parameters (weights) and processes a subset of the training data. After each iteration, each worker communicates with every other worker and updates its local weights by averaging the gradients from all workers.

Parameter Server: A dedicated server, called the parameter server, stores the global model parameters.

The workers compute gradients for a subset of the training data and send these gradients to the parameter server. The server then updates the global model parameters and sends the updated weights back to the workers.

Ring All-Reduce: Arranges n workers in a logical ring and updates the model parameters by passing messages in a circular fashion. Each worker computes gradients for a subset of the training data, splits the gradients into n equally sized chunks and sends a chunk of the gradients to their neighbors in the ring. Each worker receives the gradient chunks from its neighbors, updates its local parameters, and passes the updated gradient chunks along the ring. After $n-1$ passes, all gradient chunks have been aggregated across workers, and the aggregated chunks are passed along to all workers in the next $n-1$ steps. This is illustrated in Fig.

1.

For each of the distributed training paradigms, fill in the total number of messages sent and the size of each message. Assume that there are n workers and the model has p parameters, with p divisible by n .

Here are the givens by the problem:

All-to-all Num | messages sent: FILL_HERE size of each message p
 Parameter Server | messages sent: $2n$ size of each message: FILL_HERE
 Ring All-REDUCE | messages sent : $n(2(n-1))$ size of each
 message: FILL_HERE

Please explain each answer you fill in

To address the query, we need to determine the total number of messages sent and the size of each message for three distributed training paradigms: All-to-All Communication, Parameter Server, and Ring All-Reduce. The problem provides some values and asks us to fill in the missing ones, given n workers and a model with p parameters, where p is divisible by n . Let's analyze each paradigm systematically, explaining the calculations for each answer.

1. All-to-All Communication

Description: Each of the n workers maintains a copy of the model parameters (p parameters) and processes a subset of the training data. After each iteration, every worker communicates its gradients to every other worker, and each worker updates its local weights by averaging the gradients from all workers.

Given: Size of each message is p .

To Find: Total number of messages sent.

Analysis:

- **Number of messages:** In All-to-All Communication, each worker sends its computed gradients (for all p parameters) to every other worker. There are n workers, and each worker sends messages to $n - 1$ other workers (excluding itself).
 - Total messages = n workers $\times$ ($n - 1$) recipients per worker = $n(n - 1)$.
- **Size of each message:** Each message contains the gradients for the entire model, which has p parameters. Since the problem states the size of each message is p , we interpret this as the number of parameters (or equivalently, the size of the gradient vector, assuming each parameter's gradient is one unit of data, e.g., one float).

Verification: The given size of each message is p , which aligns with sending the full gradient vector of p parameters. The number of messages $n(n - 1)$ accounts for every worker sending to every other worker.

Answer:

- Total number of messages sent: $n(n - 1)$
 - Size of each message: p
-

2. Parameter Server

Description: A dedicated parameter server stores the global model parameters (p parameters). Each of the n workers computes gradients on a subset of the training data, sends these gradients to the parameter server, which updates the global parameters and sends the updated parameters back to the workers.

Given: Total number of messages sent is $2n$.

To Find: Size of each message.

Analysis:

- **Number of messages:** Each of the n workers sends its gradients to the parameter server (that's n messages). The parameter server then sends the updated parameters back to each of the n workers (another n messages).
 - Total messages = n (workers to server) + n (server to workers) = $2n$.
- **Size of each message:** Each worker computes gradients for the entire model (p parameters) and sends these to the parameter server. Thus, the message from each worker to the server contains p parameters (gradients). The parameter server, after updating the global parameters, sends the updated parameter vector (also p parameters) back to each worker. Therefore, every message (both worker-to-server and server-to-worker) contains the full set of p parameters.

Verification: The given number of messages ($2n$) matches our calculation. The size of each message is p , as both the gradients sent to the server and the updated parameters sent back are for the entire model.

Answer:

- Total number of messages sent: $2n$

- Size of each message: p
-

3. Ring All-Reduce

Description: n workers are arranged in a logical ring. Each worker computes gradients for a subset of the training data, splits the gradients into n equally sized chunks, and sends a chunk to its neighbor. Over $n - 1$ passes, each worker receives and aggregates gradient chunks from others, updating its local parameters. After aggregation, the aggregated chunks are distributed back to all workers in another $n - 1$ steps.

Given: Total number of messages sent is $n(2(n - 1))$.

To Find: Size of each message.

Analysis:

- **Number of messages:**
 - The process has two phases: gradient aggregation and distribution of aggregated gradients.
 - **Aggregation phase:** Each worker splits its gradients into n chunks and sends one chunk to its neighbor in each of $n - 1$ steps. In each step, all n workers send one chunk to their neighbor, resulting in n messages per step. Over $n - 1$ steps, the number of messages is $n \times (n - 1)$.
 - **Distribution phase:** After aggregation, each worker has one fully aggregated chunk, which it sends to its neighbor over another $n - 1$ steps to ensure all workers receive all aggregated chunks. This also results in $n \times (n - 1)$ messages.
 - Total messages = $n(n - 1)$ (aggregation) + $n(n - 1)$ (distribution) = $n \times (n - 1) \times 2 = 2n(n - 1)$.
 - The given expression is $n(2(n - 1))$, which simplifies to $2n(n - 1)$, matching our calculation.
- **Size of each message:** The model has p parameters, and each worker's gradients are split into n equal chunks (since p is divisible by n).
 - Size of each chunk = p/n .

- In each step of the aggregation phase, each worker sends one chunk of size p/n to its neighbor. In the distribution phase, each worker sends an aggregated chunk, which is also of size p/n , as it represents the aggregated gradients for p/n parameters.

Verification: The given number of messages $n(2(n - 1)) = 2n(n - 1)$ matches the $n(n - 1)$ messages for aggregation plus $n(n - 1)$ for distribution. The size of each message is p/n , as each message involves one chunk of the gradients (or aggregated gradients).

Answer:

- Total number of messages sent: $2n(n - 1)$
 - Size of each message: p/n
-

Final Answer

- **All-to-All Communication:**
 - Total number of messages sent: $n(n - 1)$
 - Size of each message: p
- **Parameter Server:**
 - Total number of messages sent: $2n$
 - Size of each message: p
- **Ring All-Reduce:**
 - Total number of messages sent: $2n(n - 1)$
 - Size of each message: p/n

Explanations:

- **All-to-All:** Each worker sends its full gradient vector (p parameters) to $n - 1$ other workers, resulting in $n(n - 1)$ messages, each of size p .
- **Parameter Server:** Each of the n workers sends gradients (p parameters) to the server, and the server sends updated parameters (p parameters) back to each worker, totaling $2n$ messages, each of size p .

- **Ring All-Reduce:** Over $n - 1$ steps, each of the n workers sends one gradient chunk per step (aggregation phase), and then sends one aggregated chunk per step (distribution phase), totaling $2n(n - 1)$ messages. Each chunk, containing gradients for p/n parameters, has size p/n .

This completes the analysis for all three paradigms.