

hw10.pdf

403 lines

PDF

You are an AI assistant that can solve Deep Neural Network problems. I will upload a pdf of all the problems, and then give you instructions for which problems to solve. Please make sure to show all your work and calculations in your output (including all equations you use to solve, etc.). If you are uncertain about a problem, please state what you are confused on rather than just hallucinating.

I've received your homework PDF for EECS 182 Deep Neural Networks (Homework 10). I can see it contains problems on:

1. Kernelized Linear Attention (Part II)
2. Coding Question: Hand-Design Transformers
3. Coding Question: Summarization
4. Example Difficulty and Early Exit
5. Read a Research Paper: FaceNet
6. Homework Process and Study Group

Please let me know which specific problem(s) you'd like me to help you solve, and I'll provide detailed solutions with all work shown, including equations and step-by-step calculations.

For the coding questions (problems 2, 3, and 4), I should note that I can help explain the concepts and provide guidance on implementation, but I won't be able to run the actual notebooks or generate screenshots since those require access to your local files and computational environment.



Which problem(s) would you like me to work on?

example_difficulty.i

867 lines

IPYNB

I've uploaded the example difficulty notebook for all the parts in question 4a. please take a moment to familiarize yourself with this notebook to answer the conceptual questions

7:52 PM ↺ ✎ 📄

I've reviewed the example difficulty notebook carefully. It's a fascinating exploration of how different examples in a dataset vary in difficulty for neural networks to classify. Let me summarize the key concepts to help you understand what's happening:

Overview of the Notebook

The notebook explores **example difficulty** using a shapes dataset (circles, squares, rectangles, triangles, hearts, ellipses) with the following methodology:

1. **Train a ResNet-18** to convergence on shape classification
2. **Attach KNN classifiers** at each layer of the network
3. **Define "Prediction Depth"** - the earliest layer where KNN predictions stabilize and remain constant through the final layer
4. **Analyze patterns** in easy vs. hard examples

Key Concepts

What is Example Difficulty?

The notebook defines difficulty through **prediction depth**:

- **Easy examples** have low prediction depth (predictions stabilize early in the network)
- **Hard examples** have high prediction depth (need deeper layers to stabilize predictions)

Why Training to Convergence Matters

The notebook emphasizes training to convergence is critical because:

- We need the network to learn optimal representations at each layer
- Analyzing example difficulty on an undertrained model would give misleading results about what makes examples intrinsically difficult vs. what the model hasn't learned yet

The Hook Mechanism

The notebook uses PyTorch **forward hooks** to capture intermediate activations at each layer without modifying the model architecture directly. This is necessary because ResNet is a pre-built model where we can't directly access internal representations.

Key Findings Shown

1. **Most examples exit early** - The histogram shows many examples have prediction depth 0 or 1, meaning they're classified correctly with minimal computation
2. **Negative correlation between prediction depth and accuracy** - Examples that exit later tend to have lower final accuracy, suggesting they're genuinely harder
3. **Visual differences** - Easy examples tend to be clearer, more centered, and have distinct shapes, while hard examples are often occluded, rotated, or ambiguous
4. **Shape-specific patterns** - Different shapes have different difficulty distributions, with some shapes consistently exiting earlier than others

This is an excellent framework for understanding which examples require more computational resources and which could potentially be classified with early-exit mechanisms to save computation.

Would you like me to help you answer the specific conceptual questions from section 4(a)?

please answer part 4a i)

4a) i) Briefly explain what example difficulty is in your own words. What is the setup for us to analyze example difficulty?

Example difficulty is a measure of how hard it is for a neural network to correctly classify a particular example from a dataset. Not all examples are equally challenging - some images are easy to classify and require minimal computation, while others are inherently more difficult and require the full depth of the network to process.

The Setup for Analyzing Example Difficulty:

1. **Train a deep network to convergence:** We train a ResNet-18 on the shapes dataset until it achieves stable, high performance. Training to convergence is critical because we need the network to have learned optimal representations at each layer.
2. **Train KNN classifiers at each layer:** After the ResNet is trained, we pass all training data through the network and capture the intermediate activations at each layer. We then train a separate KNN classifier ($k=30$) on the activations from each layer to predict the shape class.
3. **Define prediction depth:** For each test example, we pass it through the ResNet and use the KNN classifiers to make predictions based on the activations at each layer. The **prediction depth** is the earliest layer L where the KNN prediction becomes stable - meaning all KNN predictions from layer L  ward remain constant and match the final layer's prediction.
4. **Analyze patterns:** Examples with low prediction depth (predictions stabilize early) are considered "easy," while examples with high prediction depth (need many layers to stabilize) are considered "hard." We can then examine what visual or structural properties distinguish easy from hard examples.

This setup allows us to identify which examples require the full computational power of the deep network versus which could be classified with early-exit mechanisms, providing insights into the intrinsic difficulty of different data points.

Analysis: Claude correctly defines example difficulty as the challenge a neural network faces in making a correct prediction. It also accurately describes the problem approach: using a ResNet-18 followed by KNN classifiers at each layer, and then defining, calculating, and analyzing the prediction depth.

please answer part 4a ii)

4a) ii) What kinds of properties do you think will make an example from this dataset difficult?

Based on the shapes dataset and the visual examples shown in the notebook, several properties likely make an example difficult to classify:

1. **Occlusion/Incompleteness:** Shapes that are partially cut off or have missing parts would be harder to classify because the network has less information to work with. A partially visible circle might be confused with an ellipse or arc.
2. **Small size:** Shapes that occupy only a small portion of the 32x32 image would have fewer pixels containing useful information, making feature extraction more challenging, especially in early layers with larger receptive fields.
3. **Ambiguous positioning:** Shapes that are not centered or are positioned at extreme angles might be harder to classify because the network's learned features may be biased toward centered, canonical orientations.
4. **Rotation and orientation:** Highly rotated shapes (especially triangles and rectangles) could be confused with each other or require deeper layers to recognize invariant features.

5. **Low contrast or faint edges:** Shapes with poor contrast against the background or blurred/faint edges would make boundary detection difficult in early convolutional layers.
6. **Ambiguous cases:** Some shapes are inherently similar to others:
 - Circles vs. ellipses (when the ellipse is nearly circular)
 - Squares vs. rectangles (when aspect ratios are close)
 - Hearts vs. circles (depending on drawing style)
7. **Noise or artifacts:** Images with visual noise, unusual textures, or artifacts would require more sophisticated features to separate signal from noise.
8. **Non-prototypical examples:** Shapes that don't match the "typical" or "average" representation of their class - for example, a very elongated rectangle or an irregular heart shape - would be harder because they're further from the learned class prototypes.

The hard examples likely combine several of these properties, requiring the network to use deeper, more abstract features to make correct classifications.

Analysis: Claude correctly states some of the dataset properties that cause example difficulty including noise and shape, and also states other properties including low contrast edges, rotation, occlusion, and ambiguous positioning.

please answer part 4a iii)

4a) iii) Why do we need to train the ResNet to convergence to be able to analyze the example difficulty?

We need to train the ResNet to convergence for several critical reasons:

1. **Optimal feature representations:** Training to convergence ensures that each layer of the ResNet has learned the best possible feature representations for the task. Early in training, the features at each layer are still developing and don't yet capture the patterns necessary for accurate classification. Only after convergence do we have meaningful, task-relevant features at each layer.
2. **Stable predictions:** Convergence means the network has reached a stable state where predictions are no longer changing significantly. If we analyzed example difficulty on an undertrained model, we wouldn't know whether an example appears "difficult" because it's intrinsically hard or simply because the network hasn't finished learning yet.
3. **Separating intrinsic difficulty from learning artifacts:** An undertrained network might struggle with examples that are actually easy but just happen to belong to patterns it hasn't learned yet. By training to convergence, we ensure we're measuring the **intrinsic difficulty** of examples (properties that make them fundamentally challenging) rather than **artifacts of incomplete training**.
4. **Reliable KNN classifiers:** The KNN classifiers we train on each layer's activations depend on having good, discriminative features. If the ResNet isn't converged, the activations at each layer won't be sufficiently informative, and the KNN predictions will be unreliable. This would corrupt our prediction depth measurements.
5. **Fair comparison across layers:** Convergence ensures that all layers have been optimized through backpropagation. If training is incomplete, deeper layers might not have received enough gradient updates, making it appear as if examples need fewer layers when in reality the deeper layers just haven't learned yet.

In summary, training to convergence is essential to ensure that our analysis reflects genuine properties of the data rather than the incomplete state of the learning process.

Analysis: Claude states that convergence is essential to ensuring that we have a fair comparison across layers, reliable KNN classifiers, stable predictions, and optimal feature representations. It also correctly mentions the importance of convergence in terms of having reliable prediction depth measurements.

please answer part 4a iv)

4a) iv) In PyTorch, what does adding hooks do? Why do we need to use them?

What hooks do:

In PyTorch, **hooks** are functions that can be registered to execute automatically during the forward or backward pass of a neural network. Specifically, **forward hooks** allow us to intercept and capture the intermediate outputs (activations) of specific layers as data flows through the network.

When we register a forward hook on a layer using `layer.register_forward_hook(hook_function)`, the hook function is called automatically whenever that layer processes input during a forward pass. The hook function receives three arguments:


- `module`: the layer itself
- `input`: the input tensor(s) to that layer
- `output`: the output tensor(s) from that layer

This allows us to "spy on" what's happening inside the network without modifying the actual computation.

Why we need to use them:

1. **Access to intermediate representations:** The ResNet model is a pre-built, encapsulated architecture from `torchvision.models`. We don't have direct access to the code that defines the internal layers and their activations. Hooks provide a way to extract intermediate layer outputs without modifying the model's source code.
2. **Non-invasive extraction:** Hooks let us capture activations without changing the model architecture or computation graph. The network continues to function exactly as designed, but we get to observe and store the intermediate values.
3. **Training KNN classifiers:** We need the activations at each layer to train our KNN classifiers. Without hooks, we would have to manually rewrite the ResNet forward pass to return all intermediate outputs, which would be cumbersome and error-prone.
4. **Flexibility:** Hooks can be easily added and removed (`handle.remove()`), allowing us to collect activations only when needed (e.g., on the training data) and remove them afterward to avoid unnecessary memory overhead during normal inference.

In this notebook, hooks are essential because they provide the only practical way to capture the layer-by-layer activations from a pre-built ResNet model, which we then use to train KNN classifiers and analyze prediction depth.



Analysis: Claude correctly states that adding hooks allow us to capture intermediate values in pytorch models, enabling us to have access to the activations of the internal layers.

please answer part 4a v)

4a) v) Why do we train KNN classifiers at each layer of the ResNet? What data are we training them on?

Why we train KNN classifiers at each layer:

We train KNN classifiers at each layer to determine **when the network has extracted sufficient features to correctly classify an example**. The purpose is to measure how the representations evolve through the network and identify the "prediction depth" for each example.

Specifically:

1. **Measure feature quality at each depth:** Each KNN classifier tells us how discriminative the learned features are at that particular layer. If the KNN can correctly classify based on early layer activations, it means those features are already good enough for that example.
2. **Define prediction depth:** By comparing the predictions across all layers, we can determine at which layer the predictions stabilize and remain constant through the final layer. This stability point defines the prediction depth - the minimum computational depth needed for confident classification.
3. **Proxy for difficulty:** The KNN classifiers serve as a simpler, interpretable proxy for understanding when the network has "decided" on a classification. If predictions stabilize early, the example is easy; if they only stabilize in deep layers, the example is

4. **Independent of the ResNet's classifier:** Using KNN classifiers provides an external measure of feature quality that's independent of the ResNet's final fully-connected layer. This helps us understand the representations themselves, not just the final decision boundary.

What data are we training them on:

We train the KNN classifiers on **the training data's activations:**

1. We pass all training examples through the fully-trained ResNet
2. We capture the intermediate activations at each layer using hooks
3. For layer L, we collect the flattened activation vectors from all training examples
4. We train a KNN classifier ($k=30$) using these activation vectors as features and the true shape labels as targets

Then, at test time, we pass test examples through the ResNet, capture their activations at each layer, and use the trained KNN classifiers to make predictions based on those activations. This allows us to analyze which examples can be correctly classified using features from early vs. late layers.



Analysis: Claude correctly states that we are training KNN classifiers at each layer in order to measure the prediction depth for each example. It also correctly mentions that we are training the KNN classifiers on the training data's activations (i.e. intermediate representations in ResNet).

please answer part 4a vi)

4a) vi) Why do you think so many examples exit in the earlier layers?

Several factors explain why many examples exit in the earlier layers:

1. **The dataset is relatively simple:** The shapes dataset consists of basic geometric shapes (circles, squares, rectangles, triangles, hearts, ellipses) on simple backgrounds. These are not complex natural images with intricate textures or subtle distinctions. Many examples have clear, unambiguous shapes that don't require deep semantic understanding.
2. **Low-level features are sufficient:** Early layers of CNNs learn simple features like edges, corners, and basic geometric patterns. For well-formed, prototypical shapes (e.g., a clear, centered circle), these low-level features may be entirely sufficient to distinguish between classes. A circle has smooth, curved edges everywhere, while a square has straight edges with sharp corners - features detectable in early layers.
3. **Strong class separability:** The six shape classes are fundamentally different from each other in terms of basic geometry. Unlike more challenging datasets (e.g., distinguishing between dog breeds), where fine-grained details matter, shapes have distinctive structural properties that create clear decision boundaries even in early feature spaces.
4. **Prototypical examples dominate:** The histogram shows that the majority of examples are likely **prototypical** - centered, properly oriented, clear, and unoccluded. These "textbook" examples of each shape are easy to classify and don't require the sophisticated, invariant features that deeper layers provide.
5. **ResNet architecture:** ResNet is actually quite deep (18 layers) and powerful for this relatively simple task. It's somewhat "overqualified" for classifying basic shapes, meaning even its early layers have more than enough representational power for most examples.
6. **Successful training:** The network trained to convergence with high accuracy, meaning it learned effective features throughout. This successful learning ensures that even early layers have discriminative power for the easier examples.

In essence: Most shapes in the dataset are clear, unambiguous examples where basic geometric features (edges, curves, corners) detected in early layers provide sufficient information for correct classification. Only the ambiguous, occluded, or atypical examples require deeper processing to resolve.

Analysis: Claude correctly states that a lot of examples exit in the earlier layers because the dataset has a lot of "easy" examples (simple dataset, strong class separability, etc.).

please answer part 4a vii)

4a) vii) Explain why the accuracy of the ResNet has a negative relationship with the prediction layer.

The negative relationship between ResNet accuracy and prediction layer reveals a fundamental insight: **examples that require deeper layers to classify are intrinsically more difficult.**

Here's why this relationship exists:

1. **Easy examples are classified correctly early:** Examples with low prediction depth (predictions stabilize in early layers) are the ones where the basic features captured by early convolutional layers are sufficient for correct classification. These are clear, prototypical shapes that the network confidently classifies correctly, resulting in high accuracy for this group.
2. **Hard examples need more computation:** Examples with high prediction depth require the full depth of the network to extract increasingly abstract and invariant features. These are the ambiguous, occluded, rotated, or atypical examples that don't match learned prototypes well. The network needs deeper layers to try to resolve the ambiguity.

3. **Deep layers can't always resolve difficulty:** Even with the full computational power of all layers, some examples remain genuinely difficult. They might be:
 - Inherently ambiguous (e.g., a nearly circular ellipse)
 - Poorly formed or heavily occluded
 - Edge cases that don't fit well into any class
 - Examples where even the final, most abstract features can't definitively separate classes
4. **Prediction depth correlates with intrinsic difficulty:** The fact that predictions take longer to stabilize (high prediction depth) indicates the example has properties that make it fundamentally harder. If the network struggles to form a stable prediction even with deeper features, it's more likely to ultimately classify it incorrectly.
5. **Confidence vs. difficulty trade-off:** Early exits indicate high confidence - the network has seen enough evidence to make a decision that won't change with more processing. Late exits indicate uncertainty and ambiguity that persists through multiple layers, correlating with higher error rates.

In summary: Prediction layer serves as a proxy for example difficulty. Easy examples (low prediction depth) are classified correctly, while difficult examples (high prediction depth) challenge the network throughout its depth and are more likely to be misclassified even after full processing. This validates prediction depth as a meaningful measure of example difficulty.

Analysis: Claude correctly states that there is a negative relationship between the accuracy of the ResNet and the prediction layer because examples that require deeper layers to classify are intrinsically more difficult, which aligns with the answer key stating that the examples that exited earlier were the easiest.

please answer part 4a viii)

4a) viii) Explain the distinction between the hard and easy examples. Does this surprise you?

The distinction between hard and easy examples:

Looking at the visualizations in the notebook:

Easy examples (low prediction depth):

- **Clear and well-defined:** The shapes have sharp, distinct boundaries
- **Centered and properly oriented:** Shapes are positioned centrally in the frame with canonical orientations
- **Prototypical:** They look like "textbook" versions of their respective shapes - a circle is clearly circular, a square has obvious right angles and equal sides
- **High contrast:** Strong distinction between the shape and the background
- **Complete and unoccluded:** The entire shape is visible with no missing parts
- **Standard size:** The shapes occupy a reasonable portion of the image, not too small or too large

Hard examples (high prediction depth):

- **Ambiguous or unclear boundaries:** Edges may be blurry, faint, or poorly defined
- **Off-center or rotated:** Shapes positioned at unusual locations or extreme angles
- **Non-prototypical:** Irregular or atypical versions of the shape (e.g., very elongated rectangles, asymmetric hearts)
- **Small or partial:** Shapes that are tiny or cut off at the image boundaries
- **Low contrast:** Poor separation between shape and background
- **Similar to other classes:** Examples that fall in the ambiguous zone between classes (e.g., an almost-circular ellipse, a square that's nearly a rectangle)
- **Noise or artifacts:** Visual irregularities that make classification harder

Does this surprise you?

No, this is not surprising - and that's actually a good sign! The results are highly intuitive and align with what we would expect from both human perception and machine learning theory:

1. **Matches human intuition:** If we were asked to classify these shapes ourselves, we would find the same examples easy or hard. This suggests the network is learning representations that capture genuine difficulty rather than arbitrary patterns.
2. **Validates the methodology:** The fact that prediction depth correlates with visually apparent difficulty confirms that our approach (using KNN classifiers on intermediate layers) is measuring something meaningful about example difficulty.
3. **Consistent with learning theory:** It makes sense that prototypical, high-quality examples are easier to classify while ambiguous, noisy, or atypical examples are harder. This aligns with concepts like decision boundaries, class overlap, and feature discriminability.
4. **Practical implications:** The unsurprising nature of these results actually makes them more useful - we can trust that prediction depth is a reliable measure of difficulty that could be applied to other domains and datasets.

What might be slightly surprising: The degree to which so many examples exit very early (layer 0-1) might be unexpected, suggesting that for simple datasets like this, even shallow features are remarkably powerful. This has important implications for model efficiency and early-exit architectures.

Analysis: Claude correctly states that the hard examples are smaller than the easier examples, as well as other properties that they tend to be off-center / rotated, have lower contrast, and have more noise. Claude also mentions that this is mostly not surprising since it aligns with what we would expect, but misses the point that it is interesting that the network can distinguish these examples.

please answer part 4a ix)

4a) ix) What kinds of patterns do you notice? Based on the composition of the layers, does it make sense?

Patterns noticed:

1. **Early layer dominance:** The histogram shows a massive spike at prediction layers 0-2, with the vast majority of examples exiting very early. This indicates that for most examples, the initial convolutional layers extract sufficient features.
2. **Exponential decay:** The number of examples at each successive prediction layer drops dramatically. Very few examples require the full depth of the network (layers 9-10).
3. **Shape-specific difficulty patterns:** Different shapes show different distributions:
 - Some shapes (like circles) may consistently exit earlier because they have distinctive features (continuous smooth curves)
 - More complex shapes (like hearts or triangles) may have more varied prediction depths due to greater variability in how they can appear
4. **Layer composition matters:** Looking at the ResNet architecture:
 - **Layer 0 (Input):** Just the raw pixel data
 - **Layers 1-3:** Initial convolution, batch norm, ReLU, max pooling - basic edge and texture detection
 - **Layers 4-7:** Residual blocks extracting increasingly complex patterns

- **Layers 8-10:** Deeper residual blocks, adaptive pooling, and final classification

5. **Clustering at early layers:** The fact that predictions stabilize so early suggests that basic geometric features (edges, corners, curvature) captured in the first few convolutional layers are sufficient for most shape classification.

Does this make sense based on layer composition?

Yes, absolutely! This pattern makes perfect sense:

1. **Early layers detect geometric primitives:** The first few layers (conv1, bn1, relu, maxpool, layer1) are designed to detect basic features like:

- Edges at different orientations
- Corners and junctions
- Simple curves and contours

For geometric shapes, these ARE the defining features! A circle is defined by curved edges, a square by straight edges with 90° corners, a triangle by three straight edges meeting at points.

2. **Residual blocks for refinement:** The residual blocks (layers 4-7) learn more complex combinations and spatial relationships. They are needed for:

- Handling rotated or scaled shapes



- Dealing with occlusions
 - Resolving ambiguous cases
 - Building invariance to position and orientation
3. **Deep layers for hard cases:** Only the genuinely difficult examples (ambiguous, occluded, atypical) require the deep layers (8-10) that learn highly abstract, invariant representations. These layers handle:
- Extreme rotations or deformations
 - Partial visibility
 - Distinguishing between very similar shapes (circle vs. ellipse)
4. **Overparameterization for simple tasks:** ResNet-18 has 11+ million parameters and was designed for ImageNet (1000 classes of complex natural images). For 6 simple geometric shapes, it's significantly overparameterized. This means even the early layers have more than enough capacity to solve the task for most examples.
5. **Progressive abstraction matches progressive difficulty:** The architecture's design (simple features → complex features → abstract features) naturally aligns with example difficulty (easy → moderate → hard). Easy examples only need simple features, hard examples need the full abstraction hierarchy.

Conclusion: The patterns we observe directly reflect the hierarchical feature learning of CNNs. Early layers capture the geometric primitives that define most shapes, while deeper layers are reserved for the minority of examples requiring sophisticated invariance and disambiguation. This validates both the architecture design and our difficulty analysis method.

Analysis: While Claude correctly mentions that a lot of the examples exit early on, it does not detect that many examples also exit after layer 8.

early_exit.ipynb

925 lines

IPYNB

Okay, I am uploading one more notebook that you will need to complete all the TODO comments for, please take some time to read through it.

I've carefully reviewed the early exit notebook. This is an excellent practical application of the example difficulty concepts from the previous notebook! Let me summarize the key concepts:

Overview of Early Exit Networks

The notebook implements **BranchyNet**, an early exit architecture that allows examples to "exit" the network at different depths based on their difficulty, rather than always processing through all layers.

Key Concepts:

1. **Modified Loss Function:** Instead of just the final layer loss, the network optimizes a weighted sum of losses from multiple exit points:
 - $L_{\text{early_exit}} = \sum w_i * L(\hat{y}_{\text{exit}_i}, y)$
 - The notebook suggests $w_0 = 1.0$ and $w_i = 0.3$ for $i > 0$
2. **Entropy-Based Exit Criterion:** The network uses entropy of the output distribution as a confidence measure:
 - Low entropy $\rightarrow$ high confidence $\rightarrow$ exit early
 - High entropy $\rightarrow$ low confidence $\rightarrow$ continue processing

3. **Computational Savings:** By allowing easy examples to exit early, we save computation (measured in MACs - Multiply-Accumulate operations)

Architecture Details:

The `EarlyExitResNet18` has multiple exit points (4 total based on the `exiting_points` dictionary), each with its own classifier. During inference:

- After each exit point, compute the output and its entropy
- If entropy < threshold, return prediction and stop
- Otherwise, continue to next exit point

The TODOs to Complete:

1. **Standard ResNet Training:** Call the model with `early_exit=False` and `entropy_tolerance` parameter
2. **Standard ResNet Evaluation:** Same forward pass setup
3. **Early Exit Training:** Forward pass that returns all exit outputs for the weighted loss computation



4. **Early Exit Evaluation:** Forward pass with `entropy_tolerance=0.05` and `early_exit=True`

Key Metrics:

- **Accuracy:** Should remain similar ($\geq 90\%$)
- **MACs:** Should decrease significantly with early exit
- **Exit distribution:** Shows how many examples exit at each point

This demonstrates how understanding example difficulty enables computational efficiency - easy examples don't need the full network depth!

Would you like me to help you fill in the TODO sections, or do you have specific questions about the implementation?



write the code for the standard training

Here's the code for the standard ResNet training (the first TODO section):

```
python

step = 0
resnet_losses = []
for epoch in tqdm(range(10)): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(resnet_trainloader, 0):
        step += 1
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = inputs, labels = data[0].to(device), data[1].to(device)

        inputs = inputs.unsqueeze(1)
        inputs = inputs.repeat(1, 3, 1, 1)
        inputs = inputs.to(device)

        # zero the parameter gradients
        resnet_optimizer.zero_grad()

        # forward + backward + optimize
        # Pass entropy_tolerance=0 and early_exit=False for standard training
        outputs = resnet(inputs, entropy_tolerance=0, early_exit=False)
        loss = criterion(outputs, labels)
        loss.backward()
        resnet_optimizer.step()
```

```

resnet_optimizer.step()
resnet_losses.append(loss.item())
# print statistics
running_loss += loss.item()
if i % 50 == 49:    # print every 2000 mini-batches
    print(f'[{epoch + 1}, {i + 1}] loss: {running_loss / 50:.3f}')
    running_loss = 0.0

print('Finished Training')

```

Key points:

- `entropy_tolerance=0`: This parameter doesn't really matter when `early_exit=False`, but we set it to 0
- `early_exit=False`: This tells the model to only use the final output and ignore all intermediate exit points
- The model returns a single output tensor  final layer predictions)

- The model returns a single output tensor (the final layer predictions)
- We use standard cross-entropy loss on just the final output
- This is baseline training - no early exit behavior, just training the network normally

Analysis: Claude had a couple of errors in this code, so I pasted in the error to get it to get the correct answer.

```

---> 20 outputs = resnet(inputs, entropy_tolerance=0, early_exit=False) 21
loss = criterion(outputs, labels) 22 loss.backward() File
~/miniconda3/envs/zero/lib/python3.9/site-
packages/torch/nn/modules/module.py:1553, in
Module._wrapped_call_impl(self, *args, **kwargs) 1551 return
self._compiled_call_impl(*args, **kwargs) # type: ignore[misc] 1552 else: ->
1553 return self._call_impl(*args, **kwargs) File
~/miniconda3/envs/zero/lib/python3.9/site-
packages/torch/nn/modules/module.py:1562, in Module._call_impl(self, *args,
Show more

```


I see the issue! The model's forward method doesn't accept those keyword arguments in that format. Looking at the notebook more carefully, I need to check how the `EarlyExitResNet18` model is designed to be called.

Based on the evaluation code that works, it seems the model expects the input and then positional arguments. Let me correct the training code:

```
python

step = 0
resnet_losses = []
for epoch in tqdm(range(10)): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(resnet_trainloader, 0):
        step += 1
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = inputs, labels = data[0].to(device), data[1].to(device)

        inputs = inputs.unsqueeze(1)
        inputs = inputs.repeat(1, 3, 1, 1)
        inputs = inputs.to(device)

        # zero the parameter gradients
```

```

# zero the parameter gradients
resnet_optimizer.zero_grad()

# forward + backward + optimize
# For standard training: (inputs, entropy_tolerance, early_exit)
outputs = resnet(inputs, 0, False)
loss = criterion(outputs, labels)
loss.backward()
resnet_optimizer.step()
resnet_losses.append(loss.item())
# print statistics
running_loss += loss.item()
if i % 50 == 49: # print every 2000 mini-batches
    print(f'[{epoch + 1}, {i + 1}] loss: {running_loss / 50:.3f}')
    running_loss = 0.0

print('Finished Training')

```

Key change:

- Use positional arguments: `resnet(input ↓, False)` instead of keyword arguments
- The signature appears to be: `forward(x, entropy_tolerance, early_exit=False)`

This should match the format used in the working evaluation code where it calls `get_model_macs(resnet, (inputs, 0, False))`.

now write the standard resnet evaluation code

Here's the code for the standard ResNet evaluation:

10:39 PM   

Dec 10, 2025 at 10:39 PM

python



```
resnet.eval()
total_macs = 0

for epoch in range(1): # loop over the dataset multiple times

    total_correct = 0
    with torch.no_grad():
        for i, data in tqdm(enumerate(resnet_testloader, 0)):
            # get the inputs; data is a list of [inputs, labels]

            inputs, labels = inputs, labels = data[0].to(device), data[1].to(device)

            inputs = inputs.unsqueeze(1)
            inputs = inputs.repeat(1, 3, 1, 1)
            inputs = inputs.to(device)

            # forward + backward + optimize
            # For standard evaluation: (inputs, entropy_tolerance, early_exit)
            outputs = resnet(inputs, 0, False)
            total_macs += get_model_macs(resnet, (inputs, 0, False))

            indices = torch.argmax(outputs, dim=1)

            total_correct += torch.sum(labels == indices)

    total_correct += torch.sum(labels == indices)

print(f'Accuracy: {total_correct/6000} %')
print('Total MACS: ', total_macs)
```

Key points:

- `resnet.eval()` : Set model to evaluation mode
- `with torch.no_grad()` : Disable gradient computation for efficiency
- `outputs = resnet(inputs, 0, False)` : Standard forward pass (no early exit)
- `get_model_macs(resnet, (inputs, 0, False))` : Calculate MACs for each batch
- Accumulate accuracy and total MACs across all test examples
- This gives us baseline metrics to compare against early exit performance

now write the early exit training code



Analysis: Claude was able to get this code correct in one-shot which was nice to see!

Here's the code for the early exit training:

```
python

step = 0
resnet_losses = []
for epoch in tqdm(range(10)): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(resnet_trainloader, 0):
        step += 1
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = inputs, labels = data[0].to(device), data[1].to(device)

        inputs = inputs.unsqueeze(1)
        inputs = inputs.repeat(1, 3, 1, 1)
        inputs = inputs.to(device)

        # zero the parameter gradients
        resnet_optimizer.zero_grad()

        # forward + backward + optimize
        # For early exit training, we need all exit outputs
        # The model should return a list/tuple of outputs from all exit points
        all_outputs = resnet_early(input ↓ 0, False)
```



```

# Compute weighted loss: w_0 = 1.0, w_i = 0.3 for i > 0
# all_outputs should be a list of tensors, one for each exit point
loss = criterion(all_outputs[0], labels) # Final exit with weight 1.0
for exit_output in all_outputs[1:]: # Intermediate exits with weight 0.3
    loss += 0.3 * criterion(exit_output, labels)

loss.backward()
resnet_optimizer.step()
resnet_losses.append(loss.item())
# print statistics
running_loss += loss.item()
if i % 50 == 49: # print every 2000 mini-batches
    print(f'[{epoch + 1}, {i + 1}] loss: {running_loss / 50:.3f}')
    running_loss = 0.0

print('Finished Training')

```

Key points:

- `all_outputs = resnet_early(inputs, 0, False)`: During training with `early_exit=False`, the model should return all exit point outputs (likely a list or tuple)

- **Weighted loss:** Following the BranchyNet formula: $L = w_0 * L_{final} + \sum (w_i * L_{exit_i})$
 - Final exit (index 0): weight = 1.0
 - All intermediate exits (indices 1+): weight = 0.3
- This trains all exit points simultaneously, allowing the network to make good predictions at multiple depths
- The weighted loss encourages both good final predictions and good intermediate predictions

Analysis: The initial code led to some errors, so I pasted the error to get Claude to correct its mistake.

```
Cell In[12], line 21  18 # forward + backward + optimize  19 # TODO Use w_0
= 1 and w_i = 0.3 for i > 0  20 outputs = resnet_early(inputs, 0, False) ----> 21
loss = criterion(outputs[0], labels)  22 loss.backward()  23
resnet_optimizer.step() File ~/miniconda3/envs/zero/lib/python3.9/site-
packages/torch/nn/modules/module.py:1553, in
Module._wrapped_call_impl(self, *args, **kwargs) 1551 return
self._compiled_call_impl(*args, **kwargs) # type: ignore[misc] 1552 else: ->
1553 return self._call_impl(*args, **kwargs) File
~/miniconda3/envs/zero/lib/python3.9/site-
Show more
```

The error suggests that `outputs` is not returning what we expected. The dimension mismatch `(got input: [6], target: [256])` indicates that `outputs` is a tensor with shape `[6]` (6 classes) rather than `[batch_size, 6]`.

It seems like when `early_exit=False` during training, the model might be returning something different. Looking at the working output in the notebook, it seems the model might return outputs differently. Let me check the architecture file reference...

Actually, looking at the error and the notebook pattern, I think during training with `early_exit=False`, the model should return outputs that need to be handled differently. Let me revise:



python



```
step = 0
resnet_losses = []
for epoch in tqdm(range(10)): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(resnet_trainloader, 0):
        step += 1
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = inputs, labels = data[0].to(device), data[1].to(device)

        inputs = inputs.unsqueeze(1)
        inputs = inputs.repeat(1, 3, 1, 1)
        inputs = inputs.to(device)

        # zero the parameter gradients
        resnet_optimizer.zero_grad()

        # forward + backward + optimize
        # For early exit training, we need to process each sample individually
        # to get all exit outputs, or the model returns a list of all exit outputs
        all_outputs = resnet_early(inputs, 0, False)

        # Check if all_outputs is a list/tuple of outputs from different exits
        if isinstance(all_outputs, (list, tuple)):
```



```
        # Check if all_outputs is a list/tuple of outputs from different exits
        if isinstance(all_outputs, (list, tuple)):
            # Compute weighted loss: w_0 = 1.0 for final exit, w_i = 0.3 for earlier
            loss = criterion(all_outputs[-1], labels) # Final exit with weight 1.0
            for exit_output in all_outputs[:-1]: # Intermediate exits with weight 0.3
                loss += 0.3 * criterion(exit_output, labels)
        else:
            # If it returns a single tensor, just use standard loss
            loss = criterion(all_outputs, labels)

        loss.backward()
        resnet_optimizer.step()
        resnet_losses.append(loss.item())
        # print statistics
        running_loss += loss.item()
    if i % 50 == 49: # print every 2000 mini-batches
        print(f'[epoch + 1], {i + 1}] loss: {running_loss / 50:.3f}')
        running_loss = 0.0
```



```
print('Finished Training')
```


However, without seeing the `architectures.py` file, it's hard to know exactly what the model returns. Could you share what's in the `EarlyExitResNet18` class, or try adding a debug print to see what shape/type `outputs` is?

```
python

outputs = resnet_early(inputs, 0, False)
print(f"Type: {type(outputs)}, Shape/Length: {outputs.shape if hasattr(outputs, 'sha
```

Analysis: After this fix, the code was able to run!

now write the early exit evaluation code

```
python

resnet_early.eval()
exiting_points = {0:0, 1:0, 2:0, 3:0}
entropies = []
total_early_macs = 0

for epoch in range(1): # loop over the dataset multiple times

    total_early_correct = 0
    with torch.no_grad():
        for i, data in tqdm(enumerate(resnet_testloader, 0)):
            # get the inputs; data is a list of [inputs, labels]

            inputs, labels = inputs, labels = data[0].to(device), data[1].to(device)

            inputs = inputs.unsqueeze(1)
            inputs = inputs.repeat(1, 3, 1, 1)
            inputs = inputs.to(device)

            # forward + backward + optimize
            # Use an entropy tolerance of 0.05 and early_exit=True
            outputs, num, curr_entropy = resnet_early(inputs, 0.05, True)
            entropies.append(curr_entropy)
            total_early_macs += get_model_macs(resnet_early, (inputs, 0.05, True))
```

```

total_early_macs += get_model_macs(resnet_early, (inputs, 0.05, True))

exiting_points[num] += 1

indices = torch.argmax(outputs, dim=1)

total_early_correct += torch.sum(labels == indices)

print(f'Accuracy: {total_early_correct/6000} %')
print('Num Exiting: ', exiting_points)
print('Total MACs: ', total_early_macs)
entropies = sorted(entropies)
print(len(entropies))
print('Entropies: ', entropies[0], entropies[len(entropies)//4], entropies[len(entropies)-1])

```

Key points:

- `resnet_early.eval()` : Set model to evaluation mode
- `resnet_early(inputs, 0.05, True)` :
 - `0.05` is the entropy tolerance threshold
 - `True` enables early exit behavior
- `True` enables early exit behavior
- Returns: `(outputs, exit_point_number, entropy_value)`
- `outputs` : The prediction from whichever exit point was used
- `num` : Which exit point this batch used (0, 1, 2, or 3)
- `curr_entropy` : The entropy value at the exit point
- Track which exit points are being used via `exiting_points` dictionary
- Calculate total MACs (should be lower than standard ResNet)
- Collect entropy values to analyze the distribution

This allows us to see the computational savings (reduced MACs) while maintaining accuracy through adaptive depth based on example difficulty!



Analysis: This code ran, but unfortunately, Claude wasn't able to pinpoint the bug in this code, since it outputted the same number of MACs as the standard ResNet. However, in the following questions I asked about the code, it tried to reason about possible bugs.

early_exit.ipynb
973 lines

IPYNB

here is the completed notebook with all outputs after the code was ran,
please look at this for a second and I will ask you some questions:

10:47 PM ↺ ✎ 📄

I've reviewed the completed notebook with all the outputs. Here are the key observations:

Standard ResNet Results:

- **Accuracy:** 93.33%
- **Total MACs:** 3,336,213,504,000
- **Training:** Converged well over 10 epochs (final loss ~0.003-0.005)

Early Exit ResNet Results:

- **Accuracy:** 93.33% (same as standard!)



- **Accuracy:** 93.33% (same as standard!)
- **Total MACs:** 3,336,213,504,000 (same as standard!)
- **Exit distribution:** All 47 batches exited at point 3 (the final exit)
- **Entropy range:** 0.0002 to 0.357

Key Issues I Notice:

1. **No early exit benefit:** The ratio is 1.0, meaning there's NO computational savings. All examples went through the entire network (exit point 3).
2. **Why this happened:** With entropy tolerance of 0.05, even the median entropy (0.0376) is below the threshold, but ALL batches still exited at the final point. This suggests:
 - The early exit mechanism might not be working as intended during inference
 - OR the model wasn't properly trained with the weighted loss (notice the training loss pattern is different/unstable in epochs 5-6)
3. **Training issue:** There's a suspicious spike in loss at epoch 6 (jumps from 0.001 to 0.882), suggesting the training might have issues with the weighted loss implementation or the model state.
4. **Potential bug:** The line `resnet_early = net.to(device)` should probably be a fresh `EarlyExitResNet18()` instance, not reusing the already-trained `resnet` model.

I'm ready to answer your questions about this notebook!

i. What was the accuracy with a regular ResNet? Inference Speed? Total MACS?

4b) i) What was the accuracy with a regular ResNet? Inference Speed? Total MACS?

Based on the output from the completed notebook:

Accuracy: 93.33% (5,600 out of 6,000 test examples correctly classified)

Inference Speed: The notebook processed 47 batches in approximately 3 seconds (based on the progress bar showing "47it [00:03, 12.21it/s]"), which gives us:

- **~12.21 batches/second**
- **~1,562 images/second** (since batch size is 128 for testing)

Total MACs: 3,336,213,504,000 (approximately  **.34 trillion MACs**)

Summary:

- The regular ResNet-18 achieved very high accuracy (93.33%) on this relatively simple shapes dataset
- It processed the test set relatively quickly at about 12 batches per second
- However, it required over 3.3 trillion multiply-accumulate operations to process all 6,000 test examples, which represents the full computational cost of running every example through the entire network depth



Analysis: Claude got this completely correct!

What was the accuracy with an early exit ResNet-18? Inference Speed? Total MACS?

4b) ii) What was the accuracy with an early exit ResNet-18? Inference Speed? Total MACS?

Based on the output from the completed notebook:

Accuracy: 93.33% (same as the regular ResNet)

Inference Speed: The notebook processed 47 batches in approximately 4 seconds (based on the progress bar showing "47it [00:04, 10.26it/s]"), which gives us:

- **~10.26 batches/second**
- **~1,313 images/second** (since batch size is 128 for testing)

Total MACs: 3,336,213,504,000 (approximately **3.34 trillion MACs** - exactly the same as standard ResNet!)



Exit Point Distribution:

- Exit 0: 0 batches
- Exit 1: 0 batches
- Exit 2: 0 batches
- Exit 3: 47 batches (100% of batches exited at the final layer)

Key Observation: The early exit ResNet provided **NO computational savings** in this implementation. All examples went through the entire network (exit point 3), resulting in:

- Same MACs as standard ResNet
- Actually slightly **slower** inference speed (10.26 vs 12.21 batches/sec) due to the overhead of checking entropy at each exit point
- Same accuracy maintained

This indicates that either the entropy threshold (0.05) was too strict, or there may be an implementation issue preventing examples from exiting early despite having low entropy values (median entropy was 0.038, which is below the 0.05 threshold).

Analysis: Given the bug, Claude still correctly parsed the outputs from the notebook and answered the questions using those, but wasn't able to completely resolve the issue.

iii. How did early exit do? Compare accuracy and MACS.

4b) iii) How did early exit do? Compare accuracy and MACS.

Accuracy Comparison:

- **Standard ResNet:** 93.33%
- **Early Exit ResNet:** 93.33%
- **Result:** Identical accuracy - no degradation from early exit mechanism

MACs Comparison:

- **Standard ResNet:** 3,336,213,504,000 MACs
- **Early Exit ResNet:** 3,336,213,504,000 MACs
- **Ratio:** 1.0 (no savings)
- **Result:** No computational savings achieved

Overall Performance:



Early exit did NOT perform as intended. Despite maintaining the same accuracy, it failed to provide any computational benefits because:

1. **All examples exited at the final layer** (exit point 3): 100% of the 47 test batches went through the entire network depth
2. **No early exits occurred:** Despite having examples with very low entropy (minimum 0.0002, median 0.038), none exited early even with a tolerance of 0.05
3. **Added overhead:** The inference was actually slightly slower (10.26 vs 12.21 batches/sec) due to the computational overhead of:
 - Computing predictions at multiple exit points
 - Calculating entropy at each exit point
 - Checking the exit condition

Why This Happened: The implementation likely has issues - either:

- The early exit logic isn't working correctly during inference
- The model wasn't properly trained with the weighted loss (note the unstable training with spike at epoch 6)
- The entropy threshold might need adjustment, though the entropy values suggest examples should have exited early



Ideal Scenario: A properly working early exit network should show:

- Similar accuracy (~90%+)
- Significantly reduced MACs (e.g., 40-60% of standard ResNet)
- Examples distributed across multiple exit points based on difficulty

Analysis: Despite the bug, Claude was able to correctly identify what the properly working early exit network should have looked like.

iv. What was the lowest MACs you found. What might this say?

4b) iv) What was the lowest MACs you found. What might this say?

Lowest MACs Found: 3,336,213,504,000 MACs (the same as the standard ResNet - no improvement)

What This Might Say:

1. **Implementation Issue:** The most likely explanation is that there's a bug or problem in the early exit implementation:
 - Despite low entropy values (median of 0.038, below the 0.05 threshold), no examples exited early
 - The exit mechanism may not be correctly implemented in the forward pass
 - The model assignment `resnet_early = resnet.to(device)` reuses the standard trained model rather than creating a fresh early exit model
2. **Training Problem:** The early exit network wasn't properly trained with the weighted loss:
 - The unstable training (loss spike at ↓ h 6: 0.001 → 0.882) suggests the weighted loss may not have worked correctly

- All exit heads need to be trained simultaneously to make good predictions, which may not have happened
3. **Entropy Threshold Too Strict:** The 0.05 threshold might be too low for this model's confidence calibration:
- Even though median entropy is 0.038, the network might not be exiting early due to how entropy is calculated or compared
 - Adjusting to higher thresholds (e.g., 0.1, 0.2, 0.3) might enable early exits
4. **The Concept Still Has Potential:** Despite this implementation not working, the entropy distribution (0.0002 to 0.357) shows that examples do have varying confidence levels, suggesting:
- Early exit *should* work with proper implementation
 - There's theoretical potential for computational savings if the mechanism functioned correctly
 - The problem is with the specific implementation, not the early exit concept itself

What Should Have Happened: With a properly working early exit network, we should see MACs reduced to perhaps 40-70% of the baseline (e.g., 1.3-2.3 trillion MACs) as easy

What Should Have Happened: With a properly working early exit network, we should see MACs reduced to perhaps 40-70% of the baseline (e.g., 1.3-2.3 trillion MACs) as easy examples exit early while maintaining >90% accuracy.



Analysis: Claude was able to correctly determine the lowest number of MACs we should have seen without the bug, which was cool to see!

v. When would we use early exit, versus just using a smaller model? What factors should we consider?

4b) v) When would we use early exit, versus just using a smaller model? What factors should we consider?

This is an excellent question that gets at the fundamental trade-offs in neural network design. Here are the key factors to consider:

When to Use Early Exit:

1. Variable Example Difficulty:

- When your dataset has examples with widely varying difficulty (some easy, some hard)
- Early exit allows **adaptive computation** - easy examples use fewer resources, hard examples get full network capacity
- A small model treats all examples the same, potentially underperforming on hard cases

2. **Need for High Accuracy on Hard Examples:**

- When you can't compromise on hard example performance
- Early exit maintains access to the full network depth for challenging cases
- A smaller model has a fixed, limited capacity that might fail on hard examples

3. **Dynamic/Unpredictable Workloads:**

- When you don't know the difficulty distribution in advance
- Production systems where incoming data varies in complexity
- Early exit automatically adapts to the data it sees

4. **Inference-Time Resource Constraints:**

- When you need to meet strict latency or energy budgets on average, but can handle occasional hard examples
- Mobile/edge devices with variable power availability
- Real-time systems where average-case performance matters more than worst-case

When to Use a Smaller Model:



When to Use a Smaller Model:

1. Uniform Example Difficulty:

- When most examples have similar difficulty levels
- No benefit from adaptive computation if all examples are equally easy/hard

2. Strict Worst-Case Constraints:

- When you need guaranteed maximum latency/energy per example
- Early exit has variable cost (hard examples still use full network)
- A smaller model has predictable, consistent resource usage

3. Simpler Deployment:

- Smaller models are easier to deploy and maintain
- No need for entropy calculation, exit logic, or multiple classifiers
- Lower implementation complexity and fewer potential bugs

4. Memory Constraints:

- Early exit models require multiple classifier heads, increasing memory footprint
- A smaller model might have fewer t ↓ parameters
- Important for severely resource-constrained devices

- Important for severely resource-constrained devices

5. **Training Simplicity:**

- Standard training procedures work fine
- No need for weighted loss across multiple exits
- Easier hyperparameter tuning

Key Factors to Consider:

Dataset Characteristics:

- Distribution of example difficulty
- Ratio of easy to hard examples
- Whether difficulty is predictable from input features

Performance Requirements:

- Accuracy requirements (especially on hard examples)
- Average vs. worst-case latency/energy budgets
- Tolerance for variable inference time

Computational Resources:

- Available memory (early exit needs more)
- Processing power
- Energy budget (average vs. peak)

Operational Considerations:

- Deployment complexity tolerance
- Maintenance and debugging resources
- Need for interpretability (why did this example exit early?)

Economic Trade-offs:

- Training cost (early exit is more complex to train)
- Inference cost savings (if early exit works well)
- Development time and expertise required



Relationship to Example Difficulty:

Early exit directly leverages our understanding of example difficulty:

- **Easy examples** (low prediction depth) → exit early, save computation
- **Hard examples** (high prediction depth) → use full network, maintain accuracy

This makes early exit ideal when:

- You've analyzed your dataset and found significant variation in example difficulty
- The computational cost of hard examples is justified by their rarity
- Most examples are easy, allowing substantial average-case savings

A smaller model ignores example difficulty and treats everything uniformly, which is wasteful for easy examples and potentially inadequate for hard ones.

Practical Recommendation:

Use Early Exit when:

- Dataset has clear difficulty variation (like ↓ saw: easy examples can be classified in early layers)


- Average-case performance optimization is more important than worst-case
- You need high accuracy on rare hard examples but want to save computation on common easy ones

Use a Smaller Model when:

- Examples are uniformly difficult
- Simplicity and predictability are paramount
- Memory/deployment constraints are severe
- The accuracy loss on hard examples is acceptable

Consider Hybrid Approaches:

- Ensemble of small models of different sizes
- Cascaded classifiers (progressively more complex)
- Dynamic model selection based on input complexity estimation



Analysis: Claude did a good job providing a comprehensive overview for this question. Following this, I tried prompting it to identify other possible sources of error in the code it generated, but it was ultimately not able to figure out exactly where it went wrong. I ran different ablations with all the suggestions it provided, but none of them resolved the MAC issue.