



You are a deep learning teaching assistant helping me map out conceptual relationships in a research paper.

Your task:

- Carefully read the research paper "A Spectral Condition for Feature Learning" and identify all major concepts, methods, and results.
- Build a structured concept map ("brain map") using Mermaid syntax (graph TD), where:
- Nodes represent key topics, methods, or results. Include related works and references inside the nodes.
- Edges connect related nodes, with descriptive labels explaining the relationship (not just keywords).
- Node labels are clear and concise; avoid line breaks for compatibility.
- For each node, list relevant citations and references (from the paper and cited works) below the diagram, grouped by concept.
- Flag nodes that are central (highly connected) or have many references, and highlight these for further study.
- If supported, also generate a PNG image of the diagram for visual reference.

Output Instructions

- Present the concept map first using Mermaid syntax in a Markdown code block. Use graph TD, clear node labels, and descriptive edge labels (e.g., `→ |enables parallel attention|`).
- Represent the 10 most important concepts as nodes in the Mermaid diagram. Include related works and references inside the node.
- List all references/citations grouped by node below the diagram (not inside node labels).
- If possible, generate and return a PNG image of the diagram.
- Avoid using special characters such as parentheses in labels that might not be compatible with Mermaid diagram generators.

Format example:

```
graph TD
    Transformer["Transformer Model (Vaswani et al. 2017)"]
    SelfAttn["Self-Attention (Bahdanau et al. 2014; Lin et al. 2017)"]
    MultiHead["Multi-Head Attention (Vaswani et al. 2017)"]
    PosEnc["Positional Encoding (Gehring et al. 2017)"]
```

```
Adam["Adam Optimizer (Kingma & Ba 2014)"]
```

```
Transformer -->|Directly models all pairwise token dependencies with attention| SelfAttn
Transformer -->|Requires token position information for context awareness| PosEnc
Transformer -->|Is trained efficiently using adaptive gradient optimization| Adam
SelfAttn -->|Provides the base mechanism for building rich token representations| MultiHead
MultiHead -->|Enables parallelized, diverse attention for complex relationships| SelfAttn
```

References by Node:

- Transformer Model: Vaswani et al. 2017, Gehring et al. 2017, Wu et al. 2016
- Self-Attention: Bahdanau et al. 2014, Lin et al. 2017, Parikh et al. 2016
- Multi-Head Attention: Vaswani et al. 2017, Britz et al. 2017
- Positional Encoding: Gehring et al. 2017, Sennrich et al. 2015
- Adam Optimizer: Kingma & Ba 2014

Background:

I am an undergraduate student who is interested in machine learning, deep learning, and research in these fields.

Paper abstract/intro:

The push to train ever larger neural networks has motivated the study of initialization and training at large network width. A key challenge is to scale training so that a network's internal representations evolve nontrivially at all widths, a process known as feature learning. Here, we show that feature learning is achieved by scaling the spectral norm of weight matrices and their updates like p fan-out/fan-in, in contrast to widely used but heuristic scalings based on Frobenius norm and entry size. Our spectral scaling analysis also leads to an elementary derivation of maximal update parametrization. We develop this spectral perspective and structure the text with the intent of providing the reader with a solid conceptual understanding of the scaling behavior of feature learning in neural networks.

Now I'll create a concept map based on the paper. Let me extract and organize the 10 most important concepts with their relationships and references:

```
graph TD
```

```
FeatureLearning["Feature Learning (Yang & Hu 2021b; Lee et al. 2019; Sohl-Dickstein et al. 2016)"]
SpectralScaling["Spectral Scaling Condition (Bernstein et al. 2023)"]
MuP["Maximal Update Parametrization muP (Yang & Hu 2021b; Yang et al. 2021)"]
SpectralNorm["Spectral Norm vs Frobenius Norm (Rudelson & Vershynin 2010; Vershynin 2012)"]
NTP["Neural Tangent Parametrization NTP (Jacot et al. 2018; Lee et al. 2019)"]
WidthScaling["Network Width Scaling (Brown et al. 2020; Arora et al. 2019)"]
GradientAlignment["Gradient Update Alignment (Low-Rank Structure)"]
InfiniteWidth["Infinite Width Limit (Jacot et al. 2018; Yang & Hu 2021b)"]
Optimization["Operator Structure in Optimization (Bernstein et al. 2020a; You et al. 2020)"]
HyperparameterTransfer["Hyperparameter Transfer (Dey et al. 2023; Yang et al. 2021)"]
```

```
FeatureLearning -->|Requires hidden representations to evolve nontrivially during training| SpectralScaling
SpectralScaling -->|Provides elementary derivation and implementation of| MuP
SpectralScaling -->|Ensures feature vectors and updates scale correctly to achieve| FeatureLearning
SpectralNorm -->|Measures matrix size more accurately than Frobenius norm for| SpectralScaling
```

MuP -->|Enables hyperparameter tuning on narrow models then transferring to| WidthScaling
 NTP -->|Loses feature learning at large width unlike| MuP
 GradientAlignment -->|Updates have rank-one structure that aligns with incoming vectors
 InfiniteWidth -->|Feature learning limit requires proper scaling at| WidthScaling
 Optimization -->|Frobenius-normalized updates are suboptimal compared to| SpectralNorm
 HyperparameterTransfer -->|Practical benefit enabled by proper feature learning in| MuP

References by Node:

Feature Learning

- Yang & Hu 2021b (Tensor Programs IV: Feature learning in infinite-width neural networks)
- Lee et al. 2019 (Wide neural networks of any depth evolve as linear models under gradient descent)
- Sohl-Dickstein et al. 2020 (On the infinite width limit of neural networks with a standard parameterization)
- Atanasov et al. 2022 (Neural networks as kernel learners: the silent alignment effect)
- Fort et al. 2020 (Deep learning versus kernel learning)
- Lee et al. 2020 (Finite versus infinite neural networks: an empirical study)
- Vyas et al. 2022 (Limitations of the NTK for understanding generalization in deep learning)

Spectral Scaling Condition

- Bernstein et al. 2023 (Automatic Gradient Descent: Deep Learning without Hyperparameters)
- Current paper's main contribution
- Relates spectral norm to fan-out/fan-in ratio: $\sqrt{(n_\ell/n_{\ell-1})}$

Maximal Update Parametrization (μP)

- Yang & Hu 2021b (Tensor Programs IV)
- Yang et al. 2021 (Tensor Programs V: Tuning large neural networks via zero-shot hyperparameter transfer)
- Dey et al. 2023 (Cerebras-GPT)
- Derived via tensor programs tracking feature distributions analytically

Spectral Norm vs Frobenius Norm

- Rudelson & Vershynin 2010 (Non-asymptotic theory of random matrices: Extreme singular values)
- Vershynin 2018 (High-Dimensional Probability)
- Dhillon & Tropp 2008 (Matrix nearness problems with Bregman divergences)
- Key insight: spectral norm = largest singular value, better for operator structure

Neural Tangent Parametrization (NTP)

- Jacot et al. 2018 (Neural tangent kernel: Convergence and generalization in neural networks)
- Lee et al. 2019 (Wide neural networks of any depth evolve as linear models)
- Arora et al. 2019 (On exact computation with an infinitely wide neural net)
- Loses feature learning at infinite width (lazy training regime)

Network Width Scaling

- Brown et al. 2020 (Language models are few-shot learners - GPT-3)
- Ramesh et al. 2022 (Hierarchical text-conditional image generation with CLIP latents)
- Silver et al. 2016 (Mastering the game of Go with deep neural networks)
- Canatar et al. 2021 (Spectral bias and task-model alignment)
- Challenge: maintaining feature learning as width grows

Gradient Update Alignment

- Current paper's empirical finding (Figure 1)
- Updates have low stable rank even at large batch sizes
- Rank-one structure: $\Delta W_\ell = -\eta_\ell \cdot \nabla_{\{h_\ell\}^L} L \cdot h_{\{\ell-1\}}^T$
- Perfect alignment with incoming vectors in batch size 1 case

Infinite Width Limit

- Jacot et al. 2018 (Neural tangent kernel)
- Lee et al. 2018 (Deep neural networks as Gaussian processes)
- Mei et al. 2019 (Mean-field theory of two-layers neural networks)
- Geiger et al. 2020 (Disentangling feature and lazy training)
- Yaida 2022 (Meta-principled family of hyperparameter scaling strategies)
- Different scaling regimes: NTK limit vs feature learning limit

Operator Structure in Optimization

- Bernstein et al. 2020a (On the distance between two neural networks and the stability of learning)
- Bernstein et al. 2020b (Learning compositional functions via multiplicative weight updates)
- You et al. 2017, 2020 (Large batch optimization - Frobenius-normalized updates)
- Liu et al. 2021 (Learning by turning: Neural architecture aware optimisation)
- Shazeer & Stern 2018 (Adafactor: Adaptive learning rates with sublinear memory cost)
- Amari 1998 (Natural gradient works efficiently in learning)
- Nemirovsky & Yudin 1983 (Problem complexity and method efficiency in optimization)

Hyperparameter Transfer

- Yang et al. 2021 (Zero-shot hyperparameter transfer)

- Dey et al. 2023 (Cerebras-GPT: Open compute-optimal language models)
- Key practical benefit: tune on narrow networks, transfer to wide networks
- Avoids expensive tuning on large models

Central/Highly-Connected Nodes for Further Study:

1. **Spectral Scaling Condition** - Most central node; directly connects to feature learning, μP , spectral norm theory, and gradient alignment
2. **Feature Learning** - Core desideratum; connects to width scaling, μP , NTP, and hyperparameter transfer
3. **Maximal Update Parametrization (μP)** - Key practical outcome; connects to spectral scaling, feature learning, width scaling, and hyperparameter transfer
4. **Spectral Norm** - Fundamental mathematical tool; connects to spectral scaling, optimization theory, and distinguishes from Frobenius norm approaches

Additional Key Papers to Study:

- Yang & Hu 2021b: Original μP derivation via tensor programs
- Jacot et al. 2018: Neural tangent kernel framework
- Bernstein et al. 2020a: Earlier operator structure analysis (with flawed conditioning assumption corrected by current paper)
- Vershynin 2018: Random matrix theory foundations

**