

Participation Section E:

Brief Intro

In this part, I will use Grok4 (fast) as a teacher to teach me about SSM (State Space Model). I will try to let LLM to fully help me to understand this knowledge. I will attach my prompt and LLM's response (they are exported as a separated file), and my personal comment on LLM's response.

Comment

My first prompt is:

```
Hi there! I'm a student learning deep learning, and I want to learn about State Space Model. Could you first give a brief formulation of ssm model and introduction of its design?
```

As kick off, I want Grok to first give a comprehensive introduction to SSM. From Grok's response, we can tell that its answer covered wide range of topic and knowledge of SSM. From discrete and continuous scenario, its motivation and advantage, etc.

This is normal if you ask Grok a general question, I want Grok to cut in this topic in step by step. My second prompt want to hence start with the discrete case:

```
Let's focus on discrete scenario first and leave out these complicated deduction. Please first elaborate on the motivation for SSM and what's the benefit of SSM compared to RNN?
```

I asked this question because I think it is important to grasp the motivation and advantages for every model, no matter it is SSM or other novel models. So I asked that question to ask Grok to elaborate on that. Grok covered Computation Mode, Handling Long Dependencies, Memory Efficiency, Stability & Training, and Expressivity. I'm quite satisfied with its answer as it covered most of the points covered in the lecture.

In the lecture, I didn't fully understand why FFT can degrade the computation complexity to $O(L\log L)$. So then my question is:

```
Thank you, can you elaborate on why FFT can reduce computation complexity?
```

Grok claimed that SSM can be regarded as sort of convolution, and convolution with sequence is equivalent to multiplication in frequency domain. And give a summary on the usage of FFT. Now I had a basic understanding about why FFT can help reduce the computation complexity.

Then my question is:

OK. Please tell me what's the connection between discrete scenario and continuous scenarion, and what's the point to delve into continuous scenario?

In short, the continuous scenario is the "why it works" layer—discrete is the "how to implement" layer. In the response, Grok skimmed the effect of ΔT , so added one prompt about this point:

Can you give intuitive example to demonstrate the effect of ΔT ?

Grok perfectly explained its effect, and match what's been concluded in the lecture. Lastly, one thing came to my mind that for the sake of easy computation of A^k , we may need some special design of A matrix structure:

In continuos and discrete scenario, we say that in order to simplify the computation of A^k , maybe we want A to be in some special form?

Finally, I asked Grok to give a summary about SSM:

I get that. Is there any special method for initialization, since that SSM model is quite different than other traditional model?

In summary, Grok indeed is a good tutor and can answer my questions in detail. In the following part, the chat with Grok is attached.

Hi there! I'm a student learning deep learning, and I want to learn about State Space Model. Could you first give a brief formulation of ssm model and introduction of its design?

✓ Thought for 3s

Brief Introduction to State Space Models (SSMs)

State Space Models (SSMs) are a class of sequence modeling architectures in deep learning, inspired by classical control theory and linear dynamical systems. They provide an efficient alternative to Transformers for handling long-range dependencies in sequences (e.g., time series, audio, or text), by maintaining a compact "state" that evolves over time. SSMs excel in computational efficiency for very long inputs, as they can be computed via fast parallel scans or convolutions, avoiding the quadratic scaling of attention mechanisms. Popular implementations include S4 (Structured State Space sequence models) and Mamba, which extend SSMs with nonlinearities and selective mechanisms for better expressivity.

Mathematical Formulation

An SSM models a sequence input $x = (x_1, \dots, x_L) \in \mathbb{R}^{L \times N}$ (where L is sequence length and N is input dimension) to produce outputs $y = (y_1, \dots, y_L) \in \mathbb{R}^{L \times M}$ (output dimension M). It maintains a hidden state $h_t \in \mathbb{R}^D$ (state dimension D , often small like 64).

Continuous-Time Form (Foundation)

The core is a linear ODE:

$$\dot{h}(t) = Ah(t) + Bx(t), \quad y(t) = Ch(t)$$

- $A \in \mathbb{R}^{D \times D}$: Transition matrix (controls state dynamics).
- $B \in \mathbb{R}^{D \times N}$: Input matrix (how input affects state).
- $C \in \mathbb{R}^{M \times D}$: Output matrix (how state produces output).

This describes how the state $h(t)$ evolves continuously under input $x(t)$.

Discrete-Time Form (For Sequences)

For discrete steps (e.g., time $t = 1, \dots, L$), we discretize using a step size Δ (learnable or fixed):

$$h_t = \overline{A}h_{t-1} + \overline{B}x_t, \quad y_t = Ch_t$$

Where:

$$\overline{A} = \exp(A\Delta), \quad \overline{B} = (A\Delta)^{-1}(\exp(A\Delta) - I) \cdot (B\Delta)$$

(Using zero-order hold discretization; other methods like bilinear exist.)

The full sequence can be unrolled recurrently (for training) or computed in parallel via convolution:

$$y = (C \circ \overline{B}) * (x \cdot K), \quad K_t = \overline{A}^{L-t}$$

Where $*$ is convolution and $\circ$ is matrix multiplication—enabling $O(L \log L)$ or linear time complexity.

Key Design Aspects

- **State Representation:** The fixed-size state h_t compresses the entire history, enabling linear-time inference (unlike RNNs' sequential bottleneck or Transformers' memory explosion).
- **Initialization & Structuring:** To capture long dependencies, designs like HiPPO (High-order Polynomial Projection Operators) initialize A diagonally with structured eigenvalues, ensuring the model approximates Legendre polynomials for orthogonal basis functions over long horizons.
- **Efficiency Tricks:**
 - **Parallel Scan:** Computes the recurrent form in log-linear time using associative scans.
 - **Convolution Mode:** Equivalent to a global convolution kernel, fast via FFT.
- **Extensions:** Modern SSMs (e.g., Mamba) add **selectivity** (input-dependent A, B, C) and gating/nonlinearities for Transformer-level performance on discrete data like language.

- **Training:** Parameters A, B, C, Δ are learned end-to-end; often stacked in layers like RNNs.

For deeper dives, check papers like "Efficiently Modeling Long Sequences with Structured State Spaces" (S4) or "Mamba: Linear-Time Sequence Modeling with Selective State Spaces." If you'd like code examples (e.g., in PyTorch) or comparisons to LSTMs/Transformers, let me know!

Let's focus on discrete scenario first and leave out these complicated deduction. Please first elaborate on the motivation for SSM and what's the benefit of SSM compared to RNN?

💡 Thought for 3s

Motivation for State Space Models (SSMs)

State Space Models (SSMs) emerged as a response to the limitations of existing sequence modeling architectures in deep learning, particularly for tasks involving long sequences like time series forecasting, audio processing, or natural language. The core idea draws from classical control theory, where systems are modeled using a hidden "state" that captures the system's memory and evolves over time based on inputs. In machine learning, this translates to a framework where the model maintains a compact, fixed-size hidden state h_t that summarizes the entire input history up to time t , allowing the model to "remember" relevant information without storing every past input explicitly.

The motivation stems from two key pain points in sequence modeling:

- **Scalability for Long Sequences:** Real-world data often has thousands or millions of steps (e.g., a full audio waveform or a long document). Early models like RNNs struggle with training stability over long horizons due to issues like vanishing gradients, while Transformers, though powerful, scale quadratically ($O(L^2)$ time and memory for sequence length L), becoming impractical beyond a few thousand tokens.
- **Efficiency and Parallelization:** We need models that can train and infer quickly on modern hardware (GPUs/TPUs), which excel at parallel computations, but still capture temporal dependencies like RNNs do.

SSMs address this by formulating sequence processing as a linear dynamical system in discrete time, which can be computed either recurrently (like an RNN) or in a fully parallel way (like a convolution). This makes them a "best of both worlds" approach: expressive enough for complex patterns but computationally lightweight, enabling applications on sequences up to 1M+ length without exploding resources.

Benefits of SSMs Compared to RNNs

RNNs (e.g., vanilla RNNs, LSTMs, GRUs) were the go-to for sequences before Transformers, but they have inherent flaws that SSMs mitigate. Here's a comparison:

Aspect	RNNs	SSMs
Computation Mode	Strictly sequential: Each step $h_t = f(h_{t-1}, x_t)$ must wait for the previous, leading to $O(L)$ time but poor GPU utilization (no parallelism during training).	Dual-mode: Recurrent for inference ($O(L)$, sequential) or convolutional for training ($O(L \log L)$ via FFT, fully parallel). This speeds up training by 5-10x on long sequences.
Handling Long Dependencies	Prone to vanishing/exploding gradients; LSTMs/GRUs add gates but still forget info over 100s of steps.	Structured matrices (e.g., diagonal A) ensure stable evolution; HiPPO initialization projects inputs onto orthogonal polynomial bases, capturing dependencies over 10k+ steps without gradient issues.
Memory Efficiency	Hidden state is fixed-size, but unrolling for backprop requires $O(L)$ memory in practice (or approximations like truncated BPTT).	Fixed-size state h_t (e.g., dim 64) compresses history globally; convolution kernel is precomputed once per layer, using $O(1)$ memory beyond inputs. Handles 1M-length sequences on consumer GPUs.
Stability & Training	Nonlinear activations and recurrent loops make optimization tricky; often needs careful clipping or gating.	Linear core (with optional nonlinear extensions) is easier to analyze and stabilize; parameters like A, B, C can be initialized for provable long-memory properties.

Aspect	RNNs	SSMs
Expressivity	Good for short-term patterns but limited for global structure without stacking many layers.	Matches or exceeds RNNs on continuous data (e.g., audio); extensions like Mamba add input-dependent selectivity for discrete tasks (e.g., language), rivaling Transformers.

In essence, SSMs retain the RNN's elegant state-based compression but make it faster, more stable, and scalable—turning a sequential bottleneck into a parallel powerhouse. For example, in the S4 model, the discrete recurrence $h_t = \overline{A}h_{t-1} + \overline{B}x_t, y_t = Ch_t$ looks RNN-like, but the magic is in computing the entire y via a single convolution kernel K , where K encodes all past influences.

If you'd like examples with toy code (e.g., a simple discrete SSM in PyTorch) or a dive into how this compares to Transformers next, just say the word!

Thank you, can you elaborate on why FFT can reduce computation complexity?

💡 Thought for 3s

Why FFT Reduces Convolution Complexity in SSMs

In the context of State Space Models (SSMs), the discrete-time formulation can be unrolled into a **global convolution** operation. Specifically, the output sequence y_t at each step is a weighted sum of all past inputs $x_1, \dots, x_t$, where the weights are defined by the kernel K derived from the state transition matrix $\overline{A}$: $K_s = \overline{A}^s$ for lag s . Thus, $y = K * x$, where $*$ denotes convolution.

Computing this convolution naively would be quadratic in sequence length L , but the **Fast Fourier Transform (FFT)** drops it to near-linear $O(L \log L)$, making SSMs scalable for long sequences (e.g., 1M+ steps). Let's break it down step-by-step.

1. Naive Convolution: Why It's $O(L^2)$

- Convolution slides the kernel K (length up to L) over the input x (length L) to produce y (length L).

- For each output position $t = 1$ to L , you compute:

$$y_t = \sum_{s=0}^{t-1} K_s \cdot x_{t-s}$$

(Causal convolution, padded appropriately for past-only dependencies.)

- Each sum has up to L terms, and there are L outputs, so total operations: $O(L^2)$.
- This is fine for short sequences but explodes for long ones—e.g., $L = 10^6$ requires $\sim 10^{12}$ operations, infeasible on GPUs.

2. The FFT Trick: Convolution in the Frequency Domain

- **Key Insight:** Convolution in the *time domain* (sequences like x, K) is equivalent to *pointwise multiplication* in the *frequency domain* (their Fourier transforms).
- The Discrete Fourier Transform (DFT) converts a sequence to frequencies:

$$\hat{x}_k = \sum_{n=0}^{L-1} x_n e^{-2\pi i k n / L}, \quad k = 0, \dots, L-1$$

(Inverse DFT recovers the original.)

- For *circular* convolution (wrapping around the sequence), the theorem states:

$$\mathcal{F}(K * x) = \hat{K} \odot \hat{x}$$

where $\mathcal{F}$ is the DFT, $\hat{K}, \hat{x}$ are transforms, and $\odot$ is element-wise multiplication.

- To handle *linear* (non-circular) convolution like in SSMs, we pad sequences to length $2L - 1$ with zeros, then crop the result—still efficient.

3. FFT Algorithm: From $O(L^2)$ to $O(L \log L)$

- The plain DFT is $O(L^2)$ (nested sums), same as naive convolution—no gain.
- **Fast Fourier Transform (FFT)** is a divide-and-conquer algorithm (Cooley-Tukey) that computes the DFT efficiently:
 - Recursively splits the sequence into even/odd indices, reducing to two half-size DFTs.

- Base case: DFT of size 1 is trivial.
- Recurrence: Time for size $L = 2^m$ is $T(L) = 2T(L/2) + O(L)$, solving to $T(L) = O(L \log L)$.
- Full convolution pipeline:
 1. FFT of x : $O(L \log L)$
 2. FFT of K : $O(L \log L)$
 3. Pointwise multiply $\hat{K} \odot \hat{x}$: $O(L)$
 4. Inverse FFT to get y : $O(L \log L)$
- **Total: $O(L \log L)$** —a massive speedup! For $L = 10^6$, it's ~20 million operations vs. a trillion.

4. Practical Benefits in SSMs

- **Parallelism:** FFT is highly parallelizable on GPUs (e.g., via cuFFT in PyTorch), unlike sequential RNN scans.
- **Kernel Precomputation:** In SSM layers, compute K once per layer (via fast exponentiation of $\bar{A}$, also $O(L \log L)$), then convolve with any x .
- **Trade-offs:** Assumes real/complex numbers; for very short L , naive might be faster due to FFT overhead. But for SSMs' target (long audio/text), it's ideal.
- **Extensions:** Real FFTs (for real-valued signals) halve the cost; libraries like PyTorch's `torch.fft` handle padding/cropping automatically.

Quick Toy Example (Conceptual)

Imagine $L = 4$, $x = [1, 2, 3, 4]$, simple $K = [1, 0.5]$ (exponentially decaying weights).

- Naive: 4 sums, each up to 2 multiplies $\rightarrow$ ~10 ops.
- FFT: Pad to 8, FFT each (64 ops total, but scales better), multiply (8 ops), iFFT $\rightarrow$ still low, but for $L=1024$, naive=1M ops vs. FFT~10k.

If you'd like a PyTorch code snippet to demo this (with timings), or how it ties back to SSM discretization, let me know!

OK. Please tell me what's the connection between discrete scenario and continuous scenario, and what's the point to delve into continuous scenario?

Connection Between Discrete and Continuous Scenarios in SSMs

In State Space Models (SSMs), the **continuous scenario** provides the foundational mathematical framework, while the **discrete scenario** is a practical approximation tailored for sequence data (like discrete time steps in audio samples or text tokens). The two are tightly linked: the discrete form is derived directly from the continuous one via **discretization**, which "samples" the continuous dynamics at fixed intervals. This makes the discrete model a bridge between elegant continuous theory and efficient discrete computation.

Core Mathematical Link

- **Continuous Form** (time as a real variable $t \in \mathbb{R}$):

$$\dot{h}(t) = Ah(t) + Bx(t), \quad y(t) = Ch(t)$$

Here, $\dot{h}(t)$ is the instantaneous rate of change of the state $h(t)$, governed by an ordinary differential equation (ODE). This models smooth, ongoing evolution, ideal for truly continuous signals.

- **Discretization to Discrete Form** (time as integers $t = 1, 2, \dots, L$, with step size $\Delta > 0$):

To apply this to discrete sequences, we approximate the ODE by sampling at intervals Δ (e.g., $\Delta = 1$ for unit steps). A common method (zero-order hold) assumes $x(t)$ is constant between steps, yielding:

$$h_t = \bar{A}h_{t-1} + \bar{B}x_t, \quad y_t = Ch_t$$

Where the discrete parameters are transformed from the continuous ones:

$$\bar{A} = \exp(A\Delta), \quad \bar{B} = (A\Delta)^{-1}(\exp(A\Delta) - I)B\Delta$$

- $\exp(\cdot)$ is the matrix exponential, solving the ODE over one step.
- This preserves the "physics" of the continuous system: if you solve the continuous ODE exactly between samples, you get exactly the discrete recurrence.

Other discretization schemes exist (e.g., Euler: $h_t \approx h_{t-1} + \Delta(Ah_{t-1} + Bx_t)$, simpler but less accurate for large Δ), but they all aim to mimic continuous behavior on a grid.

In practice:

- Train/learn the **continuous parameters** A, B, C, Δ .
- Convert to discrete $\overline{A}, \overline{B}$ once (cheap, $O(D^3)$ where D is state dim), then use the discrete form for forward passes (recurrent or convolutional).

This connection ensures the discrete model inherits desirable traits from the continuous one, like stability and long-memory properties, without losing much accuracy for small Δ .

Why Delve into the Continuous Scenario?

While the discrete form is sufficient for most implementations (and what we focused on earlier), starting from continuous isn't just theoretical—it's a deliberate design choice with practical payoffs. Here's why it's worth exploring:

1. Structured Initialization for Better Performance:

- Continuous math allows principled ways to initialize A (e.g., HiPPO: sets A to mimic projection onto Legendre polynomials). This ensures the model captures *any* polynomial of degree up to D over long horizons, providing a strong inductive bias for smooth, long-range dependencies.
- In discrete-only designs, it's harder to enforce such structure—discretization "bakes it in," so you get these benefits "for free" without hand-tuning.

2. Handles Variable Time Steps and Continuous Data:

- Real data like audio (continuous waveforms) or physics sims isn't perfectly discrete. Continuous formulation naturally interpolates between irregular samples (e.g., vary Δ per step), avoiding artifacts from fixed grids.
- Discrete models can alias or lose resolution for high-frequency signals; continuous avoids this by modeling the underlying dynamics.

3. Theoretical Insights and Analysis:

- ODEs have rich tools (e.g., stability via eigenvalues of A , frequency response via Laplace transforms) for proving properties like gradient flow or

forgetfulness. This guides why SSMs beat RNNs on long sequences.

- Enables hybrids: Solve the ODE numerically (e.g., via ODE solvers like Runge-Kutta) during training for more flexibility, then discretize for fast inference.

4. Scalability and Generalization:

- Continuous params are low-dimensional (e.g., diagonal A for efficiency), making training stable even for large L .
- Inspires extensions like S4D (discretized S4) or Mamba, where continuous roots enable selectivity (input-dependent dynamics) without blowing up compute.

In short, the continuous scenario is the "why it works" layer—discrete is the "how to implement" layer. For a student, grasping continuous unlocks why SSMs are more than just fancy RNNs; they're grounded in dynamical systems theory. If you want a simple numerical example (e.g., simulating both forms in code to see they match), or how HiPPO fits in, let me know!

Can you give intuitive example to demonstrate the effect of ΔT ?

💡 Thought for 24s

Intuitive Example: Exponential Decay in an SSM

To make the effect of Δ (the time step size) intuitive, let's consider a simple 1D **exponential decay** system, inspired by something everyday like a **leaking bucket** or **radioactive decay**. Imagine the hidden state h_t represents the water level (or amount of material) in the bucket. Without new input ($x_t = 0$), it leaks out over time according to the continuous dynamics:

$$\dot{h}(t) = Ah(t), \quad A = -0.5$$

This means the level drops at a rate proportional to its current height—halving roughly every $\ln(2)/0.5 \approx 1.4$ units of continuous time. Starting from $h_0 = 1$, the exact continuous solution is $h(t) = e^{At} = e^{-0.5t}$.

In the **discrete SSM**, we sample this at steps of size Δ , so the update becomes $h_t = \bar{A}h_{t-1}$, where $\bar{A} = e^{A\Delta} = e^{-0.5\Delta}$. Each discrete step (e.g., each "token" or data point)

advances the **physical time** by Δ , but we fix the number of steps to $L = 10$ (like a fixed-length sequence).

- **Key Intuition:** Δ controls the "clock speed" of the system.
 - Small Δ : Each step ticks forward a tiny bit in time $\rightarrow$ slow, gradual leak $\rightarrow$ the state **remembers** the initial level longer (less total decay over 10 steps).
 - Large Δ : Each step ticks forward a lot in time $\rightarrow$ fast, aggressive leak $\rightarrow$ the state **forgets** the initial level quickly (more total decay).

The total "physical time" elapsed after 10 steps is 10Δ , so the final state is always $h_{10} = e^{-0.5 \times 10\Delta}$ (exact match to continuous, thanks to proper discretization). But for a fixed sequence length, larger Δ squeezes more decay into fewer steps, making the model evolve faster overall.

Numerical Trajectories for Different Δ

Here's the state h_t at each of the 10 discrete steps (rounded to 3 decimals for readability). Notice how the decay curve steepens with larger Δ :

Step t	$\Delta = 0.1$ (Slow Clock: Total Time = 1.0)	$\Delta = 0.5$ (Medium Clock: Total Time = 5.0)	$\Delta = 1.0$ (Fast Clock: Total Time = 10.0)
1	0.951	0.779	0.607
2	0.905	0.607	0.368
3	0.861	0.472	0.223
4	0.819	0.368	0.135
5	0.779	0.287	0.082
6	0.741	0.223	0.050
7	0.705	0.174	0.030
8	0.670	0.135	0.018
9	0.638	0.105	0.011
10	0.607	0.082	0.007

- **Small $\Delta = 0.1$:** The level drops gently—after 10 steps, ~60% remains. It's like checking the bucket every few seconds; the leak seems slow, preserving memory of the full initial amount.
- **Medium $\Delta = 0.5$:** Faster drop—down to ~8% left. Like checking every minute; you see more noticeable emptying per check.
- **Large $\Delta = 1.0$:** Rapid plunge—barely 0.7% left. Like checking hourly; by the 10th check, it's essentially empty, forgetting the start almost immediately.

Why This Matters in SSMs

In real models (e.g., S4 or Mamba), Δ is often **learnable** per layer or even per token.

This lets the model adapt its "internal clock" to the data:

- For long documents (needing long memory), it learns small Δ to stretch the dynamics.
 - For quick events (like speech phonemes), larger Δ speeds things up.
- If Δ is fixed too large, the model might "over-forget" early context; too small, it might under-react to recent inputs.

This example keeps it linear and input-free for simplicity, but the idea extends: with inputs, large Δ also "averages" inputs over bigger time chunks, smoothing sharp changes. If you'd like a plot visualization (via code), a version with inputs, or how this ties to HiPPO initialization, just ask!

In continuous and discrete scenario, we say that in order to simplify the computation of A^k , maybe we want A to be in some special form?

💡 Thought for 4s

Simplifying A^k Computation via Special Forms in SSMs

Yes, exactly! In both continuous and discrete SSM scenarios, computing powers like A^k (or $\bar{A}^k$ in discrete) is a computational bottleneck when building the convolution kernel for efficient parallel training. The kernel $K_k = C \bar{A}^k \bar{B}$ (for lag k) needs to be evaluated for all $k = 0$ to $L - 1$, and with sequence lengths L up to millions, naive matrix exponentiation per k ($O(D^3)$ per power, times L) would be $O(L D^3)$ —prohibitively slow even for modest state dimensions $D=64$.

The key insight is to impose a **structured form** on the transition matrix A (or $\overline{A}$) that allows fast, closed-form or $O(D \log D + L)$ computation of all powers at once. This structure is typically designed in the **continuous domain** (for theoretical guarantees) and preserved through discretization. Let's break it down.

Why A^k Matters

- **Discrete Case:** Unrolling the recurrence gives $h_t = \sum_{k=0}^{t-1} \overline{A}^k \overline{B} x_{t-k}$, so the full output is a convolution with kernel $K_k = C \overline{A}^k \overline{B}$. To compute K efficiently:
 - Naive: Iterative multiplication (still $O(L D^2)$ if done smartly, but quadratic in D).
 - With FFT: We need K as a vector first, so fast powers are essential.
- **Continuous Case:** The impulse response is $K(t) = C e^{At} B$, and discretization yields $\overline{A}^k = e^{Ak\Delta}$, so structuring A simplifies the matrix exponential e^{At} too (e.g., via eigendecomposition).

Without structure, e^{At} or A^k requires full SVD or exponentiation, which is $O(D^3)$ per t/k —fine for small D, but SSMs shine with $D \sim 16\text{--}128$.

Common Special Forms for A and Their Benefits

We design A to be **diagonalizable with simple eigenvalues** or **low-rank**, often in continuous space (e.g., A diagonal with negative real parts for stability). Discretization $\overline{A} = \exp(A\Delta)$ inherits the structure (e.g., diagonal stays diagonal). Here's a comparison of forms:

Form of A	How It Simplifies A^k or e^{At}	Pros in SSMs	Cons/When Used	Example in Literature
Diagonal (A = $\text{diag}(\lambda_1, \dots, \lambda_D)$)	$A^k_{ii} = \lambda_i^k$ (elementwise scalar power); full vector K in $O(L D)$.	Extremely fast; no matrix ops. Enables per-mode (eigenvalue) independence.	Assumes uncoupled modes; simple dynamics.	HiPPO (S4): A initialized as $\text{diag}(-k)$ for polynomial approximation.

Form of A	How It Simplifies A^k or e^{At}	Pros in SSMs	Cons/When Used	Example in Literature
Triangular (Lower/Upper, e.g., $A = \Lambda + N$ where Λ diagonal, N nilpotent)	Closed-form via binomial expansion: $A^k = \sum_{m=0}^{D-1} \binom{k}{m} \Lambda^{k-m} N^m$, but truncates for large k .	Handles mild coupling; $O(D^2 L)$ but parallelizable.	More compute than diagonal; for hierarchical modes.	Early SSMs; Jordan form approximations.
Low-Rank ($A = D + UV^T$, D diagonal, U, V low-dim)	$A^k \approx (D^k) + \sum$ low-rank corrections (via Sherman-Morrison or series).	Compact storage ($O(D)$ params); good for selective SSMs.	Approximation error; needs tuning.	Mamba: Input-dependent low-rank updates.
Toeplitz/Structured (Circulant or Hankel)	FFT-diagonalizable: A^k via FFT of eigenvalues ($O(D \log(D + L))$).	Ties directly to convolution speedup; global structure.	Less flexible for arbitrary dynamics.	S4: HiPPO makes A such that e^{At} is orthogonal polynomial basis.

- **Intuition:** Think of A as defining "modes" of oscillation/decay (eigenvalues λ). Diagonal form decouples them—each dimension of h evolves independently as $h_j(t) = h_j(0) \lambda_j^t + \text{inputs}$. Powers become scalar multiplies, like independent RC circuits in parallel.
- **Computation Flow:**
 1. Initialize continuous A in special form (e.g., diagonal via HiPPO: $\lambda_j = -j$ for $j=1$ to D , mimicking decay rates).
 2. Discretize: $\bar{A}_{jj} = \exp(\lambda_j \Delta)$ —still diagonal.

3. Build kernel: For each mode j , $K^j_k = c_j \overline{A_{jj}}^k b_j$ (scalar!); sum over j for full K .
4. Convolve: $\text{FFT}(K) \odot \text{FFT}(x) \rightarrow O(L \log L)$.

Why Continuous for Design?

- **Preservation:** Structure like diagonality survives $\exp(\cdot)$, so we impose it on continuous A for clean math (e.g., eigenvalues control continuous frequencies).
- **Theoretical Perks:** In continuous, we can prove properties like the HiPPO matrix approximating $\int x(s) P_k(t-s) ds$, where P_k are orthogonal polynomials—ensuring uniform memory over $[0, T]$ without tuning.
- **Flexibility:** Learn continuous $A/B/C/\Delta$, discretize on-the-fly; avoids discrete-only instabilities.

In practice (e.g., S4 code), this makes a full layer $\sim O(L D)$ time, vs. $O(L D^2)$ for unstructured. For a toy demo, consider $D=1$: $A=-1$, $\Delta=0.1 \rightarrow \overline{A} \approx 0.905$, $K_k = 0.905^k$ —geometric series, trivial. Scale to $D=64$: Still scalar ops per mode.

If you'd like pseudocode for diagonal case, a deeper HiPPO example, or how this fails without structure, let me know!

I get that. Is there any special method for initialization, since that SSM model is quite different than other traditional model?

💡 Thought for 3s

Special Initialization Methods in SSMs

Yes, absolutely—initialization in State Space Models (SSMs) is a critical design choice, more so than in many traditional models like RNNs or Transformers, because SSMs rely on the **linear dynamical system's stability and expressivity** to handle long sequences without exploding/vanishing signals. Random initialization (e.g., Xavier/Glorot) can lead to poor conditioning of the transition matrix A , causing the model to either forget everything immediately (unstable eigenvalues) or oscillate wildly. Instead, SSMs use **structured, theory-driven initializations** that embed inductive biases for long-memory and smoothness, often derived from the continuous formulation.

The standout method is **HiPPO (High-order Polynomial Projection Operators)**, introduced in the S4 paper. It's not just a tweak—it's a foundational technique that

makes SSMs "special." Let's break it down.

1. HiPPO Initialization: The Core Special Method

- **What It Does:** HiPPO initializes the continuous transition matrix A (and related params) so that the SSM's impulse response $K(t) = Ce^{At}B$ approximates the projection of the input onto an **orthogonal polynomial basis** (e.g., shifted Legendre polynomials on $[0, T]$, where T is the sequence horizon).
 - In simple terms: It ensures the model can exactly represent *any polynomial* of degree up to the state dimension D over the sequence length, like a universal approximator for smooth functions.
 - Mathematically, for a diagonalizable A , HiPPO sets the eigenvalues λ_j (diagonal entries) to specific values that make e^{At} mimic polynomial decay rates. A common variant (HiPPO-LegS) uses:

$$A_{jj} = -(2j + 1), \quad j = 0, \dots, D - 1$$

(Negative reals for decay; other variants like HiPPO-T for time-delay systems use complex pairs for oscillations.)

- Then, B and C are chosen to align with the polynomial coefficients (e.g., B orthogonal to match the basis).
- **Why It's Special for SSMs:**
 - **Long-Memory Bias:** Traditional RNNs forget exponentially; HiPPO gives *uniform sensitivity* to all past timesteps, decaying polynomially (slower than exponential), ideal for capturing global context in audio or text.
 - **Stability Guarantee:** Eigenvalues are structured (e.g., real, negative) to ensure $|\overline{A}| < 1$ post-discretization, preventing explosions.
 - **Efficiency:** Keeps A diagonal or low-rank, enabling the fast A^k computations we discussed.
- **How It Differs from Traditional Models:**
 - RNNs/LSTMs: Use random orthogonal or Xavier init for gates/hidden states to preserve variance, but no built-in long-horizon bias—relies on nonlinearities to learn memory.

- Transformers: He/Kaiming init for weights, with layer norms for stability; focuses on attention scaling, not temporal dynamics.
- SSMs: HiPPO is *physics-inspired* (from control theory's observability), embedding "what to remember" directly into params, reducing the burden on training.
- **Implementation Snippet (Conceptual, in PyTorch-like pseudocode):**

python

```
def hippo_init(D, N, M, variant='legs'): # D=state dim, N=input dim,
    A = torch.diag(torch.arange(D).float() * -2 - 1) # HiPPO-LegS dia
    B = torch.eye(D, N) if N <= D else torch.eye(D).repeat(1, N // D)
    C = torch.eye(M, D) if M <= D else torch.eye(M).repeat(1, D // M)
    Delta = torch.ones(1) * 0.001 # Small initial step size, learnabl
    return A, B, C, Delta
```

(Full S4 libs implement variants like HiPPO-T for better frequency handling.)

2. Other Complementary Initializations in SSMs

While HiPPO is for A (the "backbone"), the full setup includes:

- **B and C Matrices:** Orthogonal initialization (e.g., via QR decomposition) to preserve signal norms through the linear layers, similar to RNNs but stricter for the compression/expansion steps.
 - Why? B projects high-dim inputs ($N \gg D$) to low-dim state; C expands back—random init could squash signals.
- **Δ (Step Size):** Initialized small (e.g., 0.001) to start with slow dynamics (long memory), then learned. This adapts the "resolution" of the internal clock.
- **Nonlinear Extensions (e.g., in Mamba):** Gating params (like GLU) use standard Xavier, but the linear core stays HiPPO-structured.

Quick Comparison Table: SSM Init vs. Traditional

Model Type	Key Init Method	Goal	Long-Sequence Handling
Vanilla RNN	Xavier/Orthogonal on weights	Variance preservation	Poor; exponential forget via gradients
LSTM/GRU	Orthogonal for gates, Xavier elsewhere	Gate stability	Better with cells, but still O(L) issues
Transformer	He/Kaiming on FFNs, small for QKV	Norm flow through residuals	Quadratic; positional encodings for order
SSM (S4/Mamba)	HiPPO for A + Orthogonal B/C	Polynomial memory + stability	Excellent; O(L) with structured decay

Benefits and Tips for Learning/Using

- **Empirical Wins:** HiPPO makes SSMs 10-100x faster on long audio tasks (e.g., Speech Commands dataset) by avoiding gradient clipping hacks needed in RNNs.
- **When to Use Variants:** HiPPO-LegS for general sequences; HiPPO-T for oscillatory data like audio.
- **As a Student:** Experiment by ablating HiPPO in a toy S4 impl—train on sine waves of varying frequencies; you'll see random A fails on long periods.

If you'd like a code example to toy with HiPPO (e.g., via the [s4](#) PyTorch lib), details on Mamba's selective init, or how to derive the polynomial math, let me know!

Can you give a final detailed and full-around summary about the knowledge of SSM?

🕒 Thought for 10s

Comprehensive Summary of State Space Models (SSMs) in Deep Learning

State Space Models (SSMs) represent a transformative class of architectures in deep learning for sequence modeling, bridging classical control theory with modern neural networks. Unlike dominant paradigms like Transformers, SSMs maintain a compact, evolving hidden state to capture temporal dependencies efficiently, making them particularly suited for long sequences in domains like time series, audio, vision, and language. Originating from linear dynamical systems, SSMs have evolved rapidly since

the seminal S4 (Structured State Space sequence models) in 2021, with breakthroughs like Mamba (2023) introducing selectivity and hardware awareness. By 2025, SSMs are no longer niche; they're integral to hybrid systems, rivaling Transformers in performance while slashing compute costs, and extending to robotics, medical imaging, and beyond. This summary covers their foundations, mechanics, advantages, implementations, and cutting-edge developments. [arxiv.org](#)

Motivation and Core Benefits

SSMs address key bottlenecks in sequence modeling: **scalability for ultra-long inputs** (e.g., 1M+ tokens) and **parallelizable computation** on GPUs. Traditional recurrent models like RNNs/LSTMs suffer from vanishing gradients and sequential bottlenecks, while Transformers' self-attention scales quadratically ($O(L^2)$ time/memory for length L), limiting them to ~4k-8k tokens without approximations.

- **Motivation:** Real-world sequences (e.g., genomic data, high-res audio) demand models that compress history into a fixed-size state without losing global context. SSMs draw from control theory's state-space representation, where a latent state h_t evolves linearly under inputs, enabling "global convolution" for $O(L \log L)$ or linear-time processing.
- **Benefits:**
 - **Efficiency:** Linear scaling in L ; trains 5-10x faster than Transformers on long data.
 - **Long-Range Memory:** Structured designs ensure polynomial (not exponential) decay, capturing dependencies over 10k+ steps.
 - **Simplicity:** Fewer hyperparameters than gated RNNs; stable gradients via linearity.
 - **Versatility:** Handles continuous (e.g., physics sims) and discrete data; extensible with nonlinearities.

Compared to RNNs (sequential, unstable) and Transformers (memory-hungry, parallel but quadratic), SSMs offer a "recurrent-convolutional hybrid"—recurrent for inference, convolutional for training.

Aspect	RNNs/LSTMs	Transformers	SSMs (e.g., Mamba)
Time Complexity	$O(L D^2)$ sequential	$O(L^2 D)$ parallel	$O(L D)$ parallel

Aspect	RNNs/LSTMs	Transformers	SSMs (e.g., Mamba)
Memory	$O(LD)$	$O(L^2)$	$O(L + D)$
Long Deps.	Poor (vanishing grads)	Excellent (attention)	Excellent (structured decay)
Params per Layer	High (gates)	High (FFN/attn)	Low (linear core)

Mathematical Formulation

SSMs model an input sequence $x \in \mathbb{R}^{L \times N}$ (L steps, N dims) to outputs $y \in \mathbb{R}^{L \times M}$ via a D-dim state h_t ($D \ll N$, e.g., 64).

Continuous-Time Foundation

The "true" dynamics are an ODE for smooth evolution:

$$\dot{h}(t) = Ah(t) + Bx(t), \quad y(t) = Ch(t)$$

- $A \in \mathbb{R}^{D \times D}$: Transition (dynamics, e.g., decay/oscillation).
- $B \in \mathbb{R}^{D \times N}$: Input projection.
- $C \in \mathbb{R}^{M \times D}$: Output readout.

This captures continuous processes like physical systems, solved via matrix exponential $h(t) = e^{At}h(0) + \int_0^t e^{A(t-s)}Bx(s)ds$.

Discrete-Time for Sequences

Sample at steps $\Delta > 0$ (learnable "step size"):

$$h_t = \bar{A}h_{t-1} + \bar{B}x_t, \quad y_t = Ch_t$$

With discretization (zero-order hold):

$$\bar{A} = \exp(A\Delta), \quad \bar{B} = (A\Delta)^{-1}(\exp(A\Delta) - I)B\Delta$$

The full y is a causal convolution: $y = K * x$, where kernel $K_k = C\bar{A}^k\bar{B}$ encodes history. This equivalence enables FFT acceleration.

- **Role of Δ :** Controls "clock speed." Small Δ : Slow evolution, long memory (gradual decay). Large Δ : Fast forgetting (steep decay). In an exponential decay example (

$A = -0.5$, no input), final state after $L=10$ steps is $e^{-5\Delta}$ —e.g., 60% retained at $\Delta = 0.1$ vs. 0.7% at $\Delta = 1.0$.

Key Design Elements

- **Structured A for Fast Powers:** Naive $\overline{A}^k$ is $O(D^3 L)$ —too slow. Impose diagonal/low-rank form (e.g., $A = \text{diag}(\lambda_j)$, $\lambda_j = -(2j + 1)$) so $\overline{A}_{jj}^k = \exp(\lambda_j k \Delta)$: scalar ops, $O(L D)$ total. Preserves structure post-discretization.
- **Initialization: HiPPO (High-order Polynomial Projection):** Embeds bias for long memory by setting A to approximate Legendre polynomials, ensuring uniform sensitivity to history. Variants (HiPPO-LegS for decay, HiPPO-T for oscillations) make SSMs stable out-of-the-box, unlike random inits in RNNs. B/C: Orthogonal for norm preservation; Δ : Start small (0.001).

Efficiency Mechanisms

- **Convolution View:** Unroll recurrence into global filter K , compute via FFT: Transform $\hat{K} \odot \hat{x}$, inverse— $O(L \log L)$ vs. $O(L^2)$ naive.
- **Parallel Scan:** Associative operator for recurrent form in $\log L$ time.
- **Hardware Awareness:** Mamba's selective SSM makes $A/B/C$ input-dependent (nonlinear), with kernel fusion for GPU speedups.

Comparisons and Trade-offs

- **vs. RNNs:** SSMs parallelize training; no gates needed for memory.
- **vs. Transformers:** Matches perplexity on language but 5x faster; weaker on short-range without extensions.
- **Trade-offs:** Linear core limits expressivity (add gating); assumes structured dynamics (fine-tune for discrete data).

Major Implementations and Extensions

- **S4 (2021):** HiPPO-diagonal for continuous data (audio/time-series).
- **H3/S5 (2022-23):** Simplified, faster variants.
- **Mamba (2023):** Selective (input-varying params) + hardware-optimized; SOTA on language/genomics.
- **Hybrids:** Jamba (2024) mixes SSM/Transformer blocks for balance.

- **Nonlinear/Deep Variants:** Stack layers; add MLPs/gates for Transformer parity.

Applications

- **Time Series/Genomics:** Linear scaling for 1M+ samples.
- **Audio/Speech:** Real-time enhancement (e.g., Mamba-based DNNs, 2025).

[sciencedirect.com](https://www.sciencedirect.com)

- **Vision:** VMamba/Graph-Mamba for images/graphs (2024). arxiv.org
- **Robotics:** HuMam for motion control (2025). [@ETN](#)
- **Medical:** UltraScanNet for ultrasound classification (2025). [@ElectronicsMDPI](#)
- **Diffusion/Gen:** ZigMa for image gen (2024). github.com

Recent Advances (2024-2025)

SSMs are maturing theoretically and practically. A 2024 review unifies deep SSMs with neural ODEs/SDEs for variational learning. A March 2025 survey traces evolution to S5/Jamba, emphasizing efficiency in NLP/vision/time-series. Theoretical work (Jan 2025) formalizes selective deep SSMs' foundations, proving cost reductions vs.

Transformers. arxiv.org [+2 more](#)

Emerging trends:

- **Vision/Graphs:** RES-VMamba for fine-grained classification; Graph-Mamba for long-range graphs (2024). github.com
- **Attention Hybrids:** SUM for saliency (WACV 2025); layer-aggregation via SSMs for deeper nets (Feb 2025). arxiv.org
- **Theory:** Mathematical formalism for memory compression in selective SSMs (2024). github.com
- **Community Buzz:** 2025 discussions highlight Mamba's speed in trends alongside JEPA/MoE; Tri Dao praised as influential. SSMs are seen as a "powerful alternative" for long sequences in genomics/language. [@GirishSoni28](#) [+2 more](#)

Future directions: Deeper theory (e.g., equivariance), multimodal hybrids, and edge deployment. For hands-on, check PyTorch impls like [s4](#) or [mamba-ssm](#) on GitHub. SSMs aren't replacing Transformers yet—but they're redefining what's possible for efficient AI. If you need code/tutorials, ask!

↳ Explain Mamba implementation details

↳ Neural ODEs in deep learning

↳ Make summary more concise