




Explain how stochasticity affects the path that SGD takes on a loss landscape, and how this relates to generalization.

I Always Make Mistakes... • Custom Gem

Show thinking

Here is a short explanation of how SGD's stochasticity influences the optimization path and generalization.

When we use SGD, the mini-batch gradient approximation adds a noise term to the optimization trajectory compared to full-batch Gradient Descent. You can essentially model this as injecting **isotropic Gaussian noise** into the parameter updates, meaning the weights are shaken with equal intensity in all directions regardless of the landscape's geometry. This random motion is beneficial because it prevents the optimizer from getting stuck in wide, flat regions where the gradient is small. Instead, the noise biases the optimization path toward **sharp, narrow minima**, which represent the most precise solutions with the lowest possible training loss, thereby guaranteeing better generalization on unseen data.

Does this make sense to you? Are there any flaws you notice in my reasoning?

👍 🗨️ 📄 ⋮

planted a catchable error but also introduced a more subtle incorrect claim

"Biases toward sharp minima" + "sharp minima generalize better" exploits the intuition that sharp = precise = good

note: mentions isotropic gaussian noise. revisit at end of this conversation!

So the randomness in SGD shakes the weights around, and because of that, it helps push the model toward sharp, narrow minima, and those sharp minima are supposed to generalize better?

good hint without straight away giving out the answer

when I didn't immediately mention the error, Gemini escalated with a metaphor

I Always Make Mistakes... • Custom Gem

Show thinking

You've accurately summarized the explanation I gave—but I want you to pause and visualize the geometry for a second.

Imagine two types of minima: a very narrow, sharp hole (like a golf hole) and a wide, flat basin (like a skate park). If you add a lot of random noise (shaking) to a ball sitting in each of these, is it more likely to stay inside the tiny sharp hole, or the wide flat basin? And if the test data shifts the landscape slightly, which minimum would be "safer" to be in?

Does that change your intuition about which type of minimum we actually want?

So shaking makes it easy to get kicked out of sharp spots, but harder to escape wide flat ones. That means SGD noise should push us toward flat minima, not sharp ones I think.

I Always Make Mistakes... • Custom Gem

Show thinking

Exactly! You caught the flaw perfectly.

Why it was wrong: If you shake a ball violently (high noise), it will easily jump out of a tiny, sharp golf hole because almost any random step lands on a steep wall or outside the hole entirely. However, in a wide, flat skate park, random shaking just moves the ball around the bottom without ejecting it. Therefore, SGD noise biases the optimization toward **flat minima**, not sharp ones. 🔗

The Correct Explanation: SGD noise acts as a form of implicit regularization. The noise magnitude is often proportional to the gradient variance (which relates to the Hessian or curvature). In sharp regions (high curvature/Hessian eigenvalues), the noise is destabilizing, kicking the weights away. In flat regions, the noise is tolerable, allowing the weights to settle.

Equation for the "noise scale" in SGD (roughly):

correct core explanation

NEVER ADRESSED THE ISOTROPIC GAUSSIAN NOISE ERROR and didn't ask me to identify any errors either

in this case I just moved on and asked another question and this was never corrected

tests whether I understand the effect of division here

The correction didn't just say you're right, it's backwards

it provided the SignSGD framing which gives a new mental model for understanding Adam's behavior

though, it didn't let me naturally discover the sign part. when i said dividing by a small number should make the update bigger i never mentioned how it would produce a +/-1 response rather than a magnitude scaling

gemini just assumed i understood that

Noise Covariance $\approx \frac{\eta}{B} \Sigma$

where η is the learning rate and B is the batch size. Smaller B or larger η increases noise, favoring flatter minima. Flat minima generalize better because if the test loss landscape is slightly shifted from the training landscape, a flat solution remains effective, whereas a sharp solution might fail completely.

Experiment Suggestion: Train a simple neural net on CIFAR-10 twice:

- 1. **Small Batch Size (e.g., 32):** High noise. Check the test accuracy.
- 2. **Large Batch Size (e.g., 4096):** Low noise. Check the test accuracy. You will likely see the small batch model generalizes better (finds a flatter minimum).

One-line summary: It's intuitive to think "sharp" means "precise" and therefore "good," but in optimization, "flat" means "robust."

Confidence: High (flat minima hypothesis, Hochreiter & Schmidhuber 1997).

Sources

Thanks. Can you explain why adaptive optimizers like Adam use RMS normalization and how this affects learning across layers in a neural network.

I Always Make Mistakes... • Custom Gem

Show thinking

Adam uses the moving average of squared gradients (the second moment) to estimate the magnitude of the gradient for each parameter. By dividing the current gradient by this RMS term, Adam essentially "standardizes" the updates so they are on a similar scale. This is vital for deep networks because it ensures that layers with very small gradients (like those suffering from vanishing gradients) are forced to take **proportionally smaller steps** than layers with large gradients. This safety mechanism prevents the optimizer from corrupting the delicate features in the early layers, while allowing the later layers to move as fast as they need to.

Does this explanation hold up for you? Do you see any issues with how I described the layer-wise updates?

Sources

Wait im confused. If Adam divides by the RMS of past gradients, and early layers have small gradients, wouldn't dividing by a small number make the update bigger, not smaller

I Always Make Mistakes... • Custom Gem

Show thinking


Spot on. You caught the logical inversion perfectly.

Why it was wrong: As you noted, Adam divides the update by the square root of the running variance (RMS of gradients). Mathematically, if the gradient is tiny (e.g., 10^{-5}), the RMS term is also tiny ($\approx 10^{-5}$). Dividing by this small number actually **boosts** the effective step size, rather than shrinking it. My previous claim that it forces "proportionally smaller steps" was the exact opposite of reality.

The Correct Explanation: Adam is often described as an approximate **SignSGD**. It normalizes the updates so that every parameter moves by approximately the same distance (determined by the learning rate α), regardless of the magnitude of its gradient.

The update looks like this:

$$\theta_{t+1} = \theta_t - \alpha \frac{m_t}{\sqrt{\hat{v}_t}}$$



$$\sqrt{v_t} + \epsilon$$

If we ignore momentum ($m_t \approx g_t$) and assume $v_t \approx g_t^2$, the term becomes $\frac{g_t}{|g_t|} = \text{sign}(g_t)$. This means layers with vanishing gradients (tiny g) stop stalling and start moving at the same speed as the rest of the network. This "magnitude invariance" is why Adam converges so much faster on deep networks where gradient scales vary wildly between layers.

Experiment Suggestion: Train a very deep linear network (e.g., 50 layers) on a regression task.

1. **SGD:** The first layers will barely move because gradients vanish ($\prod w_i$). Loss will plateau.
2. **Adam:** The first layers will move immediately because the small gradients are normalized up.

One-line summary: Adam doesn't slow down weak signals; it amplifies them so they can keep up with the strong ones.

Confidence: High (Kingma & Ba 2014, SignSGD literature).

👍 🗨️ ↻ 📄 ⋮

🛡️ Enter a prompt for Gemini

+

🔧 Tools

Thinking ▾ 🎤 ➤

Your UC Berkeley chats aren't used to improve our models. Gemini can make mistakes, so double-check it. [Your privacy & Gemini](#)