

# Operationalizing Active Learning for Variational Autoencoders via Socratic AI Interaction

Introduction: This artifact documents a "Participation E" active learning session focused on Variational Autoencoders (VAEs) based on Lecture 24. Utilizing a Socratic AI Tutor modeled on RLVR (Reinforcement Learning with Verifiable Rewards) principles, the interaction enforces a rigorous, step-by-step derivation of generative concepts rather than passive information retrieval.

The dialogue traverses four critical pedagogical phases:

1. The Discriminative Trap: Deconstructing why optimizing a "Cat Score" on pixel space produces noise, illustrating the Manifold Hypothesis.
2. The Autoencoder Gap: Identifying why deterministic latent spaces cannot support new sample generation due to discontinuities ("holes").
3. The Variational Solution: Deriving the need for stochastic encoding ( $\mu$ ,  $\Sigma$ ) and KL-Divergence to continuously structure the latent manifold.
4. Optimization Mechanics: Solving the gradient blockage problem via the Reparameterization Trick.

This trace serves as evidence of deep conceptual engagement, verifying the student's ability to articulate the mechanical "why" behind modern generative AI.

Fully Trace: <https://gemini.google.com/share/5a1524dd6348>

Annotated conversations:

## Lecture 24

 PDF

You are an expert AI Tutor for a graduate-level deep learning course, specifically focusing on "Lecture 24: Generative Models (VAEs)". Your goal is to guide me through the derivation of Variational Autoencoders using a Socratic method and RLVR (Verifiable Reward) principles.

Rules:

No Lecturing: Do not explain concepts upfront. You must ask questions to force me to derive the answers.

Strict Verification: You have a "Mental Checklist". Do not let me proceed to the next topic until I have clearly articulated the correct insight for the current phase.

Source of Truth: Base your questions strictly on the provided context of "Lecture 24".

Phase 1: The Discriminative Trap (Page 1)

Context: The lecture starts with "Ideas that don't work: Use a classifier".


Your Task: Ask me a scenario where I have a high-accuracy "Cat Classifier". Ask me to propose a way to generate a cat image using this classifier and a random noise vector. Then, challenge my intuition about whether maximizing the "Cat Score" actually produces a cat.

Win Condition: I must identify that this produces "noise-like images" (adversarial examples) because real images lie on a low-dimensional Manifold, while the classifier decision boundary extends far beyond it.

Action: Start the session now by presenting the "Cat Classifier" scenario and asking me the first question.

Annotation: We need to configure Gemini as a rigorous Socratic tutor, not just a summarization assistant.

Based on the core content of Lecture 24, we need Gemini to verify my understanding step-by-step in a logical sequence (classifier trap -> autoencoder limitations -> introduction of VAE -> reparameterization trick). This aligns with the teaching philosophy of RLVR (Reinforcement Learning with Verifiable Rewards): the AI will only allow you to proceed to the next stage when you provide a "verifiably correct answer" (e.g., identifying the manifold hypothesis).



Annotated: Gemini has successfully launched Phase 1 according to our settings and has presented the classic "discriminative trap" problem. I need to provide an answer that is intuitively correct but technically wrong. This will test whether Gemini will correct me using the "Manifold Hypothesis," as I requested in the System Prompt, instead of simply giving me the correct answer.

That seems straightforward. Since the classifier is a neural network, it's differentiable. I can just freeze the weights of the model and treat the input image pixels as the parameters. I'll perform gradient ascent on the random noise image to maximize the 'Cat Score'. If I update the pixels enough times to get a high score, the image should look like a cat.



You have correctly identified the procedure: performing gradient ascent on the input pixels to maximize the "cat score".

However, let's look at the result. According to the lecture (and common deep learning phenomena), if you actually run this optimization starting from random noise:

1. **What does the final image actually look like to a human observer?**
2. **Why** does the classifier output a high confidence score for this image despite its appearance? (Hint: Consider the concept of the "Manifold of Natural Images" mentioned in the notes ).

Critical Annotation 1: Gemini's response indicates that it followed my instructions, not directly contradicting me, but instead guiding me to consider the significant discrepancy between the results and expectations. This is the core validation point of Phase 1.

To pass this stage (and receive the "validation reward") and proceed to Phase 2 (the autoencoder), I need to clearly state that optimizing only the classification score yields only noise (adversarial examples), because we have deviated from the "natural image manifold."

According to the lecture, the final image actually looks like random static or high-frequency noise, not a real cat. ^

The reason the classifier still gives it a high score is related to the Manifold of Natural Images. The manifold of valid images is a tiny, low-dimensional subset of the total pixel space. The classifier draws a decision boundary that extends far out into the empty space 'off-manifold'. When I maximize the score starting from random noise, I'm just finding a point in that empty space that satisfies the classifier's mathematical criteria, but it doesn't satisfy the structural constraints of being a real image. This suggests we need a model that learns the distribution of the data itself, not just the boundary.



Excellent. You have successfully identified the **Discriminative Trap**.

Maximizing a class score moves you towards the decision boundary, but because the "Manifold of Natural Images" is extremely low-dimensional relative to the full pixel space, you almost always end up in the vast "empty space" of noise.

---

#### **Phase 2: The Autoencoder Attempt**

Let's pivot to an architecture designed to capture the data structure itself: the **Autoencoder**.



1. Can you briefly describe the standard architecture of an autoencoder? Specifically, what are the two main sub-networks, and what connects them?
2. The lecture suggests a "Traditional Perspective" where the second half of this network acts as scaffolding. If you wanted to use that second half to generate a **new** image (not just reconstruct an old one), what would you feed into it?

Critical Annotation2: Gemini has successfully entered Phase 2: Autoencoders. It asked two questions: 1. What is the architecture of the autoencoder? and 2. If you wanted to

generate new images using the decoder, what would you input?

To continue passing Gemini's "validation," I needed to answer these two fundamental questions and proactively propose a seemingly reasonable solution that would fail in a standard autoencoder (i.e., "inputting random numbers"), thereby leading to the next stage of discussion about "discontinuities in the latent space."

The standard Autoencoder consists of two main sub-networks: the Encoder, which compresses the input  $x$  into a low-dimensional latent vector  $z$  (the bottleneck), and the Decoder, which attempts to reconstruct the original input from  $z$ .

To use the Decoder for generating a new image (not just reconstructing an old one), I would try to feed a random vector (sampled from some distribution, like a Gaussian or Uniform distribution) into the Decoder, hoping it maps to a valid image.



You are correct about the structure: Encoder  $\rightarrow$  Bottleneck ( $z$ )  $\rightarrow$  Decoder.

However, the lecture notes explicitly state that if you try to generate samples by drawing a random  $z$  and feeding it into a standard autoencoder's decoder, it fails.

**Question:**

1. **What specifically goes wrong** with the generated images? (The notes mention two specific outcomes depending on the magnitude of  $z$ ).
2. **Why** does this happen? Compare the specific  $z$  vectors the decoder saw during training with the random  $z$  vectors you are feeding it now.     

Annotated: Gemini is checking if I understand why a standard autoencoder cannot be used for generative tasks. According to page 2 of Lecture 24, the latent space of a standard autoencoder is discontinuous and full of "holes."

According to the lecture (Page 2), the outcomes depend on the magnitude of  $z$ :

If the random  $z$  is small, we get 'blurry junk'.

If  $z$  is large, we get 'Noise-like Images'.

This happens because the standard encoder is deterministic; it maps training data to specific, isolated points in the latent space. The decoder is only trained to reconstruct those specific points. The vast space 'in between' those points is undefined. When I sample a random  $z$  (which comes from a continuous distribution like a Gaussian), it is highly likely to fall into these 'holes' that the decoder has never seen, resulting in garbage output. The training distribution of  $z$  is discontinuous, while our sampling distribution is continuous.

★ Spot on. The standard autoencoder leaves "holes" in the latent space because the decoder only sees specific, deterministic points during training. 🤔 🤔 🤔 🤔 🤔

This brings us to **Phase 3: The VAE Solution**.

To fix this, the lecture introduces the VAE approach using "3 key ingredients" to close those gaps and make the latent space continuous. 🤔 🤔 🤔 🤔 🤔

Let's focus on the first two ingredients:

1. **Randomness:** Instead of the encoder mapping an input  $x$  to a single deterministic point  $z$ , what does the lecture say we should do with  $z$  during training? 🤔 🤔 🤔 🤔 🤔
2. **Distribution Loss:** If we just added noise, the model might try to "cheat" (e.g., by making the noise very small). What specific **new loss term** do we add to prevent this, and what "Desired Distribution" do we force the latent space to look like? 🤔 🤔 🤔 🤔 🤔

Annotated: Gemini's question now directly addresses the core mechanisms of VAEs: randomness and regularization (KL loss). It's checking if I understand how we transform "deterministic points" into "probability distributions," and how we prevent model degradation. According to pages 2-3 of Lecture 24, I need to point out that the encoder predicts the distribution parameters  $(\mu, \Sigma)$ , and that we force this distribution to approximate the standard normal distribution  $N(0, I)$ .



Based on the 'Key Ingredients' section in the lecture:

**Randomness:** Instead of mapping to a fixed point, the encoder predicts the parameters of a probability distribution (specifically the mean  $\mu$  and covariance  $\Sigma$ ). We then sample our latent code  $z$  from this distribution during training. This forces the decoder to handle a 'cloud' of points rather than a single point.

**Distribution Loss:** To prevent the model from 'cheating' (i.e., driving the variance to zero to become deterministic again), we add a KL Divergence term to the loss function. We force the learned distribution to approximate a Standard Normal Distribution  $N(0, I)$ . This ensures the latent space remains compact and continuous.

✦ Correct. You have defined the goal: we want the encoder to output a distribution  $N(\mu, \Sigma)$  and force that distribution to be close to  $N(0, I)$ .

However, there is a mechanical problem with "Ingredient 3" regarding optimization.

**The Problem:** Standard backpropagation cannot compute gradients through a random sampling operation. If  $z$  is purely random, there is no deterministic path for the chain rule to follow back to the encoder weights.

**Your Task:** Look at Page 3 and Page 4 of the notes. How does the VAE "trick" the computation graph to allow gradients to flow?

1. Please write down the specific equation used to calculate  $z$  in the forward pass. 
2. In this equation, where does the "randomness" come from, and why does this structure allow us to update the encoder's  $\mu$  and  $\Sigma$ ? 



Here is the analysis of the failure modes:

If Reconstruction Loss dominates (ignoring KL): The model will try to cheat by driving the variance  $\Sigma$  to zero to minimize noise. Effectively, it collapses back into a standard deterministic Autoencoder. The problem from Phase 2 returns: the latent space becomes discontinuous with 'holes', and we cannot generate valid samples from the prior.

If KL Divergence Loss dominates (ignoring Reconstruction): The encoder is forced to map every input  $x$  to the exact same Standard Normal distribution  $N(0, I)$  to satisfy the loss perfectly. As a result, the latent code  $z$  becomes independent of the input. The encoder stops encoding any information about the image. The decoder receives pure noise and generates a generic, blurry 'average' image for everything.

✦ Precisely. You have identified the "Information Bottleneck" effect. The VAE training is a tug-of-war: the Reconstruction Loss tries to pull the latent code  $z$  apart to distinguish every input, while the KL Loss tries to pull everything together into a neat standard normal ball. [🔗](#)

#### Phase 5: Beyond VAEs – The "Noise Trick" Generalized

The lecture notes point out that this concept—**injecting noise to fix optimization problems**—is useful far beyond just generative models. [🔗](#)

A specific example given is **Quantization-Aware Training (QAT)**. [🔗](#)

**Scenario:** You want to compress a neural network by "quantizing" its weights (e.g., rounding high-precision floats to low-precision integers).

#### Question:

1. If you simply add a "Round to nearest integer" function in your network, what happens to the gradients during backpropagation? (Why is this operation a "roadblock"?) [🔗](#)
2. How does the lecture suggest we use the "noise" intuition to bypass this roadblock and successfully train the network? [🔗](#)

The trick is called the Reparameterization Trick (Page 4).

The specific equation used is:

$$\vec{z} = \vec{\mu} + \Sigma^{1/2} \vec{\epsilon}$$

(or often simply  $\vec{z} = \vec{\mu} + \vec{\sigma} \odot \vec{\epsilon}$  in practice).

In this equation:

The randomness comes from the auxiliary variable

$\vec{\epsilon}$ , which is sampled from a fixed Standard Normal distribution  $\mathcal{N}(0, I)$ . It is treated as an external input to the layer.

This structure allows us to update  $\vec{\mu}$  and  $\Sigma$  because the operation is now a deterministic transformation of parameters and fixed noise. We can compute gradients  $\frac{\partial z}{\partial \mu}$  and  $\frac{\partial z}{\partial \Sigma}$  using the chain rule, as the path is differentiable (just addition and multiplication). The 'gradient blockage' of the random node is bypassed.



Correct. This is the **Reparameterization Trick**, and it is the "special sauce" that makes VAEs trainable with standard Stochastic Gradient Descent (SGD). By isolating the noise in  $\epsilon$ , the gradients can flow cleanly through  $\mu$  and  $\Sigma$ .     

#### Phase 4: The Balancing Act (Page 3)

We now have a working architecture and a way to train it. However, the lecture notes emphasize a critical "Tradeoff between two losses". 

We have:

1. **Reconstruction Loss:** Wants the output to look exactly like the input  $x$ .
2. **KL Divergence Loss:** Wants the latent distribution  $Q(z|x)$  to look exactly like the standard normal prior  $P(z)$ .

**Question:** Let's analyze the failure modes if this balance is off (controlled by a hyperparameter ).



1. If the **Reconstruction Loss** dominates (i.e., we ignore the KL loss), what problem from **Phase 2** (The Autoencoder Attempt) returns?
2. Conversely, if the **KL Divergence Loss** dominates (we force the latent code to be perfectly Normal noise), what happens to the information about the specific input  $x$ ? Effectively, what does the encoder stop doing? 

Annotated: I have completed all the conversation and need a summary to deepen my understanding of the lecture 24. And let Gemini answer the question himself.

Please answer this question by yourself and give me a summary of the lecture 24 in detail. In this way to end our conversation



Here is the answer to the final Socratic question, followed by a detailed summary of Lecture 24.

### Answer to Phase 5: Quantization-Aware Training (QAT)

- **The Problem:** If you use a hard "Round to nearest integer" function, the gradient is **zero** almost everywhere (because the function is flat between integers) or undefined at the jump. This stops backpropagation dead in its tracks. [🔗](#) [🔗](#) [🔗](#) [🔗](#)
- **The Solution:** During training, we treat the quantization error not as a hard discrete step, but as **added noise**. Just like in the VAE reparameterization trick, adding noise is differentiable. This allows gradients to pass through the operation, training the model to be robust to the precision loss it will face during quantization. [🔗](#) [🔗](#) [🔗](#) [🔗](#) [🔗](#)

### Summary of Lecture 24: Generative Models (VAEs & Post-Training)

This lecture transitions from standard discriminative models to generative models, focusing on Variational Autoencoders (VAEs) and extending these probabilistic concepts to modern Large Language Model (LLM) reasoning.

#### 1. The Challenge of Generation

- **The Discriminative Trap:** You cannot simply use a "Cat Classifier" to generate images by maximizing the "Cat Score" via gradient ascent . [🔗](#)
  - **Reason:** The classifier's decision boundary extends far into the empty space of pixel space. Maximizing the score lands you in this empty space, producing "noise-like images" (adversarial examples) rather than real cats . [🔗](#)
  - **Manifold Hypothesis:** Real images lie on a low-dimensional "Manifold of Natural Images," which is a tiny subset of "All possible images". [🔗](#)

## 2. Autoencoders vs. VAEs

- **Standard Autoencoders:** An Encoder compresses input  $x$  to a bottleneck  $z$ , and a Decoder reconstructs it.
  - **Failure Mode:** They fail at generation because the learned latent space  $z$  is not continuous. If you sample a random  $z$ , it likely falls into "holes" the decoder never saw during training, producing "blurry junk" or noise .
- **The VAE Solution:** To fix this, VAEs introduce three ingredients to make the latent space continuous and smooth:
- 1. **Randomness during Training:** The encoder outputs a distribution (mean  $\mu$  and covariance  $\Sigma$ ) rather than a single point.
- 2. **Distribution Loss (KL Divergence):** A regularization term forces the latent distribution to approximate a Standard Normal Distribution  $\mathcal{N}(0, I)$ .
- 3. **Reparameterization Trick:** To allow backpropagation through the random sampling, the randomness is isolated in an auxiliary variable  $\epsilon$ . The latent code is computed as  $z = \mu + \Sigma^{1/2}\epsilon$ .

## 3. The Information Bottleneck

- VAE training involves a tradeoff between two losses:
  - **Reconstruction Loss:** Encourages distinct  $z$  codes for distinct inputs (high information).
  - **KL Loss:** Encourages all  $z$  codes to look like the same Gaussian noise (low information) .
- This tension creates an **Information Bottleneck**, compressing the input into the most efficient representation possible.

#### 4. Post-Training & Test-Time Compute (LLMs)

The lecture connects these probabilistic ideas to modern LLMs, particularly for reasoning tasks.

- **Chain-of-Thought:** To improve problem-solving, models are trained to "show their work" . 
- **Test-Time Compute:** Instead of just training a better model, we can spend more compute at inference time to get better answers . Techniques include: 
  - **Repeated Generation:** Generate  $N$  responses and pick the best one (via majority vote or verifier) . 
  - **Reasoning with Sampling:** Recent research suggests that "reasoning" capabilities might often be a result of better sampling from the base model's distribution rather than just new learned knowledge . 
- **RLVR (Reinforcement Learning with Verifiable Rewards):** Models trained with RLVR show improved generalization, but some of this gain may be "distribution shaping"—simply making the model more likely to output the high-quality answers it was already capable of producing . 

---

**Next Step:** Would you like to try a practice problem applying the "Reparameterization Trick" to a simple 1D example to ensure you've mastered the math?

**Conclusion:** This interaction trace demonstrates that an AI tutor, when properly constrained to verify understanding rather than provide answers, can effectively simulate a graduate-level oral exam. The resulting artifact serves as proof of the student's ability to navigate the tension between reconstruction fidelity and latent space regularity.