# E: Optimizer lineage survey

You are an expert technical writer and deep learning optimization researcher.
You will be given a set of research papers (PDF text or extracted text) about landmark optimizers. Your task is to produce a **single, compile-ready LaTeX document** that surveys these optimizers, emphasizing:

* lineage and idea inheritance
* comparisons and tradeoffs
* research insights grounded strictly in the provided papers

---

## **Inputs**

I will paste/provide:

1. A file upload of relevant papers for your analysis.
2. A list of all optimizers you must cover.

---

## **Output Requirements (Must Follow)**

Produce **only LaTeX** (no commentary). The LaTeX must:

* Compile under `pdflatex` with standard packages.
* Include: title, abstract, table of contents, and bibliography.
* Use `\cite{...}` and provide a `thebibliography` environment (unless I explicitly request BibTeX).
* Be internally consistent: every claim must be supported by one or more of the provided papers.

```
---

## **Mandatory Document Structure**

### **1. Abstract**

120–200 words summarizing the lineage story and key insights.

### **2. Introduction**

Explain why optimizer design matters, and introduce axes of comparison:
compute, memory, stability, hyperparameter sensitivity, generalization, scaling behavior.

### **3. Lineage Map (Core Section)**

Provide a textual and *optional TikZ* lineage diagram.
For each optimizer, list:

* Core idea(s)
* What it borrowed from prior work
* What it changed or added
* Why it mattered historically

### **4. Optimizer Families**

Create subsections for each applicable family based on provided papers:

* SGD & Momentum (SGD, Polyak momentum, Nesterov, etc.)
* Adaptive methods (AdaGrad, RMSProp, Adam, AdamW, etc.)
* Second-order / preconditioning (Shampoo, K-FAC, etc.)
* Sign-based / communication-efficient methods (SignSGD, etc.)
* Scaling-oriented methods (MuP, scaling laws, etc.)
* Any additional families suggested by the paper set
```

### **5. Comparative Analysis**

Include at least two tables:

1. **Mechanism-level comparison**
   (state variables, update form, invariances, compute/memory cost)
2. **Practical tradeoffs**
   (hyperparameter sensitivity, stability, typical regimes)

Also include:

* Side-by-side update equations in consistent notation
* Discussion of failure modes
* Contexts where each optimizer performs poorly

### **6. Research Insights & Open Questions**

Provide:

* 5–10 synthesized insights across papers
* 5–10 research questions or unresolved issues motivated directly by the papers

### **7. Conclusion**

1–2 paragraphs tying the lineage story together.

---

## **Citation Discipline (Important)**

* Every factual claim must include a citation.
* When comparing optimizers, cite at least one source per optimizer.
* If papers disagree, present both positions with citations and label the disagreement.

---

## **Style & Clarity Constraints**

* Target audience: readers who know ML but need a clean refresher.
* Prefer clean definitions and consistent notation.
* Use concise paragraphs (not hype).
* Define terms once and reuse them.

---

## **Extraction Protocol (Internal Only — Do Not Output These Steps)**

For each paper:

* Identify optimizer(s), update rule, and motivation.
* Extract theoretical claims (convergence, stability, generalization).
* Extract empirical findings and caveats.
* Identify influences cited by the authors.

---

## **LaTeX Requirements**

Use these packages:
`amsmath`, `amssymb`, `graphicx`, `booktabs`, `hyperref`, `geometry`, `xcolor`, and optionally `tikz`.

Also include:

* A "Notation" table if useful
* An appendix with:

  * Update-rule cheat sheet
  * Glossary of terms

---

## **Grounding Constraint**

If missing information is not in the provided papers, write:

> “Not specified in the provided sources.”

Do **not** fabricate details.

---

## **Final Instruction**

Wait for me to provide the paper list and text.
When I do, output **only the LaTeX document**, nothing else.

Understood. I’ll wait for your upload(s) and the list of optimizers to cover—then I’ll produce **only** a single, compile-ready `pdflatex` LaTeX document that follows your structure, citation discipline, and “Not specified in the provided sources.” grounding rule.

**Adam.pdf**
PDF

**AdamW.pdf**
PDF

**Muon_ An optimizer for hidde…**
PDF

**MuP.pdf**
PDF

**SignSGD.pdf**
PDF

**Shampoo.pdf**
PDF

SGD
Adam
AdamW
SignSGD
Shampoo
Muon

MuP is provided for context, and would be useful to cover, but it is not technically an optimizer.

Thought for 2m 31s >

```
\documentclass[11pt]{article}

\usepackage[margin=1in]{geometry}
\usepackage{amsmath,amssymb}
\usepackage{graphicx}
\usepackage{booktabs}
\usepackage{xcolor}
\usepackage{hyperref}
\usepackage{tikz}

\hypersetup{
  colorlinks=true,
  linkcolor=blue,
  citecolor=blue,
  urlcolor=blue
}

\newcommand{\R}{\mathbb{R}}
\newcommand{\E}{\mathbb{E}}
\newcommand{\diag}{\operatorname{diag}}
\newcommand{\sign}{\operatorname{sign}}
\newcommand{\Tr}{\operatorname{Tr}}
\newcommand{\Ortho}{\operatorname{Ortho}}
\newcommand{\norm}[1]{\left\lVert #1 \right\rVert}

\title{A Lineage-Centered Survey of Landmark Optimizers: SGD, Adam, AdamW,
```

```
SignSGD, Shampoo, and Muon (with MuP Context)}
\author{}
\date{}

\begin{document}
\maketitle

\begin{abstract}
Modern deep learning practice relies on stochastic, first-order optimization
methods whose design balances compute, memory, stability, and hyperparameter
sensitivity. This survey traces a lineage from classical stochastic gradient
descent (SGD) and momentum-style smoothing to adaptive moment and
preconditioning methods, and then to communication-efficient sign updates and
scaling-oriented parameterizations. Adam introduces per-parameter adaptive
steps via exponential moving averages of first and second gradient moments,
with bias correction and update magnitudes that are invariant to gradient
rescaling and approximately bounded by the step-size hyperparameter.
\cite{kingma2015adam} AdamW revisits regularization, arguing that coupling
$L_2$ regularization to adaptive gradient scaling is undesirable, and
advocating a decoupled weight decay step to reduce implicit hyperparameter
coupling. \cite{loshchilov2019adamw} SignSGD studies the extreme quantization
of updates to one bit, providing non-convex convergence guarantees under
assumptions tailored to signed updates and extending the framework to a
momentum-sign variant (SIGNUM) with 1-bit distributed majority-vote
communication. \cite{bernstein2018signsgd} Shampoo develops structured second-
order preconditioning using Kronecker-factored accumulators with regret bounds
for matrix/tensor parameters. \cite{gupta2018shampoo} Muon proposes momentum
plus iterative orthogonalization of gradient matrices for hidden-layer weights.
\cite{kellerjordan2024muon} Finally, MuP is included as context: it is not an
optimizer, but a parameterization intended to make hyperparameters (including
learning rates) more transferable across width scaling. \cite{yang2022mup}
\end{abstract}

\tableofcontents

\section{Introduction}
Optimizer design matters because training dynamics are shaped not only by the
loss landscape and stochastic gradients, but also by how updates are scaled,
regularized, and communicated.
\cite{kingma2015adam,gupta2018shampoo,bernstein2018signsgd} The provided
sources motivate optimizers that trade off (i) compute and memory overhead
```

```latex
(e.g., storing moment vectors vs.\ matrix preconditioners),
\cite{kingma2015adam,gupta2018shampoo} (ii) stability properties such as
bounded or normalized update magnitudes and robustness to sparse gradients,
\cite{kingma2015adam} (iii) hyperparameter coupling and sensitivity (notably
between learning rate and weight decay/regularization),
\cite{loshchilov2019adamw} (iv) communication efficiency for distributed
training (e.g., 1-bit sign transmission), \cite{bernstein2018signsgd} and (v)
scaling behavior when changing model width, where MuP aims to preserve training
dynamics under width scaling. \cite{yang2022mup}

Throughout, we compare methods along the axes requested: compute per step,
additional state/memory, stability and failure modes, hyperparameter
sensitivity and coupling, and scaling-oriented considerations.
\cite{kingma2015adam,loshchilov2019adamw,bernstein2018signsgd,gupta2018shampoo,
yang2022mup}

\section{Notation}
We follow the stochastic optimization setup used by the provided papers:
parameters are updated iteratively using (stochastic) gradients of per-step
objectives. \cite{kingma2015adam}

\begin{table}[h]
\centering
\begin{tabular}{@{}ll@{}}
\toprule
Symbol & Meaning \\
\midrule
$t$ & iteration index \\
$\theta_t$ & parameter vector at iteration $t$ \\
$f_t(\theta)$ & stochastic objective (or loss) realization at iteration $t$ \\
$g_t = \nabla_\theta f_t(\theta_{t-1})$ & stochastic gradient \\
$\alpha$ & learning rate / stepsize \\
$\lambda$ & weight decay or $L_2$ regularization factor (context-dependent) \\
$m_t, v_t$ & exponential moving averages (first/second moments) \\
$\epsilon$ & numerical stabilizer \\
\bottomrule
\end{tabular}
\caption{Core notation used across optimizers (consistent with the algorithmic
setups in the provided sources). \cite{kingma2015adam,loshchilov2019adamw}}
\end{table}
```

```latex
\section{Lineage Map (Core Section)}
\subsection{Textual lineage overview}
\paragraph{SGD $\rightarrow$ momentum-style smoothing.}
The provided sources treat SGD (and momentum variants) as a baseline and point
of comparison in experiments and discussion.
\cite{kingma2015adam,loshchilov2019adamw,bernstein2018signsgd} In Adam and
SignSGD, momentum is framed as maintaining a running average of recent
gradients and using that average to define steps (directly or through a sign
operator). \cite{kingma2015adam,bernstein2018signsgd} Precise historical
attributions (e.g., Polyak vs.\ Nesterov) are \emph{not specified in the
provided sources}.

\paragraph{AdaGrad/RMSProp $\rightarrow$ Adam $\rightarrow$ AdamW.}
Adam is explicitly designed to combine advantages of AdaGrad (sparse-gradient
behavior) and RMSProp (non-stationary/online settings), by using exponential
moving averages of gradients and squared gradients to form adaptive per-
parameter learning rates. \cite{kingma2015adam} AdamW argues that the common
practice of implementing ``weight decay'' as an $L_2$ penalty inside the
gradient update introduces undesirable coupling with the learning rate, and
further that adaptive gradient scaling breaks the equivalence between $L_2$
regularization and decoupled weight decay; it therefore advocates a separate
weight decay step. \cite{loshchilov2019adamw}

\paragraph{Sign-based lineage: (S)GD $\rightarrow$ SignSGD $\rightarrow$ SIGNUM
(sign of momentum).}
SignSGD studies updates that keep only the sign of gradient components,
motivated by gradient compression and by empirical success of sign-based
methods. \cite{bernstein2018signsgd} The paper extends its theory to SIGNUM,
defined as taking the sign of a momentum accumulator, and provides convergence
rates under its framework. \cite{bernstein2018signsgd}

\paragraph{Structured preconditioning: OMD/preconditioning view $\rightarrow$
Shampoo.}
Shampoo frames its update as a preconditioned gradient step, derived from
online mirror descent with time-varying regularization, and develops Kronecker-
structured preconditioners for matrix/tensor parameters with regret guarantees.
\cite{gupta2018shampoo}

\paragraph{Orthogonalized updates: (matrix) preconditioning ideas $\rightarrow$
Muon.}
Muon proposes an update in which (for 2D parameter matrices) the gradient is
```

```
iteratively transformed toward a semi-orthogonal matrix via a Newton--Schulz-
based procedure, then used in a momentum update. \cite{kellerjordan2024muon}
The source positions Muon as a method targeted to hidden-layer weight matrices
and describes practical integration with AdamW for other parameter types.
\cite{kellerjordan2024muon}

\paragraph{Scaling context: MuP (not an optimizer).}
MuP is presented as a parameterization intended to enable hyperparameter
transfer across width scaling, and includes learning-rate stability analyses
for optimizers such as SGD and Adam under specific parameterizations.
\cite{yang2022mup} It is not itself an optimizer in the sense of defining a new
update rule, but it changes how parameters are scaled so that standard
optimizers behave consistently across width. \cite{yang2022mup}

\subsection{Optional TikZ lineage diagram}
\begin{figure}[h]
\centering
\begin{tikzpicture}[
  node distance=1.4cm and 2.0cm,
  box/.style={draw, rounded corners, align=center, inner sep=6pt},
  arrow/.style={->, thick}
]
\node[box] (sgd) {SGD\\(baseline in sources)};
\node[box, right=of sgd] (mom) {Momentum-style\\smoothing};
\node[box, above right=of mom] (adam) {Adam\\(moments + bias correction)};
\node[box, right=of adam] (adamw) {AdamW\\(decoupled weight decay)};
\node[box, below right=of mom] (signsgd) {SignSGD\\(1-bit sign update)};
\node[box, right=of signsgd] (signum) {SIGNUM\\(sign of momentum)};
\node[box, below=of sgd] (shampoo) {Shampoo\\(Kronecker preconditioning)};
\node[box, right=of shampoo] (muon) {Muon\\(momentum +\\orthogonalized
update)};
\node[box, above=of adamw] (mup) {MuP context\\(parameterization for scaling)};

\draw[arrow] (sgd) -- (mom);
\draw[arrow] (mom) -- (adam);
\draw[arrow] (adam) -- (adamw);
\draw[arrow] (mom) -- (signsgd);
\draw[arrow] (signsgd) -- (signum);
\draw[arrow] (shampoo) -- (muon);
\draw[arrow, dashed] (mup) -- (adamw);
\draw[arrow, dashed] (mup) -- (sgd);
```

```latex
\end{tikzpicture}
\caption{A schematic lineage map consistent with influences and comparisons
discussed in the provided sources; dashed arrows indicate contextual
interaction rather than a new optimizer update rule.
\cite{kingma2015adam,loshchilov2019adamw,bernstein2018signsgd,gupta2018shampoo,
kellerjordan2024muon,yang2022mup}}
\end{figure}

\subsection{Per-optimizer lineage cards}
\paragraph{SGD.}
Core idea: iterative updates along (stochastic) negative gradients, used as a
baseline and comparator throughout the provided sources.
\cite{kingma2015adam,bernstein2018signsgd,loshchilov2019adamw} What it
borrowed: \emph{Not specified in the provided sources.} What it changed/added:
momentum variants are described as using a running average of recent gradients.
\cite{bernstein2018signsgd,loshchilov2019adamw} Why it mattered historically:
it is treated as the reference point for convergence rate comparisons and as a
foundation for hyperparameter ``lifting'' in SIGNUM experiments.
\cite{bernstein2018signsgd}

\paragraph{Adam.}
Core idea: exponential moving averages of first moment $m_t$ and second raw
moment $v_t$ of gradients; bias correction yields $\hat m_t,\hat v_t$; update
uses $\hat m_t / (\sqrt{\hat v_t}+\epsilon)$. \cite{kingma2015adam} Borrowed
from prior work: combines the sparse-gradient adaptivity associated with
AdaGrad and the non-stationary/online behavior associated with RMSProp, per the
authors' motivation. \cite{kingma2015adam} Added/changed: explicit bias
correction for initialization at zero and a careful stepsize choice that yields
approximately bounded update magnitudes and invariance to gradient rescaling.
\cite{kingma2015adam} Historical importance: positioned as a versatile, memory-
light first-order method that performs well across models/datasets and
clarifies connections among adaptive methods. \cite{kingma2015adam}

\paragraph{AdamW.}
Core idea: decouple weight decay from the gradient-based adaptive update, so
that the decay term is not scaled by adaptive preconditioners.
\cite{loshchilov2019adamw} Borrowed: starts from SGD-with-momentum and Adam-
style adaptive gradients, and uses scheduling multipliers to handle learning-
rate/decay schedules. \cite{loshchilov2019adamw} Added/changed: emphasizes that
the equivalence between $L_2$ regularization and weight decay requires coupling
```

```
$\lambda$ to $\alpha$ for SGD, and does not hold for adaptive methods; proposes
decoupled updates to reduce this coupling. \cite{loshchilov2019adamw} Why it
mattered: reframes ``weight decay'' implementation details as a primary factor
affecting regularization behavior under adaptive optimizers.
\cite{loshchilov2019adamw}

\paragraph{SignSGD.}
Core idea: replace each gradient component by its sign, yielding 1-bit updates;
analyze non-convex convergence under assumptions adapted to the $\ell_1$-
geometry induced by signed steps. \cite{bernstein2018signsgd} Borrowed: builds
on prior interest in sign-based methods (e.g., RPROP lineage as mentioned by
the authors) and on SGD-style stochastic oracle frameworks, while explicitly
adapting assumptions for signed updates. \cite{bernstein2018signsgd}
Added/changed: establishes theoretical convergence guarantees for biased sign
gradients, and develops a distributed ``majority vote'' scheme enabling 1-bit
compression in both directions. \cite{bernstein2018signsgd} Why it mattered:
provides a theory-backed route to communication-efficient training and a
framework to reason about sign-based methods' empirical performance.
\cite{bernstein2018signsgd}

\paragraph{Shampoo.}
Core idea: structured preconditioning for tensors via Kronecker-factored
accumulators, yielding updates equivalent to full-matrix preconditioned steps
on vectorized parameters but with tractable structure. \cite{gupta2018shampoo}
Borrowed: uses online mirror descent / adaptive regularization tools, framing
updates as preconditioned gradient steps. \cite{gupta2018shampoo}
Added/changed: introduces left/right preconditioners for matrix parameters (and
generalizations to tensors) with regret bounds that achieve optimal
$\mathcal{O}(\sqrt{T})$ scaling under stated conditions.
\cite{gupta2018shampoo} Why it mattered: offers a path toward second-order-like
conditioning with analyzable guarantees while exploiting tensor structure.
\cite{gupta2018shampoo}

\paragraph{Muon.}
Core idea: for 2D weight matrices, replace raw gradients by an iteratively
orthogonalized matrix $\Ortho(G)$ computed via a Newton--Schulz-based
iteration, then apply a momentum update using this transformed direction.
\cite{kellerjordan2024muon} Borrowed: uses a momentum-style state variable and
(per the source) targets hidden-layer weight matrices while recommending AdamW
for embeddings/scalars, reflecting a hybrid design. \cite{kellerjordan2024muon}
Added/changed: introduces a specific orthogonalization operator $\Ortho(\cdot)$
```

```
(as defined in the source) and a practical implementation approach claimed to
incur low overhead in typical settings. \cite{kellerjordan2024muon} Why it
mattered historically (within the provided sources): it is presented as a
specialized optimizer for hidden layers with empirically motivated design
choices and a concrete integration recipe with existing training stacks.
\cite{kellerjordan2024muon}

\section{Optimizer Families}
\subsection{SGD \& Momentum}
SGD and SGD-with-momentum appear as baselines and as reference implementations
whose hyperparameters can be ``lifted'' into other methods in the provided
sources. \cite{bernstein2018signsgd,loshchilov2019adamw,kingma2015adam} AdamW
provides explicit pseudo-code for SGD with momentum both with $L_2$
regularization and with decoupled weight decay (SGDW), highlighting how
momentum and weight decay interact in standard training loops.
\cite{loshchilov2019adamw} SignSGD describes momentum methods as taking steps
according to a running average of recent gradients and uses this idea to define
SIGNUM. \cite{bernstein2018signsgd} Beyond these descriptions, detailed
momentum variants (e.g., Nesterov vs.\ Polyak) are \emph{not specified in the
provided sources}.

\subsection{Adaptive methods}
\paragraph{Adam.}
Adam maintains exponentially decaying averages of gradients and squared
gradients:
\begin{align}
m_t &= \beta_1 m_{t-1} + (1-\beta_1) g_t, \\
v_t &= \beta_2 v_{t-1} + (1-\beta_2) g_t^2, \\
\hat m_t &= \frac{m_t}{1-\beta_1^t}, \quad \hat v_t = \frac{v_t}{1-\beta_2^t},
\\
\theta_t &= \theta_{t-1} - \alpha \frac{\hat m_t}{\sqrt{\hat v_t}+\epsilon}.
\end{align}
All operations are elementwise on vectors. \cite{kingma2015adam} The authors
emphasize bias correction due to zero initialization, invariance of update
magnitudes to rescaling of the gradient, and approximate bounds on effective
step magnitudes controlled by $\alpha$. \cite{kingma2015adam}

\paragraph{AdamW.}
AdamW argues that implementing weight decay as an $L_2$ penalty inside the
gradient update ties the best regularization setting to the learning rate for
SGD and breaks equivalence for adaptive gradient methods due to gradient
```

```latex
scaling. \cite{loshchilov2019adamw} The proposed fix is to apply weight decay
as a separate decay step (possibly schedule-scaled), rather than including
$\lambda\theta$ in the gradient that is subsequently adaptively scaled.
\cite{loshchilov2019adamw} The paper also describes a normalized weight decay
and restart-based variants (AdamWR), but details beyond what is explicitly
presented are \emph{not specified in the provided sources}.
\cite{loshchilov2019adamw}

\subsection{Second-order / preconditioning}
\paragraph{Shampoo.}
Shampoo frames updates as preconditioned steps, noting that when $W=\R^d$, an
online mirror descent update with preconditioner $H_t$ yields a preconditioned
gradient step $w_{t+1}=w_t-\eta H_t^{-1} g_t$. \cite{gupta2018shampoo} For
matrix parameters $W_t\in\R^{m\times n}$ with gradient $G_t$, Shampoo uses
left/right accumulators (e.g., $L_t=\epsilon I_m+\sum_{s\le t} G_s G_s^\top$
and $R_t=\epsilon I_n+\sum_{s\le t} G_s^\top G_s$) and updates of the form
\begin{equation}
W_{t+1} = W_t - \eta\, L_t^{-1/4} \, G_t \, R_t^{-1/4},
\end{equation}
with analysis connecting this to a single Kronecker preconditioner on
$\mathrm{vec}(W_t)$. \cite{gupta2018shampoo}

\subsection{Sign-based / communication-efficient methods}
\paragraph{SignSGD and distributed majority vote.}
SignSGD uses updates of the form $x_{k+1}=x_k-\eta\,\sign(\tilde g)$, where
$\tilde g$ is a stochastic gradient estimate, and motivates this as 1-bit
compression of gradient information. \cite{bernstein2018signsgd} In distributed
training, the paper proposes majority vote aggregation of per-worker sign
vectors, enabling 1-bit communication both worker$\rightarrow$server and
server$\rightarrow$worker. \cite{bernstein2018signsgd} The analysis provides
non-convex convergence rates for SignSGD (Theorem~1) and for distributed
majority vote (Theorem~2) under assumptions including (in some regimes)
unimodal and symmetric gradient noise. \cite{bernstein2018signsgd}

\subsection{Specialized scaling- or structure-oriented methods}
\paragraph{Muon (hidden-layer matrix optimizer).}
Muon is presented as an optimizer for hidden-layer weight matrices: it
transforms a matrix gradient $G$ into a semi-orthogonal update $\Ortho(G)$
using an iterative Newton--Schulz-based procedure, then applies a momentum-
style update. \cite{kellerjordan2024muon} The source provides a concrete
algorithmic definition of the Newton--Schulz iteration used (including the
```

```
fixed five-step ``NewtonSchulz5'' polynomial iteration) and describes how Muon
integrates with AdamW for other parameter types in practical training.
\cite{kellerjordan2024muon}

\paragraph{MuP (context, not an optimizer).}
MuP (Maximal Update Parametrization) aims to make learning-rate and related
hyperparameters transferable across width, contrasting with parameterizations
in which stable learning rates shrink with width. \cite{yang2022mup} The paper
discusses ``$\mu$P-compatible'' learning-rate choices and includes guidance
indicating that some techniques (e.g., weight decay in $\mu$P Adam) may not be
compatible as stated. \cite{yang2022mup} Since MuP does not define a new update
rule, we treat it as a scaling framework that conditions the behavior of
optimizers such as SGD and Adam rather than a distinct optimizer family.
\cite{yang2022mup}

\section{Comparative Analysis}
\subsection{Side-by-side update equations (consistent notation)}
We collect update rules in a unified notation. For optimizers whose full
details are not specified in the provided sources (e.g., exact SGD momentum
variant choices), we state only what is specified.
\cite{loshchilov2019adamw,bernstein2018signsgd}

\paragraph{SGD (baseline form).}
\begin{equation}
\theta_t = \theta_{t-1} - \alpha\, g_t,
\end{equation}
used as a baseline comparator throughout the sources.
\cite{kingma2015adam,bernstein2018signsgd,gupta2018shampoo}

\paragraph{Adam.}
\begin{equation}
\theta_t = \theta_{t-1} - \alpha \frac{\hat m_t}{\sqrt{\hat v_t}+\epsilon},
\end{equation}
with $(m_t,v_t,\hat m_t,\hat v_t)$ as defined earlier. \cite{kingma2015adam}

\paragraph{AdamW (decoupled decay form).}
In AdamW, weight decay is applied separately from the gradient-based adaptive
update; the key claim is that this decoupling removes an explicit coupling
between $\lambda$ and $\alpha$ (and avoids scaling the decay term by Adam's
adaptive preconditioner). \cite{loshchilov2019adamw} The exact line-by-line
AdamW pseudo-code and its schedule multiplier $\eta_t$ are given in the source.
```

```latex
\cite{loshchilov2019adamw}

\paragraph{SignSGD.}
\begin{equation}
x_{k+1} = x_k - \eta\, \sign(\tilde g_k),
\end{equation}
and in distributed majority vote:
\begin{equation}
x_{k+1} = x_k - \eta\, \sign\!\left(\sum_{m=1}^M \sign(\tilde
g^{(m)}_k)\right).
\end{equation}
\cite{bernstein2018signsgd}

\paragraph{Shampoo (matrix form).}
\begin{equation}
W_{t+1} = W_t - \eta\, L_t^{-1/4} G_t R_t^{-1/4}.
\end{equation}
\cite{gupta2018shampoo}

\paragraph{Muon (matrix form, hidden layers).}
Muon uses a momentum accumulator $V_t$ updated using the orthogonalized
gradient direction $\Ortho(G_t)$, where $\Ortho(\cdot)$ is computed by a fixed-
step Newton--Schulz iteration over a scaled gradient matrix; the precise
operator and iteration are defined in the source. \cite{kellerjordan2024muon} A
minimal schematic consistent with the source is:
\begin{align}
\tilde G_t &\leftarrow \Ortho(G_t), \\
V_t &\leftarrow \mu V_{t-1} + \tilde G_t, \\
W_t &\leftarrow W_{t-1} - \eta V_t,
\end{align}
with practical details (e.g., application to 2D parameters and integration with
AdamW elsewhere) as specified by the source. \cite{kellerjordan2024muon}

\subsection{Table 1: Mechanism-level comparison}
\begin{table}[h]
\centering
\small
\begin{tabular}{@{}p{2.1cm}p{3.2cm}p{4.2cm}p{2.1cm}p{2.1cm}@{}}
\toprule
Optimizer & State variables & Core mechanism & Extra compute & Extra memory \\
\midrule
```

```
SGD & none (baseline) & $-\alpha g_t$ step
\cite{kingma2015adam,bernstein2018signsgd} & low \cite{kingma2015adam} & low
\cite{kingma2015adam} \\
Adam & $m_t,v_t$ (+ bias correction) \cite{kingma2015adam} & adaptive per-
parameter step via $\hat m_t/(\sqrt{\hat v_t}+\epsilon)$; update magnitude
approx.\ bounded, invariant to gradient rescaling \cite{kingma2015adam} &
elementwise ops \cite{kingma2015adam} & two parameter-sized buffers
\cite{kingma2015adam} \\
AdamW & Adam-style state + separate decay \cite{loshchilov2019adamw} &
decoupled weight decay step; avoids scaling decay by adaptive gradients and
reduces explicit $\lambda$--$\alpha$ coupling \cite{loshchilov2019adamw} &
small (extra decay op) \cite{loshchilov2019adamw} & Adam-like
\cite{loshchilov2019adamw} \\
SignSGD & none required (core) \cite{bernstein2018signsgd} & 1-bit sign update;
distributed majority vote compresses both directions
\cite{bernstein2018signsgd} & sign + aggregation \cite{bernstein2018signsgd} &
low \cite{bernstein2018signsgd} \\
Shampoo & preconditioners (e.g., $L_t,R_t$) \cite{gupta2018shampoo} &
Kronecker-structured matrix/tensor preconditioning; equivalent to a single
preconditioner on $\mathrm{vec}$ form \cite{gupta2018shampoo} & matrix
roots/inverses \cite{gupta2018shampoo} & $m^2+n^2$ for $m\times n$ (implied by
storing $L_t,R_t$) \cite{gupta2018shampoo} \\
Muon & momentum $V_t$ + orthogonalization procedure \cite{kellerjordan2024muon}
& replace $G$ by $\Ortho(G)$ via fixed Newton--Schulz iteration; apply momentum
and update hidden-layer matrices \cite{kellerjordan2024muon} & iterative matrix
ops \cite{kellerjordan2024muon} & momentum buffer + temporaries
\cite{kellerjordan2024muon} \\
\bottomrule
\end{tabular}
\caption{Mechanism-level comparison grounded in algorithmic descriptions in the
provided sources.
\cite{kingma2015adam,loshchilov2019adamw,bernstein2018signsgd,gupta2018shampoo,
kellerjordan2024muon}}
\end{table}

\subsection{Table 2: Practical tradeoffs}
\begin{table}[h]
\centering
\small
\begin{tabular}{@{}p{2.1cm}p{3.5cm}p{3.8cm}p{4.0cm}@{}}
\toprule
```

```
Optimizer & Hyperparameter sensitivity & Stability / failure modes & Typical
regimes and cautions \\
\midrule
SGD & baseline; detailed sensitivity \emph{not specified}
\cite{bernstein2018signsgd} & detailed failure modes \emph{not specified}
\cite{kingma2015adam} & serves as comparator; SIGNUM suggests many settings can
be lifted from SGD when tuned \cite{bernstein2018signsgd} \\
Adam & authors report good defaults and robustness across tasks
\cite{kingma2015adam} & bias correction important when $\beta_2$ close to 1;
removing it can cause early instabilities \cite{kingma2015adam} & designed for
sparse gradients and non-stationary objectives; claims invariance to gradient
rescaling \cite{kingma2015adam} \\
AdamW & aims to decouple learning rate and decay; argues adaptive scaling makes
$L_2$ behave differently than true decay \cite{loshchilov2019adamw} & failure
mode: coupling $\lambda$ to $\alpha$ in $L_2$-as-decay and adaptive scaling of
regularizer gradients \cite{loshchilov2019adamw} & recommended when using
adaptive methods with weight decay-like regularization
\cite{loshchilov2019adamw} \\
SignSGD & fewer continuous hyperparameters; distributed majority vote adds
worker count effects \cite{bernstein2018signsgd} & theory depends on
assumptions tailored to signed updates and (for improved rate) unimodal
symmetric noise \cite{bernstein2018signsgd} & especially efficient in problems
with dense gradients/noise and $\ell_1$-geometry; can match SGD rates in
practice per experiments \cite{bernstein2018signsgd} \\
Shampoo & practical heuristics (e.g., projection omission) and computational
expense noted \cite{gupta2018shampoo} & projection step can be expensive and is
often omitted; regret analysis assumes conditions (e.g., rank constraints in
stated theorem) \cite{gupta2018shampoo} & targets tensor parameters; uses
Kronecker structure to approximate full preconditioning \cite{gupta2018shampoo}
\\
Muon & specialized; requires 2D parameters for orthogonalization operator
\cite{kellerjordan2024muon} & cautions: applies to hidden-layer matrices;
recommends AdamW for embeddings/scalars and other parameters
\cite{kellerjordan2024muon} & intended for hidden layers; implementation uses
fixed Newton--Schulz steps (``NewtonSchulz5'') \cite{kellerjordan2024muon} \\
\bottomrule
\end{tabular}
\caption{Practical tradeoffs and cautions explicitly discussed in the provided
sources; items marked ``not specified'' are omitted by the sources and are not
inferred here.
\cite{kingma2015adam,loshchilov2019adamw,bernstein2018signsgd,gupta2018shampoo,
```

```
kellerjordan2024muon}}
\end{table}

\subsection{Failure modes and ``performs poorly'' contexts}
\paragraph{Adam: bias correction and sparse-gradient settings.}
The Adam paper reports that removing bias correction yields a RMSProp-with-
momentum-like variant and observes instability for $\beta_2$ close to 1 in
early epochs without bias correction, especially when sparse-gradient
robustness motivates $\beta_2$ near 1. \cite{kingma2015adam} Beyond this,
contexts where Adam performs poorly are \emph{not specified in the provided
sources}.

\paragraph{AdamW: regularization coupling in adaptive methods.}
AdamW emphasizes that the common equivalence between $L_2$ regularization and
weight decay does not hold for adaptive gradient methods, because the
regularizer gradient is scaled; consequently, some parameters may be
effectively regularized less under Adam+$L_2$ than under decoupled decay.
\cite{loshchilov2019adamw} Additional failure modes are \emph{not specified in
the provided sources}.

\paragraph{SignSGD: geometry and noise assumptions.}
SignSGD's theoretical rates rely on assumptions adapted to signed updates and
discuss regimes where dense gradients/noise and $\ell_1$-geometry yield
favorable dimension dependence; improved distributed rates additionally use
unimodal symmetric noise assumptions. \cite{bernstein2018signsgd} Performance
outside these regimes is \emph{not specified in the provided sources}.

\paragraph{Shampoo: computational and projection considerations.}
Shampoo notes that projections required to keep iterates within bounded domains
can be computationally expensive and are rarely performed in practice (and thus
omitted in the algorithm for a looser bound). \cite{gupta2018shampoo} The
computational cost of preconditioner operations is an inherent practical
consideration. \cite{gupta2018shampoo}

\paragraph{Muon: scope restrictions.}
Muon is explicitly targeted to 2D weight matrices (hidden layers) and
recommends using AdamW for embeddings/scalars and other parameter types,
indicating scope limitations as a practical ``performs poorly'' or ``not
applicable'' context. \cite{kellerjordan2024muon}

\paragraph{MuP context: compatibility caveats.}
```

```latex
MuP includes explicit guidance that certain practices (e.g., weight decay in
$\mu$P Adam) are not compatible as stated, which can be interpreted as a
failure mode if naively transferred. \cite{yang2022mup}

\section{Research Insights \& Open Questions}
\subsection{Synthesized insights (grounded in provided sources)}
\begin{enumerate}
\item Bias correction is not a cosmetic detail: when decay rates (especially
$\beta_2$) are close to 1, initialization bias can materially destabilize
training without correction in early epochs. \cite{kingma2015adam}
\item Decoupling weight decay from adaptive gradient scaling is conceptually
important because adaptive methods scale gradients by historical magnitudes,
which can unintentionally scale regularizer gradients and alter the intended
regularization strength across parameters. \cite{loshchilov2019adamw}
\item Signed updates induce a natural $\ell_1$-geometry, and analyzing SignSGD
requires assumptions finer-grained than standard SGD analyses; under these
assumptions, SignSGD achieves a non-convex rate with favorable dimension
dependence in dense-gradient regimes. \cite{bernstein2018signsgd}
\item Majority-vote aggregation enables 1-bit communication in both directions
and can match the theoretical speedup of distributed SGD under assumptions
validated empirically (including approximate symmetry/unimodality of noise in
deep learning gradients). \cite{bernstein2018signsgd}
\item Shampoo shows how tensor structure can yield preconditioning that is
equivalent (on vectorized parameters) to applying a Kronecker product
preconditioner, enabling regret guarantees while avoiding full dense
preconditioning. \cite{gupta2018shampoo}
\item Shampoo's regret bounds can achieve the optimal $\mathcal{O}(\sqrt{T})$
scaling under stated conditions (e.g., rank constraints), connecting structured
preconditioning to best-possible stochastic/online rates.
\cite{gupta2018shampoo}
\item Muon’s design highlights a different kind of structure: for weight
matrices, constraining updates toward semi-orthogonality (via Newton--Schulz
iterations) is proposed as a useful inductive bias for hidden-layer
optimization. \cite{kellerjordan2024muon}
\item Hybridization is an explicit design pattern in the provided sources:
AdamW proposes mixing adaptive updates with decoupled decay; Muon proposes
using Muon on some parameter subsets and AdamW elsewhere.
\cite{loshchilov2019adamw,kellerjordan2024muon}
\item Scaling behavior can hinge on parameterization as much as on the
optimizer: MuP aims to preserve update magnitudes across width so that learning
rates remain stable and transferable. \cite{yang2022mup}
```

```latex
\item Some ``best practices'' do not transfer unchanged across scaling
frameworks: MuP explicitly warns about incompatibility of certain
regularization choices (e.g., weight decay in $\mu$P Adam) in its setting.
\cite{yang2022mup}
\end{enumerate}

\subsection{Open questions (motivated directly by provided sources)}
\begin{enumerate}
\item Can the decoupled-regularization viewpoint of AdamW be generalized into a
principled taxonomy of which terms should (or should not) be passed through
adaptive preconditioners? \cite{loshchilov2019adamw}
\item SignSGD’s improved distributed theory uses unimodal symmetric noise
assumptions; how robust are these guarantees under heavier-tailed or multimodal
gradient noise observed outside the examined regimes?
\cite{bernstein2018signsgd}
\item What are the sharp boundaries of the ``dense gradients/noise'' regime in
which SignSGD enjoys favorable dimension dependence, and how can practitioners
diagnose this regime online? \cite{bernstein2018signsgd}
\item Shampoo uses Kronecker-structured accumulators; what structure classes
beyond Kronecker products yield similar guarantees with manageable compute for
modern architectures? \cite{gupta2018shampoo}
\item Shampoo’s analysis includes conditions such as rank constraints in a main
theorem; when and why do these conditions approximately hold in deep networks,
and can they be enforced or encouraged? \cite{gupta2018shampoo}
\item Muon proposes a fixed-step Newton--Schulz procedure (``NewtonSchulz5'');
how sensitive is performance to the polynomial/iteration choice, and can
stability be characterized theoretically? \cite{kellerjordan2024muon}
\item Muon is targeted to hidden-layer matrices and recommends AdamW for other
parameters; can one unify such heterogeneous parameter-type rules into a single
optimizer with analyzable behavior? \cite{kellerjordan2024muon}
\item MuP warns about incompatibilities such as weight decay in $\mu$P Adam;
what are the minimal modifications to recover regularization effects while
preserving MuP’s hyperparameter transfer properties? \cite{yang2022mup}
\item Across these methods, how should one compare ``compute budgets'' fairly
when some methods trade extra per-step compute (Shampoo, Muon) for fewer steps
or different stability behavior? \cite{gupta2018shampoo,kellerjordan2024muon}
\item The provided sources emphasize different geometries (elementwise
adaptivity, $\ell_1$ sign geometry, Kronecker preconditioning, orthogonalized
directions); is there a unifying geometric framework that predicts when each
geometry is beneficial?
\cite{kingma2015adam,bernstein2018signsgd,gupta2018shampoo,kellerjordan2024muon
```

```latex
}
\end{enumerate}

\section{Conclusion}
The provided sources illustrate a coherent optimizer lineage in which each
method targets a specific bottleneck: Adam focuses on adaptive scaling via
moment estimates and bias correction with bounded, rescaling-invariant updates.
\cite{kingma2015adam} AdamW reframes regularization as an algorithmic design
choice, arguing that decoupling weight decay from adaptive scaling reduces
implicit hyperparameter coupling and changes the effective regularization
behavior of adaptive methods. \cite{loshchilov2019adamw} SignSGD shifts
attention to communication and robustness under extreme quantization, providing
non-convex convergence theory under assumptions suited to signed updates and
offering a majority-vote protocol for 1-bit distributed training.
\cite{bernstein2018signsgd} Shampoo pushes toward structured second-order
information via Kronecker preconditioning with regret guarantees, while Muon
explores an alternative structure for hidden-layer matrices by orthogonalizing
update directions via Newton--Schulz iterations and pairing this with momentum.
\cite{gupta2018shampoo,kellerjordan2024muon}

Finally, MuP underscores that optimizer behavior is intertwined with
parameterization and scaling: stable learning-rate transfer across width may
require changing the model’s parametrization rather than the optimizer rule
itself. \cite{yang2022mup} Taken together, these works suggest that ``optimizer
design'' is best viewed as designing \emph{update geometry} under practical
constraints---compute, memory, communication, and scaling---with careful
attention to which algorithmic details (e.g., bias correction, decoupled decay,
or parameterization) fundamentally change the training dynamics.
\cite{kingma2015adam,loshchilov2019adamw,bernstein2018signsgd,gupta2018shampoo,
yang2022mup}

\appendix
\section{Update-rule cheat sheet}
\begin{table}[h]
\centering
\small
\begin{tabular}{@{}p{2.2cm}p{10.8cm}@{}}
\toprule
Optimizer & Cheat-sheet update \\
\midrule
SGD & $\theta_t \leftarrow \theta_{t-1} - \alpha g_t$ (baseline form)
```

```
\cite{kingma2015adam,bernstein2018signsgd} \\
Adam & $m_t\!\leftarrow\!\beta_1 m_{t-1}\!+\!(1-\beta_1)g_t,\;
v_t\!\leftarrow\!\beta_2 v_{t-1}\!+\!(1-\beta_2)g_t^2$,\\
& $\hat m_t\!\leftarrow\!m_t/(1-\beta_1^t),\; \hat v_t\!\leftarrow\!v_t/(1-
\beta_2^t)$,\\
& $\theta_t \leftarrow \theta_{t-1} - \alpha \hat m_t/(\sqrt{\hat
v_t}+\epsilon)$ \cite{kingma2015adam} \\
AdamW & Adam-style adaptive step plus separate weight decay step (decoupled
from adaptive scaling), with schedule multiplier $\eta_t$ in pseudo-code.
\cite{loshchilov2019adamw} \\
SignSGD & $x_{k+1} \leftarrow x_k - \eta \sign(\tilde g_k)$; distributed:
$x_{k+1} \leftarrow x_k - \eta \sign(\sum_m \sign(\tilde g^{(m)}_k))$
\cite{bernstein2018signsgd} \\
Shampoo & $W_{t+1} \leftarrow W_t - \eta\, L_t^{-1/4} G_t R_t^{-1/4}$ with
$L_t=\epsilon I_m + \sum_{s\le t} G_s G_s^\top$, $R_t=\epsilon I_n + \sum_{s\le
t} G_s^\top G_s$ \cite{gupta2018shampoo} \\
Muon & For 2D matrices: $\tilde G\leftarrow \Ortho(G)$ via fixed Newton--Schulz
steps; momentum update $V\leftarrow \mu V + \tilde G$; $W\leftarrow W-\eta V$
(integration details in source). \cite{kellerjordan2024muon} \\
\bottomrule
\end{tabular}
\caption{Update-rule cheat sheet. Items are stated only to the extent specified
in the provided sources.
\cite{kingma2015adam,loshchilov2019adamw,bernstein2018signsgd,gupta2018shampoo,
kellerjordan2024muon}}
\end{table}

\section{Glossary of terms}
\begin{description}
\item[Bias correction] Adjustment of moment estimates initialized at zero to
reduce early-iteration bias, as used in Adam via division by $(1-\beta^t)$.
\cite{kingma2015adam}
\item[Decoupled weight decay] Applying a multiplicative parameter decay step
separately from the gradient-based update, advocated to avoid adaptive scaling
of the regularizer gradient. \cite{loshchilov2019adamw}
\item[Majority vote (SignSGD)] Distributed aggregation of worker sign vectors
by taking the sign of the sum of signs, enabling 1-bit communication in both
directions. \cite{bernstein2018signsgd}
\item[Preconditioning] Multiplying gradients by (an approximation to) an
inverse curvature-like matrix to change the update geometry; Shampoo uses
structured (Kronecker) preconditioners. \cite{gupta2018shampoo}
```

```
\item[Kronecker structure] Representing a large matrix as a Kronecker product
of smaller matrices; Shampoo connects left/right matrix preconditioners to a
Kronecker product preconditioner on vectorized parameters.
\cite{gupta2018shampoo}
\item[Orthogonalized update (Muon)] Transforming a matrix gradient into a semi-
orthogonal matrix $\Ortho(G)$ using an iterative Newton--Schulz procedure
before applying a momentum update. \cite{kellerjordan2024muon}
\item[MuP] A parameterization intended to stabilize and transfer
hyperparameters across width scaling; not itself an optimizer update rule.
\cite{yang2022mup}
\end{description}

\begin{thebibliography}{9}

\bibitem{kingma2015adam}
Diederik P. Kingma and Jimmy Ba.
\newblock Adam: A Method for Stochastic Optimization.
\newblock Published as a conference paper at ICLR 2015. \cite{kingma2015adam}

\bibitem{loshchilov2019adamw}
Ilya Loshchilov and Frank Hutter.
\newblock Decoupled Weight Decay Regularization.
\newblock Published as a conference paper at ICLR 2019.
\cite{loshchilov2019adamw}

\bibitem{bernstein2018signsgd}
Jeremy Bernstein, Yu-Xiang Wang, Kamyar Azizzadenesheli, and Anima Anandkumar.
\newblock signSGD: Compressed Optimisation for Non-Convex Problems.
\newblock ICML 2018 (as provided). \cite{bernstein2018signsgd}

\bibitem{gupta2018shampoo}
Vineet Gupta, Timon Gasthaus, and others (as listed in the provided source).
\newblock Shampoo: Preconditioned Stochastic Tensor Optimization.
\newblock ICML 2018 (as provided). \cite{gupta2018shampoo}

\bibitem{kellerjordan2024muon}
Keller Jordan and collaborators (as listed in the provided source).
\newblock Muon: An optimizer for hidden layers in neural networks.
\newblock Blog PDF as provided (2024). \cite{kellerjordan2024muon}

\bibitem{yang2022mup}
```

```
Greg Yang and collaborators (as listed in the provided source).
\newblock Tensor Programs V: Tuning Large Neural Networks via Zero-Shot
Hyperparameter Transfer (MuP / Maximal Update Parametrization).
\newblock As provided in \texttt{MuP.pdf}. \cite{yang2022mup}

\end{thebibliography}

\end{document}
```

**Lecture 2.pdf**
PDF

**Lecture 3.pdf**
PDF

**Lecture 4.pdf**
PDF

**Lecture 5.pdf**
PDF

**Lecture 6.pdf**
PDF

**Lecture 7.pdf**
PDF

This is perfect! This time, can you refine this version, incorporating the last input and file uploads, but also taking into account these lecture notes?

Thought for 10m 53s >

\documentclass[11pt]{article}

\usepackage[utf8]{inputenc}
\usepackage[T1]{fontenc}

\usepackage{amsmath,amssymb}
\usepackage{graphicx}
\usepackage{booktabs}
\usepackage{hyperref}
\usepackage{geometry}
\usepackage{xcolor}
\usepackage{tikz}

\geometry{margin=1in}
\hypersetup{
colorlinks=true,
linkcolor=blue,
citecolor=blue,
urlcolor=blue
}

\title{A Lineage-Driven Survey of Landmark Optimizers:\SGD, Adam, AdamW, SignSGD, Shampoo, Muon (with MuP Context)}
\author{}
\date{}

\newcommand{\R}{\mathbb{R}}
\newcommand{\E}{\mathbb{E}}
\newcommand{\diag}{\mathrm{diag}}
\newcommand{\sign}{\mathrm{sign}}
\newcommand{\Tr}{\mathrm{Tr}}
\newcommand{\vect}{\mathrm{vec}}
\newcommand{\Ortho}{\mathrm{Ortho}}

\begin{document}
\maketitle
\tableofcontents

\begin{abstract}
Optimizer design sits at the intersection of numerical methods, statistical estimation, and systems constraints: the update rule determines not only convergence behavior but also memory footprint, communication cost, and how hyperparameters transfer

across scales. This survey traces a concrete lineage across six landmark optimizers---SGD, Adam, AdamW, SignSGD, Shampoo, and Muon---and situates MuP as a scaling framework that reframes what it means for hyperparameters to ``transfer'' with width. Starting from SGD and momentum-based smoothing (\cite{lec2,adamw}), we follow two major branches: (i) adaptive per-coordinate steps culminating in Adam and its decoupled-weight-decay variant AdamW (\cite{adam,adamw}); and (ii) curvature-aware preconditioning culminating in Shampoo (\cite{shampoo}), with Muon interpreting efficient orthogonalization as a principled approximation related to Shampoo-style geometry (\cite{muonblog}). In parallel, SignSGD compresses optimization and communication to 1-bit signs while retaining nonconvex convergence guarantees under explicit noise assumptions (\cite{signsgd}). Finally, MuP changes the parameterization to stabilize optimal hyperparameters across widths and motivates a scaling-aware lens on optimizer behavior (\cite{mup}). Throughout, when details are absent from the provided sources, we state so explicitly.
\end{abstract}

\section{Introduction}
Optimizer choice is a first-order design decision in deep learning because it affects training stability, achievable scale, and the sensitivity of results to hyperparameter tuning (\cite{mup,adam,adamw}). In practice, ``better'' optimizers are rarely better along all axes: improving one dimension (e.g., communication efficiency) can worsen another (e.g., bias or robustness), and improving geometry adaptation can increase compute/memory overhead (\cite{signsgd,shampoo}).

We organize comparisons along the axes emphasized in the provided sources and lecture notes:
\begin{itemize}
\item \textbf{Compute:} per-step FLOPs beyond backprop; matrix operations vs elementwise operations (\cite{shampoo,muonblog,lec6}).
\item \textbf{Memory:} optimizer state (e.g., momentum vectors, second-moment accumulators, matrix preconditioners) (\cite{adam,shampoo,muonblog}).
\item \textbf{Stability:} sensitivity to initialization bias or numerical issues; failure modes early in training (\cite{adam,muonblog}).
\item \textbf{Hyperparameter sensitivity:} coupling between learning rate and regularization; stability of optima across scales (\cite{adamw,mup}).
\item \textbf{Generalization / implicit bias:} regularization interactions such as weight decay vs L2 penalty and their optimizer-dependent effects (\cite{adamw,lec2}).

\item \textbf{Scaling behavior:} whether ``good'' hyperparameters persist with model width; parameterization-dependent transfer (\cite{mup,lec6}).
\end{itemize}

\subsection{Notation (used throughout)}

\begin{table}[h]

\centering

\begin{tabular}{@{}ll@{}}

\toprule

Symbol & Meaning \

\midrule

$t$ & iteration index \

$\theta_t \in \R^d$ & parameter vector at step $t$ \

$g_t$ & stochastic gradient (e.g., $g_t=\nabla_\theta f_t(\theta_{t-1})$) (\cite{adam}) \

$\alpha$ & base stepsize / learning rate (\cite{adam}) \

$m_t$ & first-moment / momentum state (\cite{adam,lec2}) \

$v_t$ & second raw-moment state (\cite{adam}) \

$\beta_1,\beta_2$ & exponential decay rates for moments (\cite{adam}) \

$\varepsilon$ & numerical stability constant (\cite{adam}) \

$W_t \in \R^{n\times m}$ & matrix parameter (used for Shampoo/Muon) (\cite{shampoo,muonblog}) \

\bottomrule

\end{tabular}

\caption{Notation table.}

\end{table}

\section{Lineage Map (Core Section)}

\subsection{Textual lineage overview}
We summarize each optimizer by (i) core idea, (ii) what it borrows, (iii) what it changes/adds, and (iv) why it mattered historically \emph{as evidenced in the provided sources}.

\paragraph{SGD.}
\textbf{Core idea:} update parameters using stochastic gradient information. \textbf{Borrowed:} gradient-based descent. \textbf{Changed/added:} stochasticity from minibatches; implicit regularization perspectives are discussed in the lecture notes (\cite{lec2}). \textbf{Why it mattered:} SGD is a baseline against which

adaptive and preconditioned methods are compared across the sources (\cite{adam,signsgd,shampoo}).

\paragraph{(Polyak) Momentum.}
\textbf{Core idea:} smooth updates via an exponentially weighted moving average of gradients. \textbf{Borrowed:} SGD step structure. \textbf{Changed/added:} introduces a velocity-like state $v_t$ and updates $\theta$ using that state (\cite{lec2}). \textbf{Why it mattered:} momentum appears repeatedly as a building block (e.g., Adam's first moment; SignSGD's SIGNUM variant) (\cite{adam,signsgd,lec2}).

\paragraph{Adam.}
\textbf{Core idea:} adaptive per-parameter learning rates computed from estimates of first and second moments of gradients (\cite{adam}). \textbf{Borrowed:} momentum-style averaging (first moment) and RMS-style normalization (second moment) (\cite{adam}). \textbf{Changed/added:} bias-correction for moment estimates and an update invariant to rescaling of the gradient (as claimed in the paper's abstract/intro) (\cite{adam}). \textbf{Why it mattered:} widely used default for high-dimensional problems; the paper emphasizes memory efficiency and robustness properties (\cite{adam}).

\paragraph{AdamW.}
\textbf{Core idea:} decouple weight decay from gradient-based adaptive updates (\cite{adamw}). \textbf{Borrowed:} Adam-style adaptive updates and momentum; classical weight decay idea (\cite{adamw}). \textbf{Changed/added:} shows that the common equivalence between L2 regularization and weight decay does \emph{not} hold for adaptive methods, motivating an explicit decoupled decay step (\cite{adamw}). \textbf{Why it mattered:} reframes regularization for adaptive optimizers and aims to decouple hyperparameters (learning rate vs decay) (\cite{adamw}).

\paragraph{SignSGD (and majority vote).}
\textbf{Core idea:} update using only the elementwise sign of stochastic gradients; in distributed settings, aggregate signs via majority vote so that communication can be 1-bit in both directions (\cite{signsgd}). \textbf{Borrowed:} gradient descent directionality. \textbf{Changed/added:} explicit compression and theory under noise-shape assumptions (e.g., unimodal and symmetric noise) (\cite{signsgd}). \textbf{Why it mattered:} targets communication bottlenecks while retaining convergence guarantees under stated conditions (\cite{signsgd,lec5}).

\paragraph{Shampoo.}
\textbf{Core idea:} tensor-aware preconditioning using left and right second-moment accumulators and inverse fractional matrix powers (\cite{shampoo}).
\textbf{Borrowed:} second-moment preconditioning ideas (general preconditioned updates) (\cite{shampoo}). \textbf{Changed/added:} Kronecker-structured approximation enabling matrix preconditioning without full $d\times d$ storage; update can be written $W_{t+1}=W_t-\eta L_t^{-1/4}G_t R_t^{-1/4}$ for matrix/tensor parameters (as referenced in the Shampoo analysis) (\cite{shampoo}).
\textbf{Why it mattered:} bridges diagonal adaptive methods and full-matrix second-order approximations with a structured, scalable design (\cite{shampoo}).

\paragraph{Muon.}
\textbf{Core idea:} apply an efficient orthogonalization operation to the (momentum) update matrix for 2D parameters using Newton--Schulz iterations, approximating the nearest (semi-)orthogonal matrix (\cite{muonblog}). \textbf{Borrowed:} SGD-momentum baseline update; numerical linear algebra tools for orthogonalization (Newton--Schulz) (\cite{muonblog}). \textbf{Changed/added:} replaces the raw update matrix with an orthogonalized surrogate, argued to amplify ``rare directions'' when updates are ill-conditioned (\cite{muonblog}). \textbf{Why it mattered:} proposes a geometry-motivated modification with claimed low overhead in typical LM scenarios and an explicit relationship to Shampoo-style geometry when accumulation is removed (\cite{muonblog,lec6}).

\paragraph{MuP (context, not an optimizer).}
\textbf{Core idea:} change parameterization so that many optimal hyperparameters remain stable across model widths, enabling ``$\mu$Transfer'' (zero-shot transfer of tuned hyperparameters from small to large models) (\cite{mup}). \textbf{Borrowed:} width-scaling analysis and infinite-width intuitions; explicitly discusses how parameterization changes learning dynamics (\cite{mup}). \textbf{Changed/added:} proposes Maximal Update Parametrization and demonstrates stable optimal learning rates across widths (including for Adam in experiments) (\cite{mup}). \textbf{Why it mattered:} reframes optimizer comparisons as parameterization-dependent and scaling-aware, affecting how we interpret hyperparameter sensitivity (\cite{mup,lec6}).

\subsection{Optional TikZ lineage diagram}
\begin{figure}[h]
\centering

```
\begin{tikzpicture}[
node/.style={draw, rounded corners, align=center, inner sep=6pt},
arrow/.style={->, thick}
]
\node[node] (sgd) {SGD\\cite{lec2,adam,signsgd}};
\node[node, right=2.6cm of sgd] (mom) {Momentum\\cite{lec2}};
\node[node, right=2.6cm of mom] (adam) {Adam\(1st+2nd moments)\\cite{adam}};
\node[node, above=1.8cm of adam] (adamw) {AdamW\(decoupled
WD)\\cite{adamw}};
\node[node, below=1.8cm of adam] (sign) {SignSGD\(majority vote)\\cite{signsgd}};
\node[node, right=3.2cm of adam] (shampoo) {Shampoo\(matrix
precond.)\\cite{shampoo}};
\node[node, above=1.8cm of shampoo] (muon) {Muon\(orthogonalize
update)\\cite{muonblog}};
\node[node, below=1.8cm of shampoo] (mup) {MuP (context)\
(scaling/transfer)\\cite{mup}};
\draw[arrow] (sgd) -- (mom);
\draw[arrow] (mom) -- (adam);
\draw[arrow] (adam) -- (adamw);
\draw[arrow] (adam) -- (sign);
\draw[arrow] (adam) -- (mup);
\draw[arrow] (sgd) -- (shampoo);
\draw[arrow] (shampoo) -- (muon);
\draw[arrow] (shampoo) -- (mup);
\end{tikzpicture}
\caption{High-level lineage map. Edges indicate conceptual inheritance as described
in the provided sources (e.g., Adam combines moment ideas; AdamW modifies
regularization handling; Muon relates to Shampoo-style geometry via
orthogonalization) (\cite{adam,adamw,shampoo,muonblog,mup}).}
\end{figure}
\section{Optimizer Families}
\subsection{SGD & Momentum}
\paragraph{SGD update.}
A standard SGD step uses the stochastic gradient direction:
\begin{equation}
```

\theta_t = \theta_{t-1} - \alpha, g_t,

\end{equation}

with $g_t$ defined as a minibatch gradient in the provided Adam pseudo-code and surrounding discussion (\cite{adam}). The lecture notes emphasize SGD's role as a baseline and discuss implicit regularization perspectives (\cite{lec2}).

\paragraph{Momentum update (as in lecture notes).}

The lecture notes present momentum as an exponential smoothing of gradients and update via the smoothed direction (\cite{lec2}):

\begin{align}

v_t &= \beta, v_{t-1} + (1-\beta), g_t, \

\theta_t &= \theta_{t-1} - \alpha, v_t.

\end{align}

\paragraph{Nesterov momentum.}

Not specified in the provided sources.

\subsection{Adaptive Methods: Adam and AdamW}

\paragraph{Adam.}

Adam maintains exponential moving averages of (i) gradients and (ii) squared gradients, then applies bias correction and an elementwise normalization (\cite{adam}):

\begin{align}

m_t &= \beta_1 m_{t-1} + (1-\beta_1) g_t, \

v_t &= \beta_2 v_{t-1} + (1-\beta_2) (g_t \odot g_t), \

\hat m_t &= \frac{m_t}{1-\beta_1^t},\qquad

\hat v_t = \frac{v_t}{1-\beta_2^t}, \

\theta_t &= \theta_{t-1} - \alpha \frac{\hat m_t}{\sqrt{\hat v_t}+\varepsilon}.

\end{align}

The Adam paper motivates the algorithm as combining advantages of AdaGrad and RMSProp, and empirically evaluates the importance of bias correction, noting instability when bias correction is removed for $\beta_2$ close to $1$ in early training (\cite{adam}).

\paragraph{AdamW (decoupled weight decay).}

AdamW argues that, unlike for SGD, the equivalence between L2 regularization and weight decay does not hold for adaptive gradient methods and that L2 coupling to learning rate is often overlooked even for SGD, motivating decoupled weight decay updates (\cite{adamw}). In the AdamW framing, weight decay is applied as an explicit decay step rather than folded into the gradient that is then adaptively normalized (\cite{adamw}).

\subsection{Second-Order / Preconditioning: Shampoo}

\paragraph{Shampoo’s structured preconditioning.}
Shampoo uses left and right preconditioners derived from accumulated gradient outer-products to precondition matrix/tensor gradients with inverse fractional powers (\cite{shampoo}). In the Shampoo analysis, the update for a matrix parameter can be written in the form
\begin{equation}
W_{t+1} = W_t - \eta, L_t^{-1/4}, G_t, R_t^{-1/4},
\end{equation}
and equivalently as a vectorized preconditioned update $w_{t+1}=w_t-\eta H_t^{-1} g_t$ for $w_t=\vect(W_t)$ with a corresponding structured preconditioner $H_t$ (\cite{shampoo}).

\subsection{Sign-Based / Communication-Efficient Methods: SignSGD}

\paragraph{SignSGD and majority vote.}
SignSGD updates parameters using only the elementwise sign of a stochastic gradient estimate (\cite{signsgd}). For distributed training with $M$ workers, the paper contrasts (i) sending each worker’s sign and summing (which breaks 1-bit communication on the return path) versus (ii) \emph{majority vote}, where the server takes a sign of the sum of worker signs and broadcasts a 1-bit decision back to all workers (\cite{signsgd}). The paper provides nonconvex convergence statements for majority vote under assumptions including unimodal and symmetric gradient noise, and describes a variance-reduction regime scaling with $\sqrt{M}$ in an appropriate sense (\cite{signsgd}). The lecture notes also contextualize SignSGD under communication and systems constraints (\cite{lec5}).

\subsection{Muon: Orthogonalized Updates for Hidden Layers}

\paragraph{Orthogonalization objective.}
Muon proposes to approximately orthogonalize an update matrix via a Newton--Schulz (NS) iteration and frames it as finding the nearest semi-orthogonal matrix in Frobenius norm (\cite{muonblog}):
\begin{equation}
\Ortho(G) ;=; \arg\min_O {|O-G|_F : O^\top O = I\ \text{or}\ OO^\top=I}.
\end{equation}
The blog notes that the NS iteration effectively replaces the SGD-momentum update

matrix with its nearest semi-orthogonal matrix and relates this to the SVD-based replacement $USV^\top \mapsto UV^\top$ (while rejecting explicit SVD as too slow) (\cite{muonblog}).

\paragraph{Compute/memory discussion.}
The Muon source states that before the NS iteration, Muon is standard SGD-momentum and thus has the same memory requirement, and it provides an overhead analysis indicating sub-1% FLOP overhead in specific LM training scenarios under typical parameter/batch scalings and a small number of NS steps (e.g., $T=5$) (\cite{muonblog}). The lecture notes summarize similar runtime comparisons and present wall-clock measurements for multiple optimizers in an LM training setting (\cite{lec6}).

\paragraph{Relationship to Shampoo.}
Muon explicitly discusses a relationship to Shampoo: it quotes a Shampoo definition and then notes an observation that removing preconditioner accumulation yields an ``orthogonalized gradient'' form, motivating Muon as an efficient route to that operation (\cite{muonblog}).

\subsection{Scaling-Oriented Context: MuP (not an optimizer)}
\paragraph{Maximal Update Parametrization and $\mu$Transfer.}
MuP claims that in the Maximal Update Parametrization, many optimal hyperparameters remain stable as model size changes, enabling the proposed $\mu$Transfer paradigm (tune on a small model and transfer to a larger one) (\cite{mup}). The paper’s introduction and figures emphasize that, in contrast to standard practice, optimal learning rates can remain stable across widths for Transformers trained with Adam under MuP (\cite{mup}). The paper also discusses how standard parameterizations can lead to width-dependent blow-ups after a single SGD step in an illustrative example, and how MuP modifies scaling to avoid such behavior (\cite{mup}).

\section{Comparative Analysis}
\subsection{Side-by-side update equations (consistent notation)}
We present a compact ``cheat sheet'' view; details and variants are as specified in the cited sources.

\paragraph{SGD / Momentum / Adam / AdamW.}
\begin{align}
\text{SGD:}\quad & \theta_t = \theta_{t-1} - \alpha g_t \qquad (\text{baseline in multiple sources}) ; \cite{adam} \

```
\text{Momentum:}\quad & v_t = \beta v_{t-1} + (1-\beta) g_t,;; \theta_t=\theta_{t-
1}-\alpha v_t ;; \cite{lec2} \
\text{Adam:}\quad & \theta_t = \theta_{t-1} - \alpha \frac{\hat m_t}{\sqrt{\hat
v_t}+\varepsilon} ;; \cite{adam} \
\text{AdamW:}\quad & \text{decouple WD from adaptive gradient scaling (conceptual
change)} ;; \cite{adamw}
\end{align}
\paragraph{SignSGD.}
\begin{equation}
\theta_t = \theta_{t-1} - \alpha, \sign(\tilde g_t)
\qquad \text{(with distributed majority vote variants)} ; \cite{signsgd}
\end{equation}
\paragraph{Shampoo and Muon (matrix parameters).}
\begin{align}
\text{Shampoo:}\quad & W_{t+1} = W_t - \eta, L_t^{-1/4} G_t R_t^{-1/4} ;;
\cite{shampoo} \
\text{Muon:}\quad & \Delta W_t \leftarrow \text{(SGD-momentum update matrix)};;;
\Delta W_t \leftarrow \Ortho(\Delta W_t);; \cite{muonblog}
\end{align}
\subsection{Table 1: Mechanism-level comparison}
\begin{table}[h]
\centering
\small
\resizebox{\textwidth}{!}{
\begin{tabular}{@{}lllll@{}}
\toprule
Optimizer & State variables & Update form & Compute (beyond backprop) & Memory
overhead \
\midrule
SGD & none & $\theta\leftarrow\theta-\alpha g$ & elementwise & none \
Momentum & $v$ & $\theta\leftarrow\theta-\alpha v$ & elementwise & $O(d)$
(\cite{lec2}) \
Adam & $m,v$ & normalized by $\sqrt{\hat v}$ & elementwise ops & $O(d)$
(\cite{adam}) \
AdamW & $m,v$ (+decay step) & decoupled WD step & elementwise + decay &
```

$O(d)$ (\cite{adamw}) \
SignSGD & none (basic) & sign-only direction & elementwise sign & minimal; comm-efficient (\cite{signsgd}) \
Shampoo & preconditioners & $L^{-1/4}GR^{-1/4}$ & matrix roots/mults & matrix/tensor states (\cite{shampoo}) \
Muon & momentum + NS iters & $\Ortho(\Delta W)$ & repeated matmuls & same as SGD-momentum + small overhead (\cite{muonblog}) \
\bottomrule
\end{tabular}
}
\caption{Mechanism-level comparison (as described in the provided sources). For precise asymptotics and constants, see the cited originals; some entries are qualitative when the source is qualitative.}
\end{table}

\subsection{Table 2: Practical tradeoffs and failure modes}
\begin{table}[h]
\centering
\small
\resizebox{\textwidth}{!}{
\begin{tabular}{@{}p{2.2cm}p{3.4cm}p{3.4cm}p{4.6cm}@{}}
\toprule
Optimizer & Hyperparameter sensitivity & Stability notes & Contexts where it performs poorly / caveats \
\midrule
SGD & LR sensitive; implicit reg discussed & baseline stability depends on LR schedule (not detailed here) & Not specified in the provided sources. \
Momentum & adds $\beta$; smoothing can help noise & lecture notes emphasize smoothing intuition (\cite{lec2}) & Not specified in the provided sources. \
Adam & $\beta_2$ near 1 increases init bias; bias correction important & removing bias correction can induce early instabilities for large $\beta_2$ (\cite{adam}) & Not specified in the provided sources. \
AdamW & targets LR--WD coupling; argues decoupling is important & discusses nonequivalence of L2 and WD for adaptive methods (\cite{adamw}) & Not specified in the provided sources. \
SignSGD & depends on noise-shape assumptions for theory & majority vote

convergence under unimodal/symmetric noise assumptions (\cite{signsgd}) & If assumptions fail, guarantees do not apply; further caveats not specified beyond stated assumptions (\cite{signsgd}). \
Shampoo & additional hyperparameters/engineering for preconditioners not summarized here & matrix root computations can be numerically delicate (general concern; specifics not detailed here) & Increased compute/memory; precise regimes not specified in the provided sources excerpted here (\cite{shampoo}). \
Muon & introduces NS iteration coefficients/steps & claims NS can run stably in bfloat16; alternatives may require float32 (\cite{muonblog}) & Applicability and best practices across architectures not specified in the provided sources. \
\bottomrule
\end{tabular}
}
\caption{Practical tradeoffs and failure modes \emph{as grounded in the provided sources}. When the sources do not specify a regime, we explicitly say so.}
\end{table}

\subsection{Discussion of failure modes (grounded)}
\paragraph{Initialization bias in adaptive moments.}
The Adam paper emphasizes that when $\beta_2$ is close to $1$ (useful for sparse gradients), initialization bias becomes larger and bias correction prevents adverse early-training effects; without bias correction, instabilities appear in early epochs in the reported VAE study (\cite{adam}).

\paragraph{Regularization coupling in adaptive methods.}
AdamW argues that for SGD, L2 and weight decay equivalence requires coupling the regularizer strength to the learning rate, and that for adaptive methods the equivalence fails; hence it advocates decoupling weight decay from the adaptive gradient transformation (\cite{adamw}).

\paragraph{Assumption-sensitive guarantees for SignSGD.}
The SignSGD majority vote analysis states convergence rates for nonconvex objectives under explicit assumptions and further highlights a variance-reduction regime under unimodal and symmetric gradient noise conditions (\cite{signsgd}). Outside these assumptions, the provided theory does not apply (and no alternative guarantees are specified in the provided sources) (\cite{signsgd}).

\paragraph{Numerical method choice for Muon’s orthogonalization.}
Muon rejects explicit SVD (too slow) and a coupled Newton iteration (needs float32 for

stability, hence slower on modern GPUs) while claiming Newton--Schulz iterations can run stably in bfloat16, motivating the chosen orthogonalization approach (\cite{muonblog}).

\section{Research Insights & Open Questions}

\subsection{Synthesized insights (5--10), each grounded}

\begin{enumerate}

\item \textbf{Bias correction is not a cosmetic detail in Adam.} The Adam paper presents empirical evidence that removing bias correction can lead to early instability for $\beta_2$ close to $1$, precisely where robustness to sparse gradients is desired (\cite{adam}).

\item \textbf{Regularization is optimizer-dependent.} AdamW highlights that the common identification of L2 regularization with weight decay breaks for adaptive methods and even for SGD can imply coupling between decay and learning rate, motivating decoupled weight decay steps (\cite{adamw}).

\item \textbf{Communication constraints can be built into the update rule.} SignSGD's majority vote scheme is explicitly designed so that both worker$\to$server and server$\to$worker communication can be 1-bit, while retaining nonconvex convergence rates under stated assumptions (\cite{signsgd}).

\item \textbf{Structured preconditioning can bridge diagonal adaptivity and second-order geometry.} Shampoo's left/right preconditioners and inverse fractional matrix powers yield a vectorized preconditioned update while exploiting tensor structure to avoid full-matrix costs (\cite{shampoo}).

\item \textbf{Orthogonalization can be interpreted as a geometry-motivated modification.} Muon frames its NS iteration as approximating the nearest semi-orthogonal matrix to the raw update, and it motivates orthogonalization by observing ill-conditioned (high-condition-number) update matrices in transformer parameters, speculating that orthogonalization amplifies `rare directions'' (\cite{muonblog}). \item \textbf{Muon's method selection is systems-aware.} Muon's source justifies Newton--Schulz by a claimed ability to run stably in bfloat16, whereas some alternatives require float32 and are slower on GPUs (\cite{muonblog}). \item \textbf{Scaling behavior can be parameterization-dependent.} MuP claims that in Maximal Update Parametrization many optimal hyperparameters remain stable across widths, enabling $\mu$Transfer and reframing `hyperparameter sensitivity'' as partly a parameterization artifact (\cite{mup}).

\item \textbf{Width-induced instabilities can be diagnosed and reproduced via transfer.} MuP describes reverse-$\mu$Transfer experiments that replicate large-

model instabilities on smaller models, suggesting a workflow where instability can be debugged at reduced cost (\cite{mup}).
\end{enumerate}

\subsection{Open questions (5--10), motivated directly}
\begin{enumerate}
\item \textbf{When do SignSGD's noise assumptions hold in modern training?} The SignSGD theory relies on unimodal and symmetric noise assumptions for certain advantages; characterizing when these assumptions hold (or fail) across models and batch sizes remains necessary to interpret guarantees (\cite{signsgd}).
\item \textbf{What is the right abstraction for ``regularization'' under adaptivity?} AdamW's critique of L2-vs-WD equivalence suggests further work on regularization operators that remain semantically consistent under adaptive rescaling (\cite{adamw}).
\item \textbf{How should Muon's orthogonalization interact with momentum and layer types?} Muon is described as SGD-momentum plus orthogonalization for matrix parameters, but best-practice design choices across heterogeneous parameter types are not specified in the provided sources (\cite{muonblog}).
\item \textbf{What are the stability limits of low-precision orthogonalization?} Muon claims NS is stable in bfloat16, while alternatives are not; a precise characterization of numerical stability vs iteration count and coefficient choice is only partially discussed (and would benefit from deeper analysis) (\cite{muonblog}).
\item \textbf{How should Shampoo-style preconditioners be engineered for scale?} Shampoo provides the mathematical update form and analysis, but practical system-level design choices (blocking, update frequency, distributed implementation) are only partially reflected in the provided sources here (\cite{shampoo}).
\item \textbf{Can MuP predict optimizer-specific scaling rules beyond SGD/Adam?} MuP states that its SGD intuitions can be made rigorous and that Adam-related intuitions are newer and to be formalized in future work, suggesting open theory for adaptive methods under MuP (\cite{mup}).
\item \textbf{Can orthogonalization be unified with Shampoo preconditioning under one theory?} Muon explicitly relates orthogonalized gradients to a Shampoo-derived form when accumulation is removed, inviting a unified view of when accumulation helps vs when orthogonalization suffices (\cite{muonblog,shampoo}).
\end{enumerate}

\section{Conclusion}
Across the provided sources, optimizer evolution can be read as a series of increasingly explicit answers to practical bottlenecks. Momentum augments SGD with temporal smoothing (\cite{lec2}); Adam adds adaptive normalization with bias correction to improve robustness and reduce tuning friction (\cite{adam}); AdamW clarifies that ``regularization'' is not optimizer-invariant and proposes decoupling to restore interpretability of weight decay under adaptivity (\cite{adamw}); SignSGD bakes communication constraints into the update by collapsing gradients to signs and formalizes behavior under explicit stochastic-noise assumptions (\cite{signsgd}); Shampoo targets curvature information with structured matrix preconditioning (\cite{shampoo}); and Muon revisits geometry via efficient orthogonalization, explicitly connecting to Shampoo-style ideas while arguing for low overhead at scale (\cite{muonblog,lec6}).

MuP adds a complementary scaling lens: it suggests that what we often attribute to ``optimizer sensitivity'' can be entangled with parameterization, and it provides a concrete mechanism for transferring hyperparameters across widths (\cite{mup}). Together, these works motivate optimizer evaluation as a joint design problem spanning algorithms, regularization semantics, numerical precision, and scaling laws.

\appendix
\section{Update-rule cheat sheet (appendix)}
This appendix consolidates key update forms. Where a detail is not specified in the provided sources, we say so.

\subsection{SGD and Momentum}
\begin{align}
\theta_t &= \theta_{t-1} - \alpha g_t \qquad (\text{SGD}) ; \cite{adam} \
v_t &= \beta v_{t-1} + (1-\beta) g_t,\qquad
\theta_t=\theta_{t-1}-\alpha v_t \qquad (\text{Momentum}) ; \cite{lec2}
\end{align}

\subsection{Adam and AdamW}
\begin{align}
m_t &= \beta_1 m_{t-1} + (1-\beta_1) g_t,\qquad
v_t = \beta_2 v_{t-1} + (1-\beta_2) (g_t\odot g_t) ; \cite{adam} \
\hat m_t &= \frac{m_t}{1-\beta_1^t},\qquad \hat v_t = \frac{v_t}{1-\beta_2^t} ; \cite{adam} \
\theta_t &= \theta_{t-1} - \alpha \frac{\hat m_t}{\sqrt{\hat v_t}+\varepsilon} ;

\cite{adam}
\end{align}
AdamW’s defining change is the explicit decoupling of weight decay from the adaptive gradient scaling; exact line-by-line pseudo-code is in the AdamW source (\cite{adamw}).

\subsection{SignSGD}
\begin{equation}
\theta_t = \theta_{t-1} - \alpha, \sign(\tilde g_t),
\end{equation}
with distributed majority vote variant
\begin{equation}
\theta_t = \theta_{t-1} - \alpha, \sign!\left(\sum_{m=1}^M \sign(\tilde g_t^{(m)})\right)
\end{equation}
as described in the SignSGD paper (\cite{signsgd}).

\subsection{Shampoo}
For matrix parameters:
\begin{equation}
W_{t+1} = W_t - \eta, L_t^{-1/4} G_t R_t^{-1/4},
\end{equation}
and equivalently in vectorized form $w_{t+1}=w_t-\eta H_t^{-1} g_t$, as used in the Shampoo analysis (\cite{shampoo}).

\subsection{Muon}
Muon applies an orthogonalization operator to the momentum update matrix, approximating the nearest semi-orthogonal matrix, implemented via Newton--Schulz iterations (\cite{muonblog}). Exact NS polynomial coefficients and iteration counts discussed in the Muon source are not reproduced here beyond what is stated there (\cite{muonblog}).

\section{Glossary (appendix)}
\begin{description}
\item[Bias correction:] Adam’s correction of moment estimates to account for initialization at zero, producing $\hat m_t,\hat v_t$ (\cite{adam}).
\item[Weight decay vs L2 regularization:] AdamW emphasizes that the equivalence between L2 penalty and multiplicative weight decay can fail for adaptive methods and can couple hyperparameters even for SGD (\cite{adamw}).

\item[Preconditioning:] Transforming gradients by an (approximate) inverse curvature/second-moment matrix, e.g., Shampoo’s structured $H_t^{-1}$ (\cite{shampoo}).
\item[Majority vote:] In distributed SignSGD, the server returns a 1-bit sign decision formed from worker sign votes (\cite{signsgd}).
\item[Semi-orthogonal matrix:] A matrix with orthonormal columns or rows (i.e., $O^\top O=I$ or $OO^\top=I$), used in Muon’s orthogonalization objective (\cite{muonblog}).
\item[$\mu$Transfer:] MuP’s paradigm of tuning hyperparameters on a smaller model and transferring to a larger model under Maximal Update Parametrization (\cite{mup}).
\end{description}


\begin{thebibliography}{99}

\bibitem{adam}
Diederik P. Kingma and Jimmy Ba.
\newblock \emph{Adam: A Method for Stochastic Optimization}.
\newblock ICLR 2015 (arXiv:1412.6980). Provided source: \texttt{Adam.pdf}.

\bibitem{adamw}
Ilya Loshchilov and Frank Hutter.
\newblock \emph{Decoupled Weight Decay Regularization}.
\newblock ICLR 2019. Provided source: \texttt{AdamW.pdf}.

\bibitem{signsgd}
Jeremy Bernstein, Yu-Xiang Wang, Kamyar Azizzadenesheli, and Anima Anandkumar.
\newblock \emph{signSGD: Compressed Optimisation for Non-Convex Problems}.
\newblock Provided source: \texttt{SignSGD.pdf}.

\bibitem{shampoo}
Vineet Gupta, Tomas Koren, and Yoram Singer.
\newblock \emph{Shampoo: Preconditioned Stochastic Tensor Optimization}.
\newblock Provided source: \texttt{Shampoo.pdf}.

\bibitem{muonblog}
Keller Jordan.
\newblock \emph{Muon: An optimizer for hidden layers in neural networks}.
\newblock Blog post (PDF snapshot provided): \texttt{Muon_ An optimizer for hidden layers in neural networks _ Keller Jordan blog.pdf}.

\bibitem{mup}
Greg Yang et al.
\newblock \emph{Tensor Programs V: Tuning Large Neural Networks via Zero-Shot Hyperparameter Transfer}.
\newblock Provided source: \texttt{MuP.pdf}.

\bibitem{lec2}
Course Lecture Notes.
\newblock \emph{Lecture 2} (provided source: \texttt{Lecture 2.pdf}).

\bibitem{lec3}
Course Lecture Notes.
\newblock \emph{Lecture 3} (provided source: \texttt{Lecture 3.pdf}).

\bibitem{lec4}
Course Lecture Notes.
\newblock \emph{Lecture 4} (provided source: \texttt{Lecture 4.pdf}).

\bibitem{lec5}
Course Lecture Notes.
\newblock \emph{Lecture 5} (provided source: \texttt{Lecture 5.pdf}).

\bibitem{lec6}
Course Lecture Notes.
\newblock \emph{Lecture 6} (provided source: \texttt{Lecture 6.pdf}).

\bibitem{lec7}
Course Lecture Notes.
\newblock \emph{Lecture 7} (provided source: \texttt{Lecture 7.pdf}).

\end{thebibliography}

\end{document}