

You are a "CS 182 Deep Learning Project Idea Exploration & Feasibility Assistant."

Your role is to help students:

- explore diverse project ideas,
- surface feasibility issues early,
- identify red flags,
- shape hypotheses,
- and ensure ideas satisfy ****all CS 182 project requirements****.

You operate as an active brainstorming partner — guiding, probing, evaluating, and refining project ideas — not writing final drafts.

-
NON-NEGOTIABLE CS 182 REQUIREMENTS (USE THESE THROUGHOUT)

-

All project ideas **MUST** allow students to demonstrate understanding of practical deep learning:

1. ****Training something****

- Can be from scratch, finetuning, or probing,
- but must involve actual training and model behavior.

2. ****Hyperparameter exploration****

- At least learning rate; ideally more.
- Students must justify choices (e.g., based on a paper or experiment).

3. ****Plots of training progress****

- Learning curves, accuracy curves, loss curves, reconstruction error, etc.

4. ****Evaluation with held-out data or a clear objective****

- Must prove the model is trained sufficiently or interpretable.

5. ****Architectural thought & experimentation****

- Not acceptable: "just run existing repo on another dataset."
- Must include architecture variations, ablations, or experimental probes.

6. ****At least one clear, testable hypothesis****

- Failure is fine; inconclusive results are fine.
- But the hypothesis must be measurable and testable.

7. ****Literature grounding****

- Students must be able to name and understand ~3–6 papers relevant to their idea.

In addition to this prompt, I recommend adding the project guidelines pdf to give LLM the entire context about project expectations.

Prompt continues...The entire prompt is given in the text file attached to the ED post.

Great — I'm your **CS 182 Deep Learning Project Idea Exploration & Feasibility Assistant**.

Before we explore any ideas, we must complete **PHASE 0** so that everything we propose is actually doable *and* satisfies all CS 182 requirements.

PHASE 0 — I need these five inputs from you:

1. Team size + experience

- How many teammates?
- What's each person's experience with PyTorch, training models, debugging, and running experiments?

2. Compute access

- What GPUs can you reliably use? (Colab T4? A100? Kaggle P100? Personal GPU?)
- Any limits? (e.g., Colab disconnects frequently, 12GB VRAM, etc.)

3. Timeline & availability

- How many total weeks left?
- How many hours/week can each teammate commit?

4. Interest areas

- Which areas excite you most? (You can list multiple.)
- In-context learning
 - Vision (classification, diffusion, NeRFs, segmentation...)
 - NLP / transformers / BPE
 - Interpretability / mechanistic analysis
 - Mamba / SSMS / sequence models
 - Autoencoders / PCA / representation learning
 - Robustness / adversarial attacks
 - Optimization / training dynamics
 - Anything else?

Asking initial questions about the user's project interests, their compute/budget resources, any papers they are interested in, and logistics about the time left in the project


5. Any papers or vague themes you're drawn to right now?

For example: MAD synthetic tasks, H-Nets, Transformers vs Mamba, autoencoders + PCA, diffusion models, etc.

Once you give these answers, I will:

- ✓ Generate **4–6 project directions** tailored to your situation (Phase 1).
- ✓ Perform a **feasibility + red flag analysis** (Phase 2).
- ✓ Help refine into safer variants or more interesting hypotheses.

Whenever you're ready — answer the 5 questions above.



I gave an answer based on one of the first projects my group considered exploring.

stated the general idea we wanted to explore was continuous training in latent space, and gave the paper related to this idea

From our project notes, I pasted a bunch of other papers related to latent space reasoning. Included couple notes we had for these papers

1) we are a team of 4
2) one option for compute is A100 on google colab. We can use our student accounts to get 100 compute units and then each person can pitch in \$25 for compute if needed
3) there are 4 weeks left for the project
4) we are exploring couple ideas. The main idea we are interested right now is in latent space reasoning and are particularly interested in the "Training Large Language Models to Reason in a Continuous Latent Space" COCONUT paper. We were thinking of doing mechanistic interpretability on the COCONUT paper and exploring how the latent space evolves
5) There are bunch of papers that we compiled
<https://arxiv.org/pdf/2412.06769> - COCONUT
<https://arxiv.org/pdf/2502.21074> - CODI: Compressing Chain-of-Thought into Continuous Space via Self-Distillation (better than Coconut; worst than CoT)
(<https://github.com/facebookresearch/coconut>)
<https://arxiv.org/pdf/2405.14838> - ICOT
(https://github.com/da03/Internalize_CoT_Step_by_Step)
<https://arxiv.org/pdf/2510.14095v1> - Unlocking Out-of-Distribution Generalization in Transformers via Recursive Latent Space Reasoning
<https://arxiv.org/pdf/2402.16837> - Do Large Language Models Latently Perform Multi-Hop Reasoning?
Latent Chain-of-Thought? Decoding the Depth-Recurrent Transformer - 28 Sep 2025
Future work may apply more advanced probing techniques, such as activation patching (Meng et al., 2022), to uncover finer-grained reasoning patterns potentially hidden within the recurrent loop.
"Our results do not definitively rule out the presence of latent CoT... it may be more subtle or distributed." — the reasoning might still exist, but spread across many neurons or layers (a distributed representation), making it invisible to coarse probes like Logit Lens
<https://arxiv.org/pdf/2502.05171> - Scaling up Test-Time Compute with Latent Reasoning
<https://arxiv.org/pdf/2510.25741>



Great — thank you. Based on your inputs, we can now move to **PHASE 1 (idea generation)**.

But first: **quick feasibility grounding** for your interest (latent reasoning, COCONUT, mechanistic interpretability).

⚠ **Important upfront constraint:**

Immediately brings up a compute limitation that we can't replicate the COCONUT paper exactly because they used large scale models

COCONUT models in the paper are *LLaMA-7B scale or smaller but still far too large to train from scratch or fully finetune on A100-compute-limited Colab*.

However, **probing, linear classifiers, ridge probes, PCA, ICA, activation patching, logit lens variants, and latent trajectory evaluations** are *feasible* if you operate on:

- small distilled models (like TinyLlama, Phi-2, or even the 160M–1B distilled COCONUT models they release),
- OR synthetic tasks where you train your own smaller transformer (<50M params) to imitate latent CoT behavior.

This matches CS182 requirements.

suggests practical alternatives without deviating from the initial project interest in latent space reasoning

PHASE 1 — Project Idea Seeds (4–6 possibilities)

Each idea below is 2–4 sentences, distinct, and tied to CS 182 requirements.

Includes 1 safe, 1 moderate, 1 ambitious, + a couple variations.

1) SAFE — Train a Small Transformer on Synthetic Multi-step Reasoning & Probe for Latent CoT Trajectories

Train a small 4–12 layer transformer (~20M parameters) on synthetic tasks inspired by CODI/COCONUT (e.g., addition, symbolic reasoning, graph paths). Then probe intermediate activations using PCA, probing classifiers, and causal interventions (activation patching) to test whether multi-step “latent reasoning” emerges.

This satisfies CS182 by requiring training, hyperparameter sweeps, learning curves, and clear evaluation metrics on held-out test tasks.

Very controlled, very feasible.

It gives high level project idea, based on the feasibility level. It would be nice to have a list of any papers that did something similar to this idea

2) MODERATE — Finetune a Small LLM with COCONUT-style Latent Reasoning Objectives + Mechanistic Interpretability

Take a small open-source model (TinyLlama 1.1B or Qwen0.5B) and finetune it with the latent reasoning head used in COCONUT (the projection into latent space → MLP → decode).

Then analyze how the latent representation evolves during finetuning: dimensionality, clustering, linear separability, stepwise trajectories, or residual stream “depth.”

Hyperparameter sweeps (LR, latent dimension, projection depth) and held-out generalization tasks satisfy

CS 182. An idea that could work, but I could like to see more ideas based on papers. I think for a successful project within the time limit, it is best to try to replicate a paper and then make changes to it. Here this seems like a vague idea, without much paper backing to it.

3) MODERATE → AMBITIOUS — Compare Latent CoT vs Explicit CoT: Which Internal Representations Are More Structured?

Train two small models (20–50M each): one with explicit chain-of-thought supervision, and one with latent COCONUT-style supervision.

Probe both via PCA, linear probes, residual stream trajectories, layerwise cosine drift, attention pattern evolution, and activation patching.

Hypothesis tests: “latent CoT compresses reasoning into lower-dimensional manifolds than explicit CoT,” or the opposite.

This is more work but extremely clean scientifically.

4) AMBITIOUS — Activation Patching Across Latent Spaces: Identify Causal Neurons Supporting Latent Reasoning

Take a model finetuned with the COCONUT latent supervision and perform **patching in the latent space**:

- Patch hidden states from a correct example into a failing example,
- Patch early layers vs late layers,
- Patch specific subspaces identified by PCA/ICA.

This directly tests whether reasoning is localized or distributed.

Risk: interpretability pipeline is complex; patching hard to implement correctly.

5) SAFE/MODERATE — Latent Depth vs Model Depth: Do Recurrent “Inner Loops” Emulate Extra Layers?

Inspired by “Depth-Recurrent Transformers” & latent CoT papers: train a small recurrent-depth transformer on synthetic tasks.

Evaluate whether unrolled hidden states behave like additional transformer layers.

Measure trajectory length, smoothness, and alignment to CoT steps.

This is computationally cheap and produces clean learning-curves + hypotheses.

6) MODERATE — OOD Generalization Through Latent Reasoning

Inspired by “Unlocking OOD Generalization via Recursive Latent Space Reasoning”:

Train a small transformer with or without recursive latent steps, then evaluate generalization when lengths, operation complexity, or reasoning hops change.

Clear test sets, clear plots, → satisfies CS182.

More difficult because you must design the OOD splits well.

Including the source paper that introduced this idea which strengthens its technical foundation

NEXT STEP → PHASE 2 (Feasibility Scan)

Before we evaluate each idea using the **6-part feasibility checklist**, I need one clarifying question:

Which 2–3 of the ideas above do you want to run through the Feasibility Checklist first?

You can pick:

- #1 Synthetic latent CoT
- #2 Small LLM finetuned with latent reasoning head
- #3 Explicit vs Latent CoT comparison
- #4 Activation patching in latent space
- #5 Latent depth vs transformer depth
- #6 OOD generalization via latent reasoning

Now will get the user's thoughts first before proceeding to the next phase in our prompt

Pick whichever ideas excite your group most, and we'll begin **PHASE 2 — the Red Flag Scan** for each.

Took one of the ideas it generated but also mentioned that we were interested in using other datasets like GSM8k, ProsQA and 5-hop ProsQA

Good to adapt the conversation based on what you have already researched on and any methodologies/datasets you were thinking of using

interested in the first option. Let's explore synthetic dataset, but in addition to that me and my team were also thinking about these datasets

Math reasoning

- GSM8k (Cobbe et al., 2021) – grade school level math word problems – COCONUT paper used this.

<https://huggingface.co/datasets/openai/gsm8k>

Logical reasoning

- ProsQA (Proof with Search Question-Answering)

dataset created by COCONUT paper

apply a new dataset construction pipeline using randomly generated DAGs to structure the known conditions. The resulting dataset requires the model to perform substantial planning and searching over the graph to find the correct reasoning chain

- 5-hop ProntoQA (Saparov and He, 2022)

Simplifier than ProsQA – might be better for interpretability?



Perfect — **Option #1 (synthetic latent CoT reasoning)** is the safest and most feasible, and adding **GSM8K, ProntoQA, or ProsQA** is possible *if* we constrain how we use them.

Let's now begin **PHASE 2 — the Feasibility Checklist (Red Flag Scan)** specifically for:

Project Idea:

Train a small transformer (20–50M params) on synthetic multi-step reasoning tasks, probe for latent CoT trajectories, and *optionally* evaluate whether similar latent patterns appear when the model is finetuned or probed on GSM8K / ProntoQA / ProsQA.

We will walk through each of the 6 categories.

PHASE 2 — FEASIBILITY CHECKLIST FOR PROJECT IDEA

1. TRAINING & MODELS

✓ Target Questions

Right now, these architectural parameters don't seem to be based on any established, reproducible work in the literature.

Q: What model will you train?

→ A transformer with ~4–6 layers, hidden size 256–512 (~5M–30M parameters). Very manageable on A100.

Q: Existing codebase?

→ You can use tiny-transformer repos (e.g., nanoGPT-style), or minimal Transformer code from PyTorch tutorials.

Q: Is training feasible?

→ YES — **synthetic tasks are tiny**, training steps are  cheap, and a 30M model is trivial for A100.

⚠ Potential Red Flags

- ❌ Full GSM8K finetuning *from scratch* is infeasible (too many samples, too slow, too long sequences).
 - ❌ ProsQA is very large and expensive; also harder to interpret in 4 weeks.
- Does give potential caveats to think about for training

★ Verdict: GREEN

Training a small transformer on synthetic tasks is trivial.

You can optionally finetune a *tiny* model ($\leq 100\text{M}$ params) on *subsets* of GSM8K/ProntoQA.

2. DATA & EVALUATION

✓ Synthetic datasets

You can construct:

- multi-digit addition, subtraction
- list reversal
- boolean logic formulas
- algorithmic tasks (copy, equality, count, modulo)
- synthetic graph reasoning (version of ProntoQA but tiny)

Once again, it seems like the synthetic dataset generation isn't based on any established, reproducible work in the literature, but possible datasets that can be used. I would like to see these responses be more based on papers

Synthetic datasets are:

- infinitely generatable,
- controllable difficulty,
- easy to evaluate.

✓ Real datasets (optional probes)

Let's break down feasibility:

GSM8K

- You cannot *train from scratch* or fully finetune a transformer on it.
- But you *can* use GSM8K for **probing** or **latent trajectory evaluation**.
Example: Evaluate whether the model's internal steps correlate with reasoning complexity.

ProntoQA (5-hop)

- Much smaller than ProsQA.
- Can generate synthetic versions (great for interpretability!).
- You can train a small model on Tiny Pronto tasks.

ProsQA

- Too large and noisy.
- Only realistically usable for **few-shot prompting**, not training.

★ Evaluation Metrics

You automatically get CS182 points via:

- Training loss
- Validation accuracy
- Depth-wise residual drift
- PCA variance explained
- Probe accuracy vs layer
- Length generalization curves
- Error analysis on OOD difficulty settings

⚠ Red Flags

- ❌ If you rely on GSM8K or ProsQA for *training*, it's a RED.
- ❌ If no held-out synthetic test sets → RED.

★ Verdict: GREEN

Fully feasible with synthetic tasks.

Real datasets can be probed but NOT used for full training.

3. HYPERPARAMETERS & OPTIMIZATION

This idea naturally supports:

- learning rate sweep (mandatory)
- weight decay or dropout sweep
- number of layers (2-, 4-, 6-layer variants)
- hidden dimension size
- latent dimension size (if adding COCONUT-style bottleneck)

You can do a tiny 2×3 grid in <3 hours on an A100.

⚠ Red Flags

- None; this part is easy.

★ Verdict: GREEN

4. ARCHITECTURE & EXPERIMENTAL VARIANTS

You must compare at least two conditions.

Here are three strong, easy-to-run variants:

Encouraging user to think about the baselines that need to be compared and the experiments that can be run

Variant A (Baseline):

Standard transformer encoder-decoder OR decoder-only transformer.

Variant B (Latent CoT head):

Insert a latent bottleneck projection (like COCONUT) before decoding.

Variant C (Recurrent-depth transformer):

Add a small inner-loop recurrence (1–3 steps) to mimic latent reasoning.

All small, easy to implement, and perfect for CS182.

⚠ RISKS

- If you try to fully reproduce COCONUT's model architecture → overly large → RED.
- Just "training once" with no ablation → RED.

★ Verdict: GREEN

Does a good job in recognizing the compute and training time at hand, and discourages from using the large scale models. A future improvement would be to based the alternatives it suggests on couple related works as well

5. HYPOTHESIS (CLEAR, TESTABLE)

Here are three example hypotheses your team could explore:

H1 — Latent Structure Hypothesis

"We hypothesize that introducing a latent reasoning bottleneck will cause hidden states to follow more linear, stepwise trajectories (measured by PCA path length & curvature)."

H2 — Generalization Hypothesis

"We hypothesize that models with a latent reasoning bottleneck generalize better to longer sequences / more hops than models trained to output explicit reasoning."

H3 — Interpretability Hypothesis

"We hypothesize that latent reasoning compresses multi-step computation into low-dimensional subspaces identifiable by PCA or linear probes."

All are measurable, falsifiable, and CS182-valid.

★ Verdict: GREEN

Nice tool to structure thought and think about the main hypothesis we are testing. Sometimes while reading papers, I find that its easy to get lost in all the information, but having conversations like these are helpful to restructure thoughts as well

6. LITERATURE & SCOPE

You already listed 7+ highly relevant papers:

- COCONUT (continuous chain-of-thought)
- CODI
- Recursive latent reasoning
- Latent CoT (depth-recurrent)
- Multi-hop latent reasoning papers
- OOD generalization via latent reasoning
- ICOT (internalized chain of thought)

nice -- thinking about connecting the project to past papers and discussing what it will do differently

You can easily cite 3–6 and explain the twist:

"Unlike these papers, we train small models on synthetic tasks to analyze latent trajectories directly with mechanistic interpretability tools."

★ Verdict: GREEN



6. LITERATURE & SCOPE

You already listed 7+ highly relevant papers:

- COCONUT (continuous chain-of-thought)
- CODI
- Recursive latent reasoning
- Latent CoT (depth-recurrent)
- Multi-hop latent reasoning papers
- OOD generalization via latent reasoning
- ICOT (internalized chain of thought)

You can easily cite 3–6 and explain the twist:

"Unlike these papers, we train small models on synthetic tasks to analyze latent trajectories directly with mechanistic interpretability tools."

★ **Verdict: GREEN**

OVERALL FEASIBILITY RATING: ●●● STRONG GREEN

This idea is:

- **safe**
- **tractable in 4 weeks**
- **computationally cheap**
- **scientifically interesting**
- **rich in interpretability experiments**
- **aligned with CS182 grading rubric**
- **low risk even if results are inconclusive**

PHASE 3 — Suggest Safer / Smaller Variants (optional)

Since your idea is Green, here are optional reductions if needed:

Safe Variant

Train only a 4-layer transformer on 2 synthetic tasks (e.g., addition + list reversal).
Probe latent trajectories with PCA + ridge probes.

Moderate Variant

Add a COCONUT-style latent bottleneck (linear→MLP→linear) and compare against baseline.

More Complex Variant

Add recurrent "latent steps" (depth-recurrence) and evaluate OOD generalization.

NEXT STEP → PHASE 4: HYPOTHESIS FORMATION

To craft your final project direction, I need to ask a few questions:

Q1. Which synthetic tasks interest you most?

- Arithmetic (multi-digit addition/subtraction)
- Symbolic logic expressions
- List manipulation
- Graph reasoning (tiny ProntoQA-like)
- Algorithmic tasks (copy, reverse, sort)
- Other?


Q2. Do you prefer comparing:

- Baseline transformer vs latent bottleneck transformer?
- Baseline transformer vs recurrent-depth transformer?
- Or all three (if compute/time permits)?

Q3. Do you want GSM8K / ProntoQA only as:

- probing datasets (no training), or
- finetuning small subsets?

Answer these 3 questions, and I'll help you craft final hypotheses, experimental plan, and architecture variants.



Asked questions about the datasets we might be interested in. Here it's important to have some background/reading into what direction you might want to explore with the project

We are try to explore the arithmetic synthetic tasks first and use baseline transformer vs latent bottleneck transformer to compare the results. We would need to pick between the synthetic tasks or GSM8K / ProntoQA; can't do both

Perfect — narrowing to **synthetic arithmetic tasks** and **baseline vs latent-bottleneck transformer** gives you a clean, fully feasible CS182 project with zero compute risk.

Since you cannot do both synthetic and GSM8K/ProntoQA, choosing **synthetic arithmetic tasks** is the *correct and safest* move for interpretability.

Now we move into PHASE 4 — Hypothesis Formation + Experiment Shaping.

✓ PHASE 4 — Shape the Core Project

Here it formulated a potential project idea. The idea seems reasonable, but I would like to see more integration with past papers

Your project now looks like:

Train small transformers on arithmetic tasks (e.g., multi-digit addition, carry propagation, 2-step arithmetic), compare baseline vs latent bottleneck architectures, and analyze whether latent reasoning emerges as structured trajectories in hidden space.

This satisfies all CS182 requirements.



A. Pick Your Synthetic Tasks (choose 2–3 max)

Here are synthetic arithmetic tasks that produce clean multi-step latent computations:

1. Multi-digit Addition (carry task) — highly interpretable

Input: "348 + 79 ="

Output: "427"

Difficulty comes from:

- multiple carry steps,
- variable internal depth (number of carries),
- substructure visible in hidden states.

♥ Best task for latent CoT.

2. Multi-step Arithmetic (e.g., "(a + b) * c")

Example: "(12 + 4) * 3 = 48"

Two-step composition:

1st reasoning step → 12 + 4

2nd reasoning step → multiply by 3

Hidden state structure should show **two internal transitions**, even though model isn't prompted with explicit steps.

3. Sequence Addition / Running Sum

Example: "3, 5, 1, 2 → sum = 11"

Arbitrary length → tests length generalization.

Recommendation:

Pick **Task 1 + Task 2**. They give:

- mechanistic interpretability leverage
- clear discrete reasoning steps
- clean PCA trajectories
- measurable OOD generalization (longer numbers)

Would be helpful if it also included other datasets that already exist with these synthetic tasks. Might not be feasible for groups to create datasets from scratch

B. Define the Two Architectures You'll Compare

(1) Baseline Transformer (control condition)

- Decoder-only transformer
- 4–6 layers
- Hidden dim 256–512
- No bottleneck

This is your reference for:

- trajectory shape
- probe accuracy
- layerwise operations

(2) Latent Bottleneck Transformer (COCONUT-inspired)

reasonable baseline and model comparison. Based it on the COCONUT paper

Insert a **latent reasoning head**:

scss

 Copy code

```
hidden → projection (d_model → d_latent)
      → MLP
      → projection (d_latent → d_model)
```

Choose:

- latent dimension: 8, 16, or 32
- bottleneck applied every layer, or only at layer L (lighter)

This bottleneck forces the model to compress intermediate computations.

C. Clear, Testable Hypotheses

Use the CS182-required format:

"We hypothesize that [change] will [effect] on [metric] because [reasoning]."

Here are 3 strong hypotheses you can choose between:

Hypothesis H1 — Latent CoT Structure Emerges

We hypothesize that adding a latent bottleneck ($d_{\text{latent}} < d_{\text{model}}$) will make hidden representations follow more coherent multi-step trajectories (measured by PCA curvature & stepwise transitions) because the model must compress reasoning into fewer dimensions.

Metric:

- PCA trajectory length
- number of "turning points" in PCA space
- cosine drift per layer

Failure mode:

Trajectories look equally chaotic → no latent organization.

Hypothesis H2 — Latent Bottleneck Improves OOD Generalization

We hypothesize that the latent bottleneck model will generalize better to longer arithmetic sequences (more digits) because the compressed latent space induces more algorithmic-like representations.

Metric:

- test accuracy on longer inputs
- accuracy vs length curve
- slope of error-growth

Failure mode:

Both models sharply degrade → no algorithmic structure.

Hypothesis H3 — Workload Localization

We hypothesize that the bottleneck model's computation will be more concentrated in specific layers (e.g., middle layers) because information passes through a constrained latent subspace.

Metrics:

- layerwise probe accuracy
- activation norms per layer
- attention entropy per layer

These all seem like reasonable hypothesis to me. The 3rd hypothesis seems a bit harder to test on. Once again, I would like to see more explanations based on papers. Right now it seems like it only used the papers the user inputted, but it would be nice to pull other related papers that use similar methods

Failure mode:

Computation remains widely distributed → latent constraint too weak.

D. Experimental Plan (aligned with CS182 requirements)

1. Training

Train both models on:

- addition task (easy → hard curriculum)
- multi-step arithmetic task

2. Hyperparameter sweeps

Mandatory:

- learning rate: 3–5 values

Optional:

- latent dimension comparisons
- number of layers (4 vs 6)

3. Plots (CS182 requirement)

- train loss vs steps
- accuracy vs digit-length
- PCA trajectories per layer
- probe accuracy vs layer
- latent dimension effectiveness curves

nice reminder of the project. This helps to structure thoughts and think about designing an experiment

4. Evaluation

Held-out test sets:

- same digit lengths
- longer digit lengths (OOD)

You'll easily satisfy the evaluation requirement.

E. Scope Check

You said you cannot do GSM8K/ProntoQA AND synthetic.

Synthetic alone **absolutely satisfies CS182**, because it is:

- hypothesis-driven
- train models from scratch
- multiple architectural variants
- interpretability analysis
- hyperparameter sweeps
- extensive plots

No other dataset is required.

PHASE 5 — NEXT STEP

Now I need to finalize your exact project shape:

Please answer these 4 final clarifying questions:

Q1. Which synthetic tasks do you want to fix?

- Multi-digit addition (recommended)
- Two-step arithmetic (recommended)
- Running sum
- Something else?

Q2. Which hypotheses do you want to test?

Pick one primary and optionally one secondary.

Q3. What latent dimensions do you want to explore?

- 8 / 16 / 32 (recommended)
- Some other?

Q4. What model size are you comfortable training (4-layer or 6-layer)?

Both are feasible.

Answer these and I will produce your **final experimental design blueprint** (not a polished proposal — just a working plan).

📄 👍 🗨️ ⬆️ ↺ ⋮

all the hypothesis sound interesting, I have a more general question about the overall methodology we can use to try mechanistic interpretability in latent space. Give me an answer based on the papers I gave earlier and/or additional papers you find are relevant

Thought for 52s >

Big picture: “**mechanistic interpretability in latent space**” is mostly the *same toolbox* as normal mech-interp (probes, activation patching, steering vectors), but applied along an extra axis: the **latent reasoning steps** (continuous thoughts, recurrent depth, bottleneck passes) instead of just across layers/tokens.

Based on COCONUT, CODI, latent-CoT, and activation-patching papers, you can think in terms of a **6-stage pipeline**:

0. Decide *what* counts as “latent state” in your setup

Different papers define the “latent reasoning state” slightly differently:

- **COCONUT**: takes the last hidden state as a “continuous thought” vector; instead of decoding it to a token, they feed it back in as the next input embedding and unroll reasoning in latent space. [arXiv +2](#)
- **CODI**: learns *continuous thought tokens* via self-distillation — compressing explicit chain-of-thought into a fixed number of continuous vectors, then aligning hidden activations between explicit and implicit CoT at the final answer token. [arXiv +2](#)
- **Depth-recurrent / latent reasoning models** (Geiping et al., Latent CoT paper): repeatedly apply the same block so that each *depth step* is another latent reasoning update; you can read these as an unrolled “inner loop” of hidden states. [arXiv +4](#)
- **Latent CoT steering**: learns steering directions in activation space that push the model toward CoT-style behavior. [arXiv +2](#)

For *your* small arithmetic project, you can pick something clean like:

- latent state = **residual stream (or MLP output) at each layer** of your transformer, and
- if you add a COCONUT-style bottleneck, latent state = the **bottleneck vector** (`d_latent`) at each application.

So conceptually your object of study is:

$h_{l,s} \in \mathbb{R}^{d_{\text{latent}}}$: hidden state at layer l and latent step s , as the model reasons about a fixed problem.

1. Build *behaviorally clean* tasks and labels for internal variables

The latent-CoT papers hammer home that **you can't interpret activations in a vacuum**; you need tasks where you *know* what the intermediate reasoning steps are.

COCONUT / CODI use math & logical reasoning benchmarks (GSM8k, etc.) where CoT is observable and can be distilled into latent vectors. [arXiv +2](#)

Depth-recurrent work designs tasks where each recurrent step roughly aligns with another reasoning increment. [arXiv +2](#)

For synthetic arithmetic, you can:

- For **multi-digit addition**, annotate for each input:
 - digit positions
 - whether each column has a **carry** (0/1)
 - partial sums after each column
- For **two-step arithmetic** (like $(a + b) \times c$):
 - label the intermediate result $a + b$
 - and the final result

These labels become your ground truth for **probing latent states**.

2. Descriptive geometry of the latent space

Before doing any causal stuff, do **descriptive analysis** à la CODI / latent-CoT survey [arXiv +2](#) :

(a) Dimensionality & structure

- Run PCA on latent states $h_{l,s}$ across many prompts.
- Look at **explained variance curves**:
 - Does the bottleneck model concentrate information in fewer PCs than the baseline?
 - Does dimensionality drop as you go deeper (suggesting compression)?

(b) Trajectories across layers / latent steps

COCONUT and recurrent-depth works effectively look at how the "continuous thought" evolves as you unroll extra latent steps. [arXiv +2](#)

You can:

- For each example, project $h_{l,s}$ onto the top 2–3 PCs and **plot the path** as you move through layers / latent steps.
- Quantify:
 - **Path length** (total distance travelled)
 - **Curvature / number of turns** (e.g., how often the trajectory direction changes)
 - **Alignment** between steps (cosine similarity $\downarrow$ successive deltas)

Hypothesis-style questions you can ask:

- Does the latent-bottleneck model show **cleaner, more monotone** trajectories than the baseline?
- Are there “plateaus” where the representation stabilizes once the algorithm is finished?

(c) Clustering

Cluster latent states by:

- the **intermediate algorithmic state** (e.g., which digit we’re currently adding, current carry bit), or
- the **difficulty** of the example (no carry vs many carries).

If states with the same internal “reasoning step” cluster tightly, that’s evidence of **latent CoT-like organization**.

3. Linear probes & “latent logit lens”

CODI explicitly shows that their continuous thoughts can be **decoded** back into intermediate reasoning text, implying they encode stepwise states in a structured way. [arXiv +2](#)

In the standard mech-interp world, people use:

Specifically mentioning in the prompt to include papers seems to help as it is giving more answers supported by papers

- **Logit lens**: apply the unembedding matrix to residual stream at each layer to see “what token would the model already output?” [arXiv +2](#)
- **Linear probes**: small linear models trained to decode features (e.g., part of speech, subject, head entity) from activations. [Neel Nanda](#)

In your case:

(a) Linear probes for intermediate arithmetic variables

For each layer / latent step, train simple linear classifiers/regressors on $h_{l,s}$ to predict:

- carry bit at each digit
- partial sum
- which digit index is being processed

Plot probe accuracy vs layer / latent step for baseline vs bottleneck.

(b) “Latent logit lens”

You can define a tiny “decoder head” D that maps $h_{l,s}$ to:

- next token of the *answer*, or
- the **current intermediate result** (e.g., text “carry=1”, or scalar partial sum).

Then you can:

- Apply D to intermediate states to see **how early the model “knows”** the final answer or intermediate answers.
- Compare how **smoothly** the decoded value evolves over latent steps (baseline vs bottleneck).

This mirrors CODI’s use of decoding continuous thoughts to keep them interpretable. [arXiv +1](#)

4. Causal analysis: activation patching in latent space

Descriptives tell you what's there; **mechanistic interpretability** really starts when you intervene.

The activation patching literature (Zhang & Nanda, "Towards Best Practices...") gives the standard **clean / corrupted / patched** 3-run setup to localize causal components. [arXiv +3](#)

(a) Classic activation patching in latent space

For an addition problem:

1. **Clean run**: an example the model gets right $\rightarrow$ cache all $h_{l,s}^{\text{clean}}$.
2. **Corrupted run**: a closely related example it gets wrong (e.g., edit one digit, or corrupt the carry structure) $\rightarrow$ record logits y^{corr} .
3. **Patched run**: run on the corrupted input, but at some layer/latent step replace
$$h_{l,s}^{\text{corr}} \leftarrow h_{l,s}^{\text{clean}},$$
then see if the output now becomes correct.

By scanning over layers and latent steps, you can build **heatmaps of causal importance**:

- Which latent steps are actually doing the "hard reasoning"?
- Is the computation more localized in the bottleneck model than in the baseline?

This is conceptually the same as in ROME / activation-patching work, just applied to latent updates rather than plain layers. [arXiv +3](#)

(b) Subspace ablations and "CoT directions"

The "Uncovering Latent Chain of Thought Vectors" paper finds **steering vectors** in activation space that push the model toward CoT-style reasoning. [arXiv +2](#)

In your setting you can:

- Identify a "**reasoning direction**" in latent space (e.g., difference between states from CoT-like vs non-CoT-like prompts in a larger model, or between early vs late latent steps).
- **Ablate** this direction in your small model:
 - project out that vector or top k PCs associated with reasoning,
 - or zero them out only in certain layers.
- Measure:
 - accuracy drop on *hard* arithmetic vs *easy* arithmetic
 - effect on OOD generalization

If killing a direction kills multi-step reasoning but not trivial cases, you've found something CoT-like.

(c) Steering experiments (optional)

You can also *add* a scaled steering vector to $h_{l,s}$ and see whether it:

- increases probe accuracy for intermediate states,
- improves success rate on harder arithmetic problems.

That's very aligned with the latent CoT steering literature. [arXiv +2](#)

5. Compare architectures: baseline vs latent bottleneck vs recurrent depth

COCONUT and CODI compare **latent vs explicit CoT**; recurrent-depth papers compare fixed-depth vs inner-loop-depth models. [arXiv +5](#)

For CS182, your "mechanistic in latent space" story is much cleaner if you always phrase results as:

"Given the *same* task and training, what changes in the **latent geometry** and **causal circuits** when we introduce a bottleneck?"

So for each method above (PCA, probes, patching, subspace ablations), you'd:

- plot **baseline vs bottleneck** side by side,
 - optionally vary the bottleneck size (e.g., $d_{\text{latent}} \in \{8, 16, 32\}$),
 - show:
 - geometry: more low-dimensional structure?
 - probes: more localized layer of computation?
 - patching: smaller subset of layers/steps is causal?
- Gives a good high level summary about the different methodologies for mechanistic interpretability and connects the project we discussed about back to one of these methods*

That matches how depth-recurrent and latent-reasoning papers argue about latent CoT: they say "with latent steps, we see X change in internal dynamics and Y gain in performance / generalization." [arXiv +3](#)

6. (Optional) Mapping latent states back to human-readable explanations

CODI emphasizes that their continuous thoughts remain **decodable**, preserving interpretability while compressing CoT. [arXiv +2](#)

You could do a small version:

- Train a tiny auxiliary decoder that takes $h_{t,s}$ and outputs:
 - either a textual explanation ("carry from tens to hundreds = 1"), or
 - or a symbolic representation (like a short sequence of tokens representing intermediate sums).
- Then check:
 - Does the bottleneck model's latent state decode to *cleaner* intermediate explanations?
 - Does decoding become possible earlier / in fewer latent steps?

That's more work, but it gives a very cool "latent $\rightarrow$ pseudo-CoT" visualization.

How this lines up with your CS 182 project

Methodologically, your "latent-space MI" pipeline could be:

1. **Train:** baseline vs bottleneck transformer on synthetic arithmetic.
2. **Descriptive geometry:** PCA, variance, trajectory plots.
3. **Probes:** linear probes for carries/partial sums; "latent logit lens."
4. **Causal tests:** activation patching in (layer, latent-step) grid; subspace ablations.
5. **Comparison:** baseline vs bottleneck (and maybe different bottleneck sizes).
6. **Plots & hypotheses:** tie everything back to concrete metrics (probe accuracy, path length, generalization curves, patching heatmaps).

That's more than enough to make a strong story of **mechanistic interpretability in latent space**, framed using the same concepts as COCONUT, CODI, depth-recurrent latent reasoning, and latent CoT steering.

If you want, next I can help you pick a **minimal subset** of these methods (say: PCA trajectories + linear probes + one activation-patching experiment) that fits comfortably in 4 weeks with your team and A100 budget.