

One-Page Cheatsheet: When the Quadratic Approximation Breaks Down

Core Idea:

The update rule $\Delta = -P_t \cdot g_t$ assumes the loss is well-approximated locally by a quadratic:

$$f(\theta + \Delta) \approx f(\theta) + g^T \Delta + \frac{1}{2} \Delta^T H \Delta$$

This only works if curvature (the Hessian) changes slowly in a small neighborhood.

Failure Condition:

The quadratic model fails when the Hessian is not stable:

When $H(\theta + \delta)$ differs sharply from $H(\theta)$ even for tiny δ .

This means third derivatives are large and the bowl shape bends unpredictably.

Why This Breaks Optimizers:

- $H^{-1} g$ becomes unreliable.
- All P_t approximations (SGD, Adam, K-FAC, Shampoo) become misleading.
- Even simple SGD can overshoot because curvature varies too quickly.

Mnemonic:

"Quadratic $\approx$ Bowl. If the bowl twists fast, curvature math collapses."

Practical Red Flags:

- Sharp jagged loss landscape.
- Very large learning rates.
- Adam/RMSprop variance estimates jump around.
- Second-order methods diverge immediately.
- Training unstable even with small steps.

Key Condition to Remember:

Quadratic approximation $\approx$ good only when the Hessian is Lipschitz (locally smooth).

Fast-changing curvature kills quadratic reasoning.

2020 Second Summary:

For quadratic modeling and linear-transform optimizers to work, curvature must be stable.

If H changes too fast $\rightarrow$ higher-order terms dominate $\rightarrow$ the quadratic view breaks.